Cumulus NetQ CLI User Guide



Table of Contents

Cumulus NetQ CLI User Guide	8
NetQ Command Line Overview	9
CLI Access	9
Command Line Basics	10
Command Line Structure	11
Command Syntax	1]
Command Output	12
Command Prompts	13
Command Completion	13
Command Help	14
Command History	15
Command Categories	15
Validation Commands	15
Monitoring Commands	17
Configuration Commands	21
NETQ AGENT CONFIGURATION	21
CLI CONFIGURATION	23
EVENT NOTIFICATION COMMANDS	24
Trace Commands	25
Command Changes	30
New Commands	30
Modified Commands	31
Deprecated Commands	32
Monitor Overall Network Health	33
Validate Network Health	33
Validate the Network Fabric	33
Validate Device Status and Configuration	37
Validate Interface Status and Configuration	39
Create Filters for Provisioning Exceptions	40
View Network Details	43
Monitor Switch Hardware and Software	52
Monitor Switch and Host Hardware Information	52

View a Summary of Your Network Inventory	55
View Information about the ASIC on all Switches	56
View Information about the Motherboard in a Switch	58
View Information about the CPU on a Switch	61
View Information about the Disk on a Switch	63
View Memory Information for a Switch	64
View a Summary of Physical Inventory for the NetQ or NetQ Cloud Appliance	67
View Memory for the NetQ or NetQ Cloud Appliance	68
View Fan Health for All Switches	68
View PSU Health for All Switches	71
View the Temperature in All switches	72
View All Sensor Data	74
View All Sensor-related Events	76
View Interface Statistics and Utilization	76
Monitor Switch Software Information	81
View OS Information for a Switch	83
View License Information for a Switch	86
View Summary of Operating System on a Switch	88
Validate NetQ Agents are Running	89
Monitor Software Services	91
View All Services on All Devices	92
View Information about a Given Service on All Devices	96
View Events Related to a Given Service	99
Monitor Physical Layer Components	101
Monitor Physical Layer Inventory	101
View Detailed Cable Information for All Devices	103
View Detailed Module Information for a Given Device	106
View Ports without Cables Connected for a Given Device	108
View Ports with Cables Connected for a Given Device	108
View Components from a Given Vendor	110
View All Devices Using a Given Component	111
View Changes to Physical Components	112
Validate Physical Layer Configuration	115
Confirm Peer Connections	115
Discover Misconfigurations	117

Identify Flapping Links	122
Monitor Data Link Layer Devices and Protocols	123
Monitor LLDP Operation	123
View LLDP Information for All Devices	124
Monitor Interface Health	126
View Status for All Interfaces	127
View Interface Status for a Given Device	129
View All Interfaces of a Given Type	131
View the Total Number of Interfaces	133
View the Total Number of a Given Interface Type	133
View Changes to Interfaces	134
Check for MTU Inconsistencies	135
Monitor VLAN Configurations	137
View VLAN Information for All Devices	138
View VLAN Interface Information	139
View MAC Addresses Associated with a VLAN	140
View MAC Addresses Associated with an Egress Port	143
View the MAC Addresses Associated with VRR Configurations	144
Monitor MLAG Configurations	145
View MLAG Configuration and Status for all Devices	146
View MLAG Configuration and Status for Given Devices	148
Monitor Time Synchronization Status for Devices	149
Monitor Spanning Tree Protocol Configuration	152
Validate Paths between Devices	153
View Paths between Two Switches with Pretty Output	154
View Paths between Two Switches with Detailed Output	156
Monitor Network Layer Protocols	159
Monitor IP Configuration	159
View IP Address Information	161
View IP Neighbor Information	171
View IP Routes Information	174
Monitor BGP Configuration	179
View BGP Configuration Information	181
Validate BGP Operation	198
Monitor OSPF Configuration	200
View OSPF Configuration Information	201

Validate OSPF Operation	205
View Paths between Devices	206
Monitor Virtual Network Overlays	211
Monitor Virtual Extensible LANs	212
View All VXLANs in Your Network	213
View the Interfaces Associated with VXLANs	218
Monitor EVPN	219
View the Status of EVPN	220
View the Status of EVPN for a Given VNI	222
View EVPN Events	223
Monitor LNV	224
View LNV Status	224
View LNV Status in the Past	225
Monitor Linux Hosts	227
Monitor Container Environments	230
Use NetQ with Kubernetes Clusters	231
Requirements	231
Command Summary	231
Enable Kubernetes Monitoring	234
View Status of Kubernetes Clusters	234
View Changes to a Cluster	236
View Kubernetes Pod Information	237
View Kubernetes Node Information	247
View Container Connectivity	255
View Kubernetes Service Connectivity and Impact	256
View Kubernetes Cluster Configuration in the Past	258
Manage NetQ Agents	262
View NetQ Agent Status	262
Modify the Configuration of the NetQ Agent on a Node	264
Disable the NetQ Agent on a Node	266
Remove the NetQ Agent from a Node	266
Configure Logging for a NetQ Agent	267
Investigate NetQ Issues	271
Browse Configuration and Log Files	271
Check NetQ Agent Health	272
Diagnose an Event after It Occurs	274

Use NetQ as a Time Machine	276
Trace Paths in a VRF	278
Generate a Support File	278
Resolve MLAG Issues	280
Scenario: All Nodes Are Up	280
Scenario: Dual-connected Bond Is Down	283
Scenario: VXLAN Active-active Device or Interface Is Down	287
Scenario: Remote-side clagd Stopped by systemctl Command	291
CLI Early Access Features	295
Enable/Disable Early Access Features	295
What the NetQ Validation System Checks	299
NetQ Agent Validation Tests	299
BGP Validation Tests	299
CLAG Validation Tests	300
Cumulus Linux Version Tests	302
EVPN Validation Tests	302
Interface Validation Tests	303
License Validation Tests	304
LNV Validation Tests	304
Link MTU Validation Tests	304
NTP Validation Tests	305
OSPF Validation Tests	305
Sensor Validation Tests	306
VLAN Validation Tests	306
VXLAN Validation Tests	307
Validation Examples	307
Perform a NetQ Agent Validation	307
Perform a BGP Validation	308
Perform a BGP Validation for a Particular VRF	308
Perform a BGP Validation with Selected Tests	309
Perform a BGP Validation and Output Results to JSON File	310
Perform a CLAG Validation	314
Perform a CLAG Validation with Selected Tests	316
Perform a Cumulus Linux Version Validation	319
Perform an EVPN Validation	320
Perform an EVPN MAC Consistency Validation	321

Perform an EVPN Validation for a Time in the Past	322
Perform an EVPN Validation with Selected Tests	323
Perform an Interfaces Validation	326
Perform an Interfaces Validation for a Time in the Past	327
Perform an Interfaces Validation with Selected Tests	328
Perform a License Validation	329
Perform a Link MTU Validation	330
Perform an NTP Validation	330
Perform an OSPF Validation	331
Perform a Sensors Validation	333
Perform a VLAN Validation	334
Perform a VXLAN Validation	334

Cumulus NetQ CLI User Guide

This guide is intended for network administrators who are responsible for monitoring and troubleshooting the network in their data center environment. NetQ 2.x offers the ability to easily monitor and manage your data center network infrastructure and operational health. This guide provides instructions and information about monitoring individual components of the network, the network as a whole, and the NetQ software itself using the NetQ command line interface (CLI). If you prefer to use a graphical interface, refer to the Cumulus NetQ UI User Guide.

NetQ Command Line Overview

The NetQ CLI provides access to all of the network state and event information collected by the NetQ Agents. It behaves the same way most CLIs behave, with groups of commands used to display related information, the ability to use TAB completion when entering commands, and to get help for given commands and options. The commands are grouped into four categories: check, show, config, and trace.



(i) NOTE

The NetQ command line interface only runs on switches and server hosts implemented with Intel x86 or ARM-based architectures. If you are unsure what architecture your switch or server employs, check the Cumulus Hardware Compatibility List and verify the value in the **Platforms** tab > **CPU** column.

CLI Access

When NetQ is installed or upgraded, the CLI may also be installed and enabled on your NetQ server or appliance and hosts. Refer to the Install NetQ and Upgrade NetQ topics for details.

To access the CLI from a switch or server:

1. Log in to the device. This example uses the default username of *cumulus* and a hostname of *switch*.

<computer>:~<username>\$ ssh cumulus@switch

2. Enter your password to reach the command prompt. The default password is *CumulusLinux!* For example:

Enter passphrase for key '/Users/<username>/.ssh/id_rsa': <enter

CumulusLinux! here>

Welcome to Ubuntu 16.04.3 LTS (GNU/Linux 4.4.0-112-generic x86_64)

- * Documentation: https://help.ubuntu.com
- * Management: https://landscape.canonical.com
- * Support: https://ubuntu.com/advantage

Last login: Tue Sep 15 09:28:12 2019 from 10.0.0.14

cumulus@switch:~\$

3. Run commands. For example:

cumulus@switch:~\$ netq show agents

cumulus@switch:~\$ netq check bqp

Command Line Basics

This section describes the core structure and behavior of the NetQ CLI. It includes the following:

- Command Line Structure
- Command Syntax
- Command Output
- Command Prompts

- Command Completion
- Command Help
- Command History

Command Line Structure

The Cumulus NetQ command line has a flat structure as opposed to a modal structure. This means that all commands can be run from the primary prompt instead of only in a specific mode. For example, some command lines require the administrator to switch between a configuration mode and an operation mode. Configuration commands can only be run in the configuration mode and operational commands can only be run in operation mode. This structure requires the administrator to switch between modes to run commands which can be tedious and time consuming. Cumulus NetQ command line enables the administrator to run all of its commands at the same level.

Command Syntax

NetQ CLI commands all begin with **netq**. Cumulus NetQ commands fall into one of four syntax categories: validation (check), monitoring (show), configuration, and trace.

```
netq check <network-protocol-or-service> [options]
netq show <network-protocol-or-service> [options]
netq config <action> <object> [options]
netq trace <destination> from <source> [options]
```

Symbols	Meaning
Parentheses (Grouping of required parameters. Choose one.
Square brackets []	Single or group of optional parameters. If more than one object or keyword is available, choose one.

Symbols	Meaning
Angle brackets < >	Required variable. Value for a keyword or option; enter according to your deployment nomenclature.
Pipe	Separates object and keyword options, also separates value options; enter one object or keyword and zero or one value.

For example, in the **netq check** command:

- [<hostname>] is an optional parameter with a variable value named *hostname*
- <network-protocol-or-service> represents a number of possible key words, such as
 agents, bgp, evpn, and so forth
- <options> represents a number of possible conditions for the given object, such as around, vrf, or json

Thus some valid commands are:

- netq leaf02 check agents json
- netq show bgp
- netq config restart cli
- netq trace 10.0.0.5 from 10.0.0.35

Command Output

The command output presents results in color for many commands. Results with errors are shown in red, and warnings are shown in yellow. Results without errors or warnings are shown in either black or green. VTEPs are shown in blue. A node in the *pretty* output is shown in bold, and a router interface is wrapped in angle brackets (< >). To view the output with only black text, run the netq config del color command. You can view output with colors again by running netq config add color.

All check and show commands are run with a default timeframe of now to one hour ago, unless you specify an approximate time using the around keyword. For example,

NetQ Command Line Overview

CLI Access

running netq check bgp shows the status of BGP over the last hour. Running netq show bgp around 3h shows the status of BGP three hours ago.

Command Prompts

NetQ code examples use the following prompts:

- cumulus@switch:~\$ Indicates the user cumulus is logged in to a switch to run the example command
- cumulus@host:~\$ Indicates the user cumulus is logged in to a host to run the example command
- cumulus@netq-appliance:~\$ Indicates the user cumulus is logged in to either the
 NetQ Appliance or NetQ Cloud Appliance to run the command

The switches must be running the Cumulus Linux operating system (OS), NetQ Platform software, and the NetQ Agent. The hosts must be running CentOS, RHEL, or Ubuntu OS and the NetQ Agent. Refer to the Install NetQ topic for details.

Command Completion

As you enter commands, you can get help with the valid keywords or options using the **Tab** key. For example, using Tab completion with **netq check** displays the possible objects for the command, and returns you to the command prompt to complete the command.

cumulus@switch:~\$ netq check <<pre>cress Tab>>

agents : Netq agent

bgp : BGP info

cl-version: Cumulus Linux version

clag : Cumulus Multi-chassis LAG

evpn : EVPN

interfaces: network interface port

license : License information

Inv : Lightweight Network Virtualization info

mtu : Link MTU

ntp : NTP

ospf : OSPF info

sensors : Temperature/Fan/PSU sensors

vlan : VLAN

vxlan : VXLAN data path

cumulus@switch:~\$ netq check

Command Help

As you enter commands, you can get help with command syntax by entering help at various points within a command entry. For example, to find out what options are available for a BGP check, enter help after entering a portion of the netq check command. In this example, you can see that there are no additional required parameters and two optional parameters, vrf and around, that can be used with a BGP check.

```
cumulus@switch:~$ netq check bgp help
```

Commands:

netq check bgp [vrf <vrf>] [around <text-time>] [json]

cumulus@switch:~\$

To see an exhaustive list of commands, run:

cumulus@switch:~\$ netq help list verbose

Command History

The CLI stores commands issued within a session, which enables you to review and rerun commands that have already been run. At the command prompt, press the **Up Arrow** and **Down Arrow** keys to move back and forth through the list of commands previously entered. When you have found a given command, you can run the command by pressing **Enter**, just as you would if you had entered it manually. Optionally you can modify the command before you run it.

Command Categories

While the CLI has a flat structure, the commands can be conceptually grouped into four functional categories:

- Validation Commands
- Monitoring Commands
- Configuration Commands
- Trace Commands

Validation Commands

The netq check commands enable the network administrator to validate the current or historical state of the network by looking for errors and misconfigurations in the network. The commands run fabric-wide validations against various configured protocols and services to determine how well the network is operating. Validation checks can be performed for the following:

- agents: NetQ Agents operation on all switches and hosts
- bgp: BGP (Border Gateway Protocol) operation across the network fabric
- clag: Cumulus Multi-chassis LAG (link aggregation) operation
- cl-version: Cumulus Linux version
- evpn: EVPN (Ethernet Virtual Private Network) operation

- interfaces: network interface port operation
- license: License status
- Inv: Lightweight Network Virtualization operation
- mtu: Link MTU (maximum transmission unit) consistency across paths
- **ntp**: NTP (Network Time Protocol) operation
- **ospf**: OSPF (Open Shortest Path First) operation
- **sensors**: Temperature/Fan/PSU sensor operation
- vlan: VLAN (Virtual Local Area Network) operation
- vxlan: VXLAN (Virtual Extensible LAN) data path operation

The commands take the form of netq check <network-protocol-or-service>
[options], where the options vary according to the protocol or service.

This example shows the output for the **netq check bgp** command, followed by the same command using the **json** option. If there had been any failures, they would be have been listed below the summary results or in the *failedNodes* section, respectively.

```
cumulus@switch:~$ netq check bgp

Total Nodes: 8, Failed Nodes: 0, Total Sessions: 30, Failed Sessions: 0

cumulus@switch:~$ netq check bgp json

{
    "failedNodes":[
    ],
    "summary":{
        "checkedNodeCount":8,
        "failedSessionCount":0,
        "failedNodeCount":30
```

```
}
}
```

Monitoring Commands

The **netq show** commands enable the network administrator to view details about the current or historical configuration and status of the various protocols or services. The configuration and status can be shown for the following:

• agents: NetQ Agents status on switches and hosts

• **bgp**: BGP status across the network fabric

• clag: CLAG status

• events: Display changes over time

• evpn: EVPN status

• interface-stats: Interface statistics

• interface-utils: Interface statistics plus utilization

• **interfaces**: network interface port status

• **inventory**: hardware component information

• ip: IPv4 status

• ipv6: IPv6 status

 kubernetes: Kubernetes cluster, daemon, pod, node, service and replication status

• IIdp: LLDP status

• Inv: Lightweight Network Virtualization status

• macs: MAC table or address information

• **notification**: Slack or PagerDuty notification configurations

• **ntp**: NTP status

• opta-health: Display health of apps on the OPTA

• ospf: OSPF status

• platform Appliance version info

• sensors: Temperature/Fan/PSU sensor status

• **services**: System services status

• vlan: VLAN status

• vxlan: VXLAN data path status

The commands take the form of netq [<hostname>] show

<network-protocol-or-service> [options], where the options vary according to the protocol or service. The commands can be restricted from showing the information for *all* devices to showing information for a selected device using the hostname option.

This example shows the standard and restricted output for the **netq show** agents command.

	switch:~\$ ne	•	w agents		
			TP Sync Version	Sys U	ptime
			ize Time Last Cha		
exit01	Fresh	yes	2.3.0-cl3u21~15692	46310.30858c3	15h:34m:
15s	15h:34m:5s		15h:34m:5s	Mon Sep 23 22	2:44:49 2019
exit02	Fresh	yes	2.3.0-cl3u21~15692	46310.30858c3	15h:35m:
57s	15h:35m:47	's	15h:35m:47s	Mon Sep 23 2	22:43:09 2019
leaf01	Fresh	yes	2.3.0-cl3u21~15692	46310.30858c3	15h:35m:
10s	15h:35m:1s		15h:35m:1s	Mon Sep 23 22	2:43:55 2019
leaf02	Fresh	yes	2.3.0-cl3u21~15692	46310.30858c3	15h:35m:
53s 15h:35m:43s 15h:35m:43s Mon Sep 23 22:44:17 2019					
leaf03	Fresh	yes	2.3.0-cl3u21~15692	46310.30858c3	15h:35m:
0s	15h:34m:51	5	15h:34m:51s	Mon Sep 23 2	2:44:01 2019
leaf04	Fresh	yes	2.3.0-cl3u21~15692	46310.30858c3	15h:36m:

```
33s
         15h:36m:24s
                          15h:36m:24s
                                           Mon Sep 23 22:43:03 2019
                          2.3.0-ub18.04u21~1569246309.30858c3 15h:14m:
server01
           Fresh
                     no
         15h:14m:34s
                          15h:14m:34s
                                           Mon Sep 23 22:48:56 2019
46s
                    yes 2.3.0-ub18.04u21~1569246309.30858c3 15h:
server02
         Fresh
14m:46s
             15h:14m:34s
                             15h:14m:34s
                                               Mon Sep 23 22:49:24
2019
server03 Fresh
                          2.3.0-ub18.04u21~1569246309.30858c3 15h:
                    yes
14m:46s
             15h:14m:34s
                              15h:14m:34s
                                               Mon Sep 23 22:49:24
2019
server04
         Fresh
                    yes
                          2.3.0-ub18.04u21~1569246309.30858c3 15h:
14m:45s
             15h:14m:33s
                              15h:14m:33s
                                               Mon Sep 23 22:49:24
2019
spine01 Fresh
                          2.3.0-cl3u21~1569246310.30858c3 15h:34m:
                    yes
6s
        15h:33m:57s
                         15h:33m:57s
                                           Mon Sep 23 22:44:27 2019
                    yes 2.3.0-cl3u21~1569246310.30858c3 15h:34m:
spine02
           Fresh
12s
         15h:34m:2s
                        15h:34m:2s
                                          Mon Sep 23 22:43:30 2019
```

```
"agentUptime":1569277757.0
},
  "status":"Fresh",
  "lastChanged":1569278589.0,
  "reinitializeTime":1569277655.0,
  "hostname":"exit02",
  "version": "2.3.0-cl3u21~1569246310.30858c3",
  "sysUptime":1569277645.0,
  "ntpSync":"yes",
  "agentUptime":1569277655.0
},
  "status":"Fresh",
  "lastChanged":1569278635.0,
  "reinitializeTime":1569277701.0,
  "hostname":"leaf01",
  "version": "2.3.0-cl3u21~1569246310.30858c3",
  "sysUptime":1569277692.0,
  "ntpSync":"yes",
  "agentUptime":1569277701.0
},
```

```
cumulus@switch:~$ netq leaf01 show agents

Matching agents records:

Hostname Status NTP Sync Version Sys Uptime

Agent Uptime Reinitialize Time Last Changed
```

leaf01	Fresh	yes	2.3.0-cl3u21~1569	 246310.30858c3	15h:57m:
24s	15h:57m:1	5s	15h:57m:15s	Mon Sep 23 2	2:43:55 2019

Configuration Commands

The **netq config** and **netq notification** commands enable the network administrator to manage NetQ Agent and CLI server configuration, set up container monitoring, and event notification.

NETQ AGENT CONFIGURATION

The agent commands enable the network administrator to configure individual NetQ Agents. Refer to Cumulus NetQ Components for a description of NetQ Agents, to Manage NetQ Agents, or to Install NetQ Agents and CLI on Switches for more detailed usage examples.

The agent configuration commands enable you to add and remove agents from switches and hosts, start and stop agent operations, add and remove Kubernetes container monitoring, add or remove sensors, debug the agent, and add or remove FRR (FRRouting).



Commands apply to one agent at a time, and are run from the switch or host where the NetQ Agent resides.

The agent configuration commands include:

```
netq config (add|del|show) agent
netq config (start|stop|status|restart) agent
```

This example shows how to configure the agent to send sensor data.

```
cumulus@switch~:$ netq config add agent sensors
```

This example shows how to start monitoring with Kubernetes.

cumulus@switch:~\$ netq config add agent kubernetes-monitor poll-period 15

This example shows how to view the NetQ Agent configuration.

```
cumulus@switch:~$ netq config show agent
netq-agent value default
------
enable-opta-discovery True True
exhibitport
agenturl
server 127.0.0.1 127.0.0.1
exhibiturl
vrf default default
agentport 8981 8981
port 31980 31980
```



(i) NOTE

After making configuration changes to your agents, you must restart the agent for the changes to take effect. Use the netq config restart agent command.

CLI CONFIGURATION

The CLI commands enable the network administrator to configure and manage the CLI component. These commands enable you to add or remove CLI (essentially enabling/ disabling the service), start and restart it, and view the configuration of the service.



(i) NOTE

Commands apply to one device at a time, and are run from the switch or host where the CLI is run.

The CLI configuration commands include:

netq config add cli server netq config del cli server netq config show cli premises [json] netq config show (cli|all) [json] netq config (status | restart) cli

This example shows how to restart the CLI instance.

cumulus@switch~:\$ netq config restart cli

This example shows how to enable the CLI on a NetQ Platform or NetQ Appliance.

cumulus@switch~:\$ netq config add cli server 10.1.3.101

This example shows how to enable the CLI on a NetQ Cloud Appliance with a single premise.

netq config add cli server api.netq.cumulusnetworks.com access-key <user-access-key> secret-key <user-secret-key> port 443

EVENT NOTIFICATION COMMANDS

The notification configuration commands enable you to add, remove and show notification application integrations. These commands create the channels, filters, and rules needed to control event messaging. The commands include:

netq (add|del|show) notification channel netq (add|del|show) notification rule netq (add|del|show) notification filter netq (add|del|show) notification proxy

An integration includes at least one channel (PagerDuty, Slack, or syslog), at least one filters (defined by rules you create), and at least one rule.

This example shows how to configure a PagerDuty channel:

cumulus@switch:~\$ netq add notification channel pagerduty pd-netq-events integration-key c6d666e210a8425298ef7abde0d1998

Successfully added/updated channel pd-netq-events

Refer to Integrate NetQ with Notification Applications for details about using these commands and additional examples.

Trace Commands

The trace commands enable the network administrator to view the available paths between two nodes on the network currently and at a time in the past. You can perform a layer 2 or layer 3 trace, and view the output in one of three formats (*json*, *pretty*, and *detail*). JSON output provides the output in a JSON file format for ease of importing to other applications or software. Pretty output lines up the paths in a pseudo-graphical manner to help visualize multiple paths. Detail output is useful for traces with higher hop counts where the pretty output wraps lines, making it harder to interpret the results. The detail output displays a table with a row for each path.

The trace command syntax is:

netq trace <mac> [vlan <1-4096>] from (<src-hostname>|<ip-src>) [vrf <vrf>]
[around <text-time>] [json|detail|pretty] [debug]
netq trace <ip> from (<src-hostname>|<ip-src>) [vrf <vrf>] [around <text-time>]
[json|detail|pretty] [debug]

Example Running a trace based on the destination IP address, in *pretty* output with a small number of resulting paths:

```
cumulus@switch:~$ netq trace 10.0.0.11 from 10.0.0.14 pretty
```

Number of Paths: 6

Inconsistent PMTU among paths

Number of Paths with Errors: 0

Number of Paths with Warnings: 0

Path MTU: 9000

leaf04 swp52 -- swp4 spine02 swp2 -- swp52 leaf02 peerlink.4094 -- peerlink.4094

leaf01 lo

peerlink.4094 -- peerlink.4094 leaf01 lo

leaf04 swp51 -- swp4 spine01 swp2 -- swp51 leaf02 peerlink.4094 -- peerlink.4094

leaf01 lo

peerlink.4094 -- peerlink.4094 leaf01 lo

leaf04 swp52 -- swp4 spine02 swp1 -- swp52 leaf01 lo

leaf04 swp51 -- swp4 spine01 swp1 -- swp51 leaf01 lo

Example Running a trace based on the destination IP address, in *detail* output with a small number of resulting paths:

cumulus@switch:~\$ netq trace 10.0.0.11 from 10.0.0.14 detail

Number of Paths: 6

Inconsistent PMTU among paths

Number of Paths with Errors: 0

Number of Paths with Warnings: 0

Path MTU: 9000

Id Hop Host	name InPort	InVlan InTunnel	InRtrIf	InVRF
OutRtrIf	OutVRF OutTu	nnel OutPort	OutVlan	
1 1 leaf04			swp52	
default	swp52			
2 spine02	swp4	swp4	default	swp2
default	swp2			
3 leaf02	swp52	swp52	default	peerlink.
4094 defaul	t p	eerlink.4094		
4 leaf01	peerlink.4094	peerl	ink.4094	
default		lo		
2 1 leaf04			swp52	
default	swp52			
2 spine02	swp4	swp4	default	swp2
default	swp2			
3 leaf02	swp52	swp52	default	peerlink.
4094 defaul	t p	eerlink.4094		
4 leaf01	peerlink.4094	peerl	ink.4094	
default		lo		
3 1 leaf04			swp51	
default	swp51			
2 spine01	swp4	swp4	default	swp2
default	swp2			
3 leaf02	swp51	swp51	default	peerlink.

4094 default	peerlir	nk.4094		
4 leaf01	peerlink.4094	peerlir	nk.4094	
default		lo		
4 1 leaf04			swp51	
default	swp51			
2 spine01	swp4	swp4	default	swp2
default	swp2			
3 leaf02	swp51	swp51	default	peerlink.
4094 default	peerlir	nk.4094		
4 leaf01	peerlink.4094	peerlir	nk.4094	
default		lo		
5 1 leaf04			swp52	
default	swp52			
2 spine02	swp4	swp4	default	swp1
default	swp1			
3 leaf01	swp52	swp52		
default		lo		
6 1 leaf04			swp51	
default	swp51			
2 spine01	swp4	swp4	default	swp1
default	swp1			
3 leaf01	swp51	swp51		
default		lo		



Example Running a trace based on the destination MAC address, in *pretty* output:

```
cumulus@switch:~$ netq trace A0:00:00:00:00:11 vlan 1001 from Server03 pretty
Number of Paths: 6
Number of Paths with Errors: 0
Number of Paths with Warnings: 0
Path MTU: 9152
Server03 bond1.1001 -- swp7 <vlan1001> Leaf02 vni: 34 swp5 -- swp4 Spine03
swp7 -- swp5 vni: 34 Leaf04 swp6 -- swp1.1001 Server03 <swp1.1001>
                            swp4 -- swp4 Spine02 swp7 -- swp4 vni: 34 Leaf04
swp6 -- swp1.1001 Server03 <swp1.1001>
                            swp3 -- swp4 Spine01 swp7 -- swp3 vni: 34 Leaf04
swp6 -- swp1.1001 Server03 <swp1.1001>
     bond1.1001 -- swp7 <vlan1001> Leaf01 vni: 34 swp5 -- swp3 Spine03 swp7 --
swp5 vni: 34 Leaf04 swp6 -- swp1.1001 Server03 <swp1.1001>
                            swp4 -- swp3 Spine02 swp7 -- swp4 vni: 34 Leaf04
swp6 -- swp1.1001 Server03 <swp1.1001>
                            swp3 -- swp3 Spine01 swp7 -- swp3 vni: 34 Leaf04
swp6 -- swp1.1001 Server03 <swp1.1001>
```

Command Changes

A number of commands have changed in this release to accommodate the addition of new keywords and options or to simplify their syntax. Additionally, new commands have been added and others have been removed. A summary of those changes is provided here.

New Commands

The following table summarizes the new commands available with this release.

Command	Summary	Version
netq check cl-version [match-version <cl-ver> min-version <cl-ver>] [include <version-number-range-list> exclude <version-number-range-list>] [around <text-time>] [json summary]</text-time></version-number-range-list></version-number-range-list></cl-ver></cl-ver>	When no options are used, validates that the Cumulus Linux version running on all of the monitored nodes is consistent. When min-version is provided, validates CL version is equal or greater than specified version. When match-version is provided, validates all nodes are using specified version.	2.3.0
netq [<hostname>] show interface-utils [<text-port>] [tx rx] [around <text-time>] [json]</text-time></text-port></hostname>	Similar to netq show interface- stats, this command displays NetQ server or appliance interface receive and transmit statistics, but it also shows utilization and port speed.	2.3.0
netq show platform [json]	Displays the NetQ software version installed on your NetQ server or appliance.	2.3.0

Modified Commands

The following table summarizes the commands that have been changed with this release.

Updated Command	Old Command	What Changed	Version
netq config add cli server <text- gateway-dest=""> [access-key <text- access-key=""> secret- key <text-secret- key=""> premises</text-secret-></text-></text->	netq config add cli server <text- gateway-dest=""> [access-key <text- access-key=""> secret-key <text- secret-key=""> cli- keys-file <text- key-file="">] [premise <text- premise-name="">] [port <text- gateway-port="">] [vrf <text-vrf- name="">]</text-vrf-></text-></text-></text-></text-></text-></text->	Adds the CLI daemon to the switch or host where this command is run. Split command into two: the first for when you want to specify the premises, the second for when you want to use the first premises in your list of premises. The access-key, secret-key, and port are required for cloud deployments. The text-gateway-dest is the only required value for on-premises deployments.	2.3.0

Deprecated Commands

The following table summarizes the commands that have been removed and a recommended alternative, if appropriate.

Command	Alternative Command
N/A	N/A

Monitor Overall Network Health

NetQ provides the information you need to monitor the health of your network fabric, devices, and interfaces. You are able to easily validate the operation and view the configuration across the entire network from switches to hosts to containers. For example, you can monitor the operation of routing protocols and virtual network configurations, the status of NetQ Agents and hardware components, and the operation and efficiency of interfaces. When issues are present, NetQ makes it easy to identify and resolve them. You can also see when changes have occurred to the network, devices, and interfaces by viewing their operation, configuration, and status at an earlier point in time.

Validate Network Health

NetQ check commands validate the various elements of your network fabric, looking for inconsistencies in configuration across your fabric, connectivity faults, missing configuration, and so forth, and then and display the results for your assessment. They can be run from any node in the network.

Validate the Network Fabric

You can validate the following network fabric elements:

cumulus@switch:/\$ netq check

agents : Netq agent

bap : BGP info

cl-version: Cumulus Linux version

clag : Cumulus Multi-chassis LAG

evpn: EVPN

interfaces: network interface port

license : License information

Inv : Lightweight Network Virtualization info

mtu : Link MTU

ntp : NTP

ospf : OSPF info

sensors : Temperature/Fan/PSU sensors

vlan : VLAN

vxlan : VXLAN data path

For example, to determine the status of BGP running on your network:

cumulus@switch:~\$ netq check bgp

Total Nodes: 15, Failed Nodes: 0, Total Sessions: 16, Failed Sessions: 0

You can see from this output that NetQ has validated the connectivity and configuration of BGP across all of the nodes in the network and found them all to be operating properly. If there were issues with any of the nodes, NetQ would provide information about each node to aid in resolving the issues.

There is a check command for each of the supported fabric elements. They all behave in a similar manner, checking for connectivity, configuration, and other problems, indicating the number of nodes that they have checked and indicating the number that have failed.

Some additional examples—

Validate that EVPN is running correctly on all nodes:

cumulus@switch:~\$ netq check evpn

Total Nodes: 15, Failed Nodes: 0, Total Sessions: 0, Failed Sessions: 0, Total VNIs: 0

Confirm all monitored nodes are running the NetQ Agent:

cumulus@switch:~\$ netq check agents

Checked nodes: 25, Rotten nodes: 1

Hostname Status Last Changed

leaf01 Rotten 8d:13h:34m:51s

Validate that all corresponding interface links have matching MTUs. The first shows no mismatches, the second shows an error.

cumulus@switch:~\$ netq check mtu

Checked Nodes: 15, Checked Links: 138, Failed Nodes: 0, Failed Links: 0

No MTU Mismatch found

cumulus@switch:~\$ netq check mtu

Checked Nodes: 13, Checked Links: 173, Failed Nodes: 1, Failed Links: 1

MTU mismatch found on following links

Hostname Interface MTU Peer Peer Interface Peer

MTU Error

leaf01 - - - Rotten Agent

Validate that VXLANs are configured and operating properly:

cumulus@switch:~\$ netq check vxlan Checked Nodes: 6, Warning Nodes: 0, Failed Nodes: 6 Nodes with error Hostname Reason inconsistent replication list for vni 104001 exit01 exit02 inconsistent replication list for vni 104001 leaf01 inconsistent replication list for vni 104001 leaf02 inconsistent replication list for vni 104001 leaf03 inconsistent replication list for vni 104001 leaf04 inconsistent replication list for vni 104001



TIP

Both asymmetric and symmetric VXLAN configurations are validated with this command.

You can be more granular in your validation as well, using the additional options available for each of the check commands. For example, validate BGP operation for nodes communicating over a particular VRF:

cumulus@switch:~\$ netq check bgp vrf DataVrf1081

Total Nodes: 25, Failed Nodes: 1, Total Sessions: 52, Failed Sessions: 1

Hostname VRF Peer Name Peer Hostname

Reason Last Changed

exit-1 DataVrf1081 swp7.3 firewall-2 BGP session with peer

firewall-2 (swp7.3 vrf 1d:5h:47m:31s

DataVrf1081) failed,

reason: Peer not configured

Each of the check commands provides a starting point for troubleshooting configuration and connectivity issues within your network in real time. They provide an additional option of viewing the network state at an earlier time, using the around option.

For example, if you were notified of an issue on your VLANs that appears to have occurred about 10 minutes ago, you could run:

cumulus@switch:~\$ netq check vlan around 10m

Checked Nodes: 15, Checked Links: 138, Failed Nodes: 0, Failed Links: 0

No VLAN or PVID Mismatch found

Validate Device Status and Configuration

You can validate the following device elements:

NTP

- Sensors
- License

It is always important to have your devices in time synchronization to ensure configuration and management events can be tracked and correlations can be made between events. To validate time synchronization, run:

cumulus@	cumulus@switch:~\$ netq check ntp							
Total Node	Total Nodes: 15, Checked Nodes: 15, Rotten Nodes: 0, Unknown Nodes: 0, failed							
NTP Nodes	5: 8							
Hostname	Ν	TP Sync Connect Time						
exit01	no	2018-09-12 16:30:39						
exit02	no	2018-09-12 16:30:45						
leaf01	no	2018-09-12 16:30:43						
leaf02	no	2018-09-12 16:30:36						
leaf03	no	2018-09-12 16:30:36						
leaf04	no	2018-09-12 16:30:34						
spine01	no	2018-09-12 16:30:44						
spine02	no	2018-09-12 16:30:40						

This example shows eight nodes that are not in time synchronization. You can now continue to investigate these nodes, validating that the NetQ Agents are active, whether an NTP server has become unreachable, and so forth.

Hardware platforms have a number sensors to provide environmental data about the switches. Knowing these are all within range is a good check point for maintenance. For example, if you had a temporary HVAC failure and you are concerned that some of your nodes are beginning to overheat, you can run:

cumulus@switch:~\$ netq check sensors

Total Nodes: 25, Failed Nodes: 0, Checked Sensors: 221, Failed Sensors: 0

You can also check for any nodes that have invalid licenses without going to each node. Because switches do not operate correctly without a valid license you might want to verify that your Cumulus Linux licenses on a regular basis:

cumulus@switch:~\$ netq check license

Total Nodes: 15, Failed Nodes: 0, Checked Licenses: 10, Failed Licenses: 0



This command checks every node, meaning every switch and host in the network. Hosts do not require a Cumulus Linux license, so the number of licenses checked might be smaller than the total number of nodes checked.

Validate Interface Status and Configuration

As with other netq check commands, you can validate the proper operation of your interfaces across the network:

cumulus@switch:~\$ netq check interfaces

Checked Nodes: 15, Failed Nodes: 8

Checked Ports: 118, Failed Ports: 8, Unverified Ports: 94

Interface Hostname Peer Hostname Peer Interface Message

leaf01	swp7	firewall02	swp3	Speed mismatch
(10G, n/a),				
			Autoneg r	nismatch (off, n/a)
leaf02	swp2	server02	eth2	Autoneg mismatch
(off, on)				
leaf03	swp1	server03	eth1	Autoneg mismatch
(off, on)				
leaf04	swp2	server04	eth2	Autoneg mismatch
(off, on)				
server01	eth1	leaf01	swp1	Autoneg mismatch
(on, off)				
server02	eth2	leaf02	swp2	Autoneg mismatch
(on, off)			·	
server03	eth1	leaf03	swp1	Autoneg mismatch
(on, off)			•	Ü
server04	eth2	leaf04	swp2	Autoneg mismatch
(on, off)			- 1	J
(011, 011)				

When failures are seen, additional information is provided to start your investigation. In this example, some reconfiguration is required for auto-negotiation with peer interfaces.

Create Filters for Provisioning Exceptions

With this release, you are able to configure filters to change validation errors to warnings that would normally occur due to the default expectations of the **netq check** commands. This applies to all protocols and services, except for Agents and LNV. For example, if you have provisioned BGP with configurations where a BGP peer is not

expected or desired, you will get errors that a BGP peer is missing. By creating a filter, you can remove the error in favor of a warning.

To create a validation filter:

- 1. Navigate to the /etc/netq directory.
- 2. Create or open the *check_filter.yml* file using your text editor of choice.

This file contains the syntax to follow to create one or more rules for one or more protocols or services. Create your own rules, and/or edit and un-comment any example rules you would like to use.

```
# Netq check result filter rule definition file. This is for filtering
# results based on regex match on one or more columns of each test result.
# Currently, only action 'filter' is supported. Each test can have one or
# more rules, and each rule can match on one or more columns. In addition,
# rules can also be optionally defined under the 'global' section and will
# apply to all tests of a check.
# syntax:
#
# <check name>:
# tests:
    <test name, as shown in test list when using the include/exclude and tab>:
     - rule:
#
       match:
#
        <column name>: regex
        <more columns and regex.., result is AND>
       action:
#
#
        filter
```

```
# - <more rules..>
# global:
# - rule:
  . . .
# - rule:
#
  . . .
# <another check name>:
# ...
#
# e.g.
#
# bgp:
# tests:
# Address Families:
# - rule:
    match:
#
  Hostname: (^exit*|^firewall)
#
  VRF: DataVrf1080
#
  Reason: AFI/SAFI evpn not activated on peer
   action:
     filter
#
   - rule:
#
    match:
#
   Hostname: exit-2
#
#
       Reason: SAFI evpn not activated on peer
#
      action:
      filter
# Router ID:
```

```
- rule:
  match:
#
      Hostname: exit-2
      action:
       filter
#
# evpn:
# tests:
  EVPN Type 2:
   - rule:
#
#
      match:
#
      Hostname: exit-1
#
      action:
        filter
#
#
```

View Network Details

The **netq show** commands display a wide variety of content about the network and its various elements. You can show content for the following:

```
cumulus@switch:~$ netq show

agents : Netq agent

bgp : BGP info

clag : Cumulus Multi-chassis LAG

events : Display changes over time

evpn : EVPN

interfaces : network interface port
```

inventory : Inventory information

ip : IPv4 related infoipv6 : IPv6 related info

kubernetes: Kubernetes Information

Ildp : LLDP based neighbor info

Inv : Lightweight Network Virtualization info

macs : Mac table or MAC address info

notification: Send notifications to Slack or PagerDuty

ntp : NTP

ospf : OSPF info

sensors : Temperature/Fan/PSU sensors

services : System services

vlan : VLAN

vxlan : VXLAN data path

For example, to validate the the status of the NetQ agents running in the fabric, run **netq show agents**. A *Fresh* status indicates the Agent is running as expected. The Agent sends a heartbeat every 30 seconds, and if three consecutive heartbeats are missed, its status changes to *Rotten*.

cumulus@switch:~\$ netq show agents

Matching agents records:

Hostname Status NTP Sync Version Sys Uptime

Agent Uptime Reinitialize Time Last Changed

exit01 Fresh yes 2.3.0-cl3u21~1569246310.30858c3 1d:21h:30m:

19s 1d:21h:30m:9s 1d:21h:30m:9s Tue Sep 24 22:45:03 2019

exit02	Fresh	yes	2.3.0-cl3u21~156924	46310.30858c3	1d:21h:32m:
1s	1d:21h:31m:5	51s	1d:21h:31m:51s	Tue Sep 24 22	2:43:35 2019
leaf01	Fresh	yes	2.3.0-cl3u21~15692	46310.30858c3	1d:21h:31m:
14s	1d:21h:31m:	5s	1d:21h:31m:5s	Tue Sep 24 22	:44:25 2019
leaf02	Fresh	yes	2.3.0-cl3u21~15692	46310.30858c3	1d:21h:31m:
57s	1d:21h:31m:	47s	1d:21h:31m:47s	Tue Sep 24 2	2:44:20 2019
leaf03	Fresh	yes	2.3.0-cl3u21~15692	46310.30858c3	1d:21h:31m:
4s	1d:21h:30m:5	55s	1d:21h:30m:55s	Tue Sep 24 22	2:44:19 2019
leaf04	Fresh	yes	2.3.0-cl3u21~15692	46310.30858c3	1d:21h:32m:
37s	1d:21h:32m:	28s	1d:21h:32m:28s	Tue Sep 24 2	2:43:05 2019
server01	Fresh	no	2.3.0-ub18.04u21~	1569246309.308	58c3 1d:21h:
10m:50s	1d:21h:	10m:38s	1d:21h:10m:3	8s Tue Sep	24 22:49:10
2019					
server02	2 Fresh	yes	2.3.0-ub18.04u21~	1569246309.308	58c3 1d:21h:
10m:50s	1d:21h:	10m:38s	1d:21h:10m:3	8s Wed Se	p 25 18:35:44
2019					
server03	3 Fresh	yes	2.3.0-ub18.04u21~	1569246309.308	58c3 1d:21h:
10m:50s	1d:21h:	10m:38s	1d:21h:10m:3	8s Wed Se	p 25 15:37:10
2019					
server04	l Fresh	yes	2.3.0-ub18.04u21~	1569246309.308	58c3 1d:21h:
10m:49s	1d:21h:	10m:37s	1d:21h:10m:3	7s Tue Sep	24 22:49:33
2019					
spine01	Fresh	yes	2.3.0-cl3u21~15692	246310.30858c3	1d:21h:30m:
10s	1d:21h:30m:	1s	1d:21h:30m:1s	Tue Sep 24 22	:44:40 2019
spine02	Fresh	yes	2.3.0-cl3u21~15692	246310.30858c3	1d:21h:30m:
16s	1d:21h:30m:	6s	1d:21h:30m:6s	Tue Sep 24 22	:43:32 2019

Some additional examples follow.

View the status of BGP:

	switch:~\$ netq show bgr		A CAL	D 461	I D(D	
Hostname	e Neighbor	VRF	ASN	Peer ASN	l PfxRx	Last
Changed						
					-	
exit01		vrf1	65041	25253	2/-/- F	ri Apr
19 16:00:4	0 2019					
exit01	swp51(spine01)	default	65041	65020	8/-/59	Fri
Apr 19 16:	00:40 2019					
exit01	swp52(spine02)	default	65041	65020	8/-/59	Fri
Apr 19 16:	00:40 2019					
exit02	swp44(internet)	vrf1	65042	25253	7/-/- F	ri Apr
19 16:00:4	0 2019					
exit02	swp51(spine01)	default	65042	65020	8/-/59	Fri
Apr 19 16:	00:40 2019					
exit02	swp52(spine02)	default	65042	65020	8/-/59	Fri
Apr 19 16:	00:40 2019					
leaf01	peerlink.4094(leaf02)	default	65011	1 65011	9/-/34	Fri
Apr 19 16:	00:40 2019					
leaf01	swp51(spine01)	default	65011	65020	6/-/34	Fri
Apr 19 16:	00:40 2019					
leaf01	swp52(spine02)	default	65011	65020	6/-/34	Fri
Apr 19 16:	00:40 2019					
leaf02	peerlink.4094(leaf01)	default	65011	1 65011	9/-/34	Fri
Apr 19 16:	00:40 2019					
leaf02	swp51(spine01)	default	65011	65020	6/-/34	Fri

leaf02	swp52(spine02)	default	65011	65020	6/-/34	Fri
Apr 19 16:0	00:40 2019					
leaf03	peerlink.4094(leaf04)	default	65012	65012	9/-/34	Fri
Apr 19 16:0	00:40 2019					
leaf03	swp51(spine01)	default	65012	65020	6/-/34	Fri
Apr 19 16:0	00:40 2019					
leaf03	swp52(spine02)	default	65012	65020	6/-/34	Fri
Apr 19 16:0	00:40 2019					
leaf04	peerlink.4094(leaf03)	default	65012	65012	9/-/34	Fri
Apr 19 16:0	00:40 2019					
leaf04	swp51(spine01)	default	65012	65020	6/-/34	Fri
Apr 19 16:0	00:40 2019					
leaf04	swp52(spine02)	default	65012	65020	6/-/34	Fri
Apr 19 16:0	00:40 2019					
spine01	swp1(leaf01)	default	65020	65011	3/-/14	Fri
Apr 19 16:0	00:40 2019					
spine01	swp2(leaf02)	default	65020	65011	3/-/14	Fri
Apr 19 16:0	00:40 2019					
spine01	swp29(exit02)	default	65020	65042	1/-/3	Fri
Apr 19 16:0	00:40 2019					
spine01	swp3(leaf03)	default	65020	65012	3/-/14	Fri
Apr 19 16:0	00:40 2019					
spine01	swp30(exit01)	default	65020	65041	1/-/3	Fri
Apr 19 16:0	00:40 2019					
spine01	swp4(leaf04)	default	65020	65012	3/-/14	Fri
Apr 19 16:0	00:40 2019					
	swp1(leaf01)	default	65020	65011	3/-/12	Fri
Apr 19 16:0	00:40 2019					
spine02	swp2(leaf02)	default	65020	65011	3/-/12	Fri

Apr 19 16:0	0:40 2019					
spine02	swp29(exit02)	default	65020	65042	1/-/3	Fri
Apr 19 16:0	0:40 2019					
spine02	swp3(leaf03)	default	65020	65012	3/-/12	Fri
Apr 19 16:0	0:40 2019					
spine02	swp30(exit01)	default	65020	65041	1/-/3	Fri
Apr 19 16:0	0:40 2019					
spine02	swp4(leaf04)	default	65020	65012	3/-/12	Fri
Apr 19 16:0	0:40 2019					

View the status of your VLANs:

Matching vl	an records:		
Hostname	VLANs	SVIs	Last Changed
server11	1	Thu Fel	o 7 00:17:48 2019
server21	1	Thu Fel	o 7 00:17:48 2019
server11	1	Thu Fel	o 7 00:17:48 2019
server13	1	Thu Fel	7 00:17:48 2019
server21	1	Thu Fel	o 7 00:17:48 2019
server23	1	Thu Fel	7 00:17:48 2019
leaf01	100-106,1000-1009	100-106 1000	-1009 Thu Feb 7 00:17:49
2019			
leaf02	100-106,1000-1009	100-106 1000-	-1009 Thu Feb 7 00:17:49
2019			
leaf11	100-106,1000-1009	100-106 1000	-1009 Thu Feb 7 00:17:49
2019			

leaf12	100-106,1000-1009	100-106 1000-1009	Thu Feb 7 00:17:50
2019			
leaf21	100-106,1000-1009	100-106 1000-1009	Thu Feb 7 00:17:50
2019			
leaf22	100-106,1000-1009	100-106 1000-1009	Thu Feb 7 00:17:50
2019			

View the status of the hardware sensors:

cumulus@	စ္တswitch:~\$ n	etq show sensors all		
Matching	sensors reco	ords:		
Hostname	e Name	Description	State	
Message		Last Changed		
exit01	fan1	fan tray 1, fan 1	ok	Wed
Feb 6 23:	02:35 2019			
exit01	fan2	fan tray 1, fan 2	ok	Wed
Feb 6 23:	02:35 2019			
exit01	fan3	fan tray 2, fan 1	ok	Wed
Feb 6 23:	02:35 2019			
exit01	fan4	fan tray 2, fan 2	ok	Wed
Feb 6 23:	02:35 2019			
exit01	fan5	fan tray 3, fan 1	ok	Wed
Feb 6 23:	02:35 2019			
exit01	fan6	fan tray 3, fan 2	ok	Wed
Feb 6 23:	02:35 2019			

exit01	psu1fan1	psu1 fan	ok	Wed
Feb 6 23:0	02:35 2019			
exit01	psu2fan1	psu2 fan	ok	Wed
Feb 6 23:0	02:35 2019			
exit02	fan1	fan tray 1, fan 1	ok	Wed
Feb 6 23:0	03:35 2019			
exit02	fan2	fan tray 1, fan 2	ok	Wed
Feb 6 23:0	03:35 2019			
exit02	fan3	fan tray 2, fan 1	ok	Wed
Feb 6 23:0	03:35 2019			
exit02	fan4	fan tray 2, fan 2	ok	Wed
Feb 6 23:0	03:35 2019			
exit02	fan5	fan tray 3, fan 1	ok	Wed
Feb 6 23:0	03:35 2019			
exit02	fan6	fan tray 3, fan 2	ok	Wed
Feb 6 23:0	03:35 2019			
exit02	psu1fan1	psu1 fan	ok	Wed
Feb 6 23:0	03:35 2019			
exit02	psu2fan1	psu2 fan	ok	Wed
Feb 6 23:0	03:35 2019			
leaf01	fan1	fan tray 1, fan 1	ok	Wed
Feb 6 23:0	01:12 2019			
leaf01	fan2	fan tray 1, fan 2	ok	Wed
Feb 6 23:0	01:12 2019			
leaf01	fan3	fan tray 2, fan 1	ok	Wed
Feb 6 23:0	01:12 2019			
leaf01	fan4	fan tray 2, fan 2	ok	Wed
Feb 6 23:0	01:12 2019			
leaf01	fan5	fan tray 3, fan 1	ok	Wed

Feb 6 23:0	1:12 2019			
leaf01	fan6	fan tray 3, fan 2	ok	Wed
Feb 6 23:0	1:12 2019			
leaf01	psu1fan1	psu1 fan	ok	Wed
Feb 6 23:0	1:12 2019			
leaf01	psu2fan1	psu2 fan	ok	Wed
Feb 6 23:0	1:12 2019			
leaf02	fan1	fan tray 1, fan 1	ok	Wed
Feb 6 22:59	9:54 2019			
leaf02	fan2	fan tray 1, fan 2	ok	Wed
Feb 6 22:59	9:54 2019			
leaf02	fan3	fan tray 2, fan 1	ok	Wed
Feb 6 22:59	9:54 2019			
leaf02	fan4	fan tray 2, fan 2	ok	Wed
Feb 6 22:59	9:54 2019			
leaf02		fan tray 3, fan 1	ok	Wed
Feb 6 22:59	9:54 2019			
•••				

Monitor Switch Hardware and Software

With NetQ, a network administrator can monitor both the switch hardware and software components for misconfigurations. NetQ helps answer questions such as:

- What switches do I have in the network?
- What hardware and software are installed on my switches?
- · How many transmit and receive packets have been dropped?
- Are all switches licensed correctly?
- Do all switches have NetQ agents running?

NetQ uses LLDP (Link Layer Discovery Protocol) to collect port information. NetQ can also identify peer ports connected to DACs (Direct Attached Cables) and AOCs (Active Optical Cables) without using LLDP, even if the link is not UP.

The NetQ CLI provides the netq show inventory, netq show sensors, and netq show events commands to monitor switches.

Monitor Switch and Host Hardware Information

You can view summary information about all switches and hosts along with their key components, including the motherboard, ASIC, microprocessor, disk and memory information.

To view the switch and host information with the CLI, use the following **netq show** commands:

```
netq [<hostname>] show inventory brief [opta] [json]
netq [<hostname>] show inventory asic [vendor <asic-vendor>|model <asic-model>|model-id <asic-model-id>] [opta] [json]
netq [<hostname>] show inventory board [vendor <board-vendor>|model <board-model>] [opta] [json]
netq [<hostname>] show inventory cpu [arch <cpu-arch>] [opta] [json]
netq [<hostname>] show inventory disk [name <disk-name>|transport <disk-transport>|vendor <disk-vendor>] [opta] [json]
netq [<hostname>] show inventory license [cumulus] [status ok | status missing]
[around <text-time>] [opta] [json]
netq [<hostname>] show inventory memory [type <memory-type>|vendor <memory-vendor>] [opta] [json]
netq [<hostname>] show inventory os [version <os-version>|name <os-name>]
[opta] [json]
```

netq [<hostname>] show sensors all [around <text-time>] [json]
netq [<hostname>] show sensors psu [<psu-name>] [around <text-time>] [json]
netq [<hostname>] show sensors temp [<temp-name>] [around <text-time>] [json]
netq [<hostname>] show sensors fan [<fan-name>] [around <text-time>] [json]

netq [<hostname>] show events [level info | level error | level warning | level critical | level debug] [type sensors] [between <text-time> and <text-endtime>] [json]



(i) NOTE

When entering a time value, you must include a numeric value and the unit of measure:

- w: week(s)
- d: day(s)
- h: hour(s)
- m: minute(s)
- s: second(s)
- now

For time ranges, the <text-time> is the most recent time and the <textendtime> is the oldest time. The values do not have to have the same unit of measure.



(i) NOTE

The keyword values for the vendor, model, model-id, arch, name, transport, type, version, psu, temp, and fan keywords are specific to your deployment. For example, if you have devices with CPU architectures of only one type, say Intel x86, then that is the only option available for the cpu-arch keyword value. If you have multiple CPU architectures, say you also have ARMv7, then that would also be an option for you.

To view the switch and host information with the GUI, use the Devices Inventory card workflow which contains a small card with a count of each device type in your network, a medium card displaying the operating systems running on each set of devices, large

cards with component information statistics, and full-screen cards displaying tables with attributes of all switches and all hosts in your network.

View a Summary of Your Network Inventory

While the detail can be very helpful, sometimes a simple overview of the hardware inventory is better. This example shows the basic hardware information for all devices.

cumulus@switch:~\$ netq show inventory brief								
Matching in	Matching inventory records:							
Hostname	Switch	OS		CP	U	ASIC		Ports
edge01	N/A	Ubuntu	· 	x8(5_64	N/A		N/A
exit01	VX	CL	x86	_64	VX		N/A	
exit02	VX	CL	x86	_64	VX		N/A	
leaf01	VX	CL	x86	_64	VX		N/A	
leaf02	VX	CL	x86	_64	VX		N/A	
leaf03	VX	CL	x86	_64	VX		N/A	
leaf04	VX	CL	x86	_64	VX		N/A	
server01	N/A	Ubuntu	I	x8	6_64	N/A		N/A
server02	N/A	Ubuntu	I	x8	6_64	N/A		N/A
server03	N/A	Ubuntu	l	x8	6_64	N/A		N/A
server04	N/A	Ubuntu	I	x8	6_64	N/A		N/A
spine01	VX	CL	x8	6_64	l VX		N/A	
spine02	VX	CL	x8	6_64	l VX		N/A	

View Information about the ASIC on all Switches

You can view the vendor, model, model identifier, core bandwidth capability, and ports of the ASIC installed on your switch motherboard. This example shows all of these for all devices.

	ntory records:	N4 1 1	M 1 1 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7	6
Hostname	Vendor	Model	Model ID	Core
BW Ports				
dell_79100_05		Tomahawk	BCM56960	
2.0T 32 x		TOTTIATIAWK	DCIVISOSOO	
	Mellanox	Spectrum	MT52132	N/
A 16 x 10		5 p 5 - 11 - 11 - 11		
	Mellanox	Spectrum	MT52132	N/
A 48 x 25	5G-SFP28 & 8 x	100G-QSFP28		
mlx-2700-11	Mellanox	Spectrum	MT52132	N/
A 32 x 10	00G-QSFP28			
qct-ix1-08	Broadcom	Tomahawk	BCM56960	
2.0T 32 x	100G-QSFP28			
qct-ix7-04	Broadcom	Trident3	BCM56870	N/
A 32 x 10	00G-QSFP28			
qct-ix7-04	N/A N	/A	N/A N/A	N/A
st1-l1 Br	oadcom	Trident2	BCM56854	720G
48 x 10G-SFP+	& 6 x 40G-QSF	P+		
st1-l2 Br	oadcom	Trident2	BCM56854	720G

st1-l3	Broadcom	Trident2	BCM56854	720G
48 x 10G-	SFP+ & 6 x 40G-0	SFP+		
st1-s1	Broadcom	Trident2	BCM56850	960G
32 x 40G-	QSFP+			
st1-s2	Broadcom	Trident2	BCM56850	960G
32 x 40G-	QSFP+			

You can filter the results of the command to view devices with a particular characteristic. This example shows all devices that use a Broadcom ASIC.

·	ow inventory asic ve	endor Broddeon	
Matching inventory records			
Hostname Vendor	Model	Model ID	Core
BW Ports			
dell-z9100-05 Broadcom	Tomahawk	BCM56960	
2.0T 32 x 100G-QSFP28	3		
qct-ix1-08 Broadcom	Tomahawk	BCM56960	
2.0T 32 x 100G-QSFP28	3		
qct-ix7-04 Broadcom	Trident3	BCM56870	N/
A 32 x 100G-QSFP28			
st1-l1 Broadcom	Trident2	BCM56854	720G
48 x 10G-SFP+ & 6 x 40G-QS	FP+		
st1-l2 Broadcom	Trident2	BCM56854	720G
48 x 10G-SFP+ & 6 x 40G-QS	FP+		
st1-l3 Broadcom	Trident2	BCM56854	720G
48 x 10G-SFP+ & 6 x 40G-QS	FP+		

st1-s1	Broadcom	Trident2	BCM56850	960G
32 x 40G-	QSFP+			
st1-s2	Broadcom	Trident2	BCM56850	960G
32 x 40G-	QSFP+			

You can filter the results of the command view the ASIC information for a particular switch. This example shows the ASIC information for *st1-11* switch.

	vitch:~\$ netq le	eaf02 show inven	tory asic	
Hostname	Vendor	Model	Model ID	Core BW
Ports				
st1-l1 E	Broadcom	Trident2	BCM56854	720G
48 x 10G-SFP	°+ & 6 x 40G-Q	SFP+		

View Information about the Motherboard in a Switch

You can view the vendor, model, base MAC address, serial number, part number, revision, and manufacturing date for a switch motherboard on a single device or on all devices. This example shows all of the motherboard data for all devices.

cumulus@switch:~\$ netq show inventory board

Matching inventory records:

Hostname Vendor Model Base MAC Serial No

Part No Rev Mfg Date

dell-z9100-05 DELL Z9100-ON 4C:76:25:E7:42:C0 CN03GT5N779315C20001 03GT5N A00 12/04/2015 mlx-2100-05 Penguin Arctica 1600cs 7C:FE:90:F5:61:C0 MT1623X10078 MSN2100-CB2FO N/A 06/09/2016 mlx-2410a1-05 Mellanox SN2410 EC:0D:9A:4E:55:C0 MT1734X00067 MSN2410-CB2F QP3 N/A 08/24/2017 mlx-2700-11 Penguin Arctica 3200cs 44:38:39:00:AB:80 MT1604X21036 MSN2700-CS2FO N/A 01/31/2016 qct-ix1-08 QCT QuantaMesh BMS T7032-IX1 54:AB:3A:78:69:51 QTFCO7623002C 1IX1UZZ0ST6 H3B 05/30/2016 qct-ix7-04 QCT IX7 D8:C4:97:62:37:65 QTFCUW821000A 1IX7UZZ0ST5 B3D 05/07/2018 gct-ix7-04 QCT T7032-IX7 D8:C4:97:62:37:65 QTFCUW821000A 1IX7UZZ0ST5 B3D 05/07/2018 st1-l1 CELESTICA Arctica 4806xp 00:E0:EC:27:71:37 D2060B2F044919GD000011 R0854-F1004-01 Redsto 09/20/2014 ne-XP st1-l2 CELESTICA Arctica 4806xp 00:E0:EC:27:6B:3A D2060B2F044919GD000060 R0854-F1004-01 Redsto 09/20/2014 ne-XP st1-l3 Penguin Arctica 4806xp 44:38:39:00:70:49 N/A N/A N/A N/A st1-s1 Dell S6000-ON 44:38:39:00:80:00 N/A N/A N/A N/A st1-s2 Dell S6000-ON 44:38:39:00:80:81 N/A N/A N/A N/A

You can filter the results of the command to capture only those devices with a particular motherboard vendor. This example shows only the devices with *Celestica* motherboards.

You can filter the results of the command to view the model for a particular switch. This example shows the motherboard vendor for the *st1-s1* switch.

View Information about the CPU on a Switch

You can view the architecture, model, operating frequency, and the number of cores for the CPU on a single device or for all devices. This example shows these CPU characteristics for all devices.

```
cumulus@nswitch:~$ netg show inventory cpu
Matching inventory records:
Hostname
             Arch Model
                                     Freq
                                             Cores
dell-z9100-05 x86_64 Intel(R) Atom(TM) C2538 2.40GHz 4
mlx-2100-05 x86 64 Intel(R) Atom(TM) C2558 2.40GHz 4
mlx-2410a1-05 x86 64 Intel(R) Celeron(R) 1047UE 1.40GHz 2
             x86 64 Intel(R) Celeron(R) 1047UE 1.40GHz 2
mlx-2700-11
qct-ix1-08
            x86 64 Intel(R) Atom(TM) C2558
                                            2.40GHz 4
qct-ix7-04 x86_64 Intel(R) Atom(TM) C2558
                                            2.40GHz 4
st1-l1
          x86 64 Intel(R) Atom(TM) C2538
                                          2.41GHz 4
st1-l2
          x86_64 Intel(R) Atom(TM) C2538
                                          2.41GHz 4
st1-l3
          x86_64 Intel(R) Atom(TM) C2538
                                          2.40GHz 4
          x86 64 Intel(R) Atom(TM) S1220
                                          1.60GHz 4
st1-s1
          x86_64 Intel(R) Atom(TM) S1220
                                          1.60GHz 4
st1-s2
```

You can filter the results of the command to view which switches employ a particular CPU architecture using the *arch* keyword. This example shows how to determine which architectures are deployed in your network, and then shows all devices with an *x86_64* architecture.

	vitch:~\$ netq show inventory cpu arch		
cumulus@sv	vitch:~\$ netq show inventory cpu arch x86_64		
	ventory records:		
Hostname	Arch Model Freq Cores		
leaf01	x86_64 Intel Core i7 9xx (Nehalem Cla N/A ss Core i7)	1	
leaf02	x86_64 Intel Core i7 9xx (Nehalem Cla N/A ss Core i7)	1	
leaf03	x86_64 Intel Core i7 9xx (Nehalem Cla N/A ss Core i7)	1	
eaf04	x86_64 Intel Core i7 9xx (Nehalem Cla N/A ss Core i7)	1	
oob-mgmt-s	erver x86_64 Intel Core i7 9xx (Nehalem Cla	N/A	1
	ss Core i7)		
server01	x86_64 Intel Core i7 9xx (Nehalem Cla N/A ss Core i7)	1	
server02	x86_64 Intel Core i7 9xx (Nehalem Cla N/A ss Core i7)	1	
server03	x86_64 Intel Core i7 9xx (Nehalem Cla N/A ss Core i7)	1	
server04	x86_64 Intel Core i7 9xx (Nehalem Cla N/A ss Core i7)	1	
spine01	x86_64 Intel Core i7 9xx (Nehalem Cla N/A ss Core i7)	1	
spine02	x86_64 Intel Core i7 9xx (Nehalem Cla N/A ss Core i7)	1	

You can filter the results to view CPU information for a single switch, as shown here for *server02*.

View Information about the Disk on a Switch

You can view the name or operating system, type, transport, size, vendor, and model of the disk on a single device or all devices. This example shows all of these disk characteristics for all devices.

cumulus@s	switch:~\$ ne	etq sho	w inventor	y disk			
Matching ir	nventory re	cords:					
Hostname	Name	-	Туре	Transport	Size	Vendor	
Model							
leaf01	vda	disk	N/A	6G	0x1af4	N/A	
leaf02	vda	disk	N/A	6G	0x1af4	N/A	
leaf03	vda	disk	N/A	6G	0x1af4	N/A	
leaf04	vda	disk	N/A	6G	0x1af4	N/A	
oob-mamt-	-server vda	a	disk	N/A	256G	0x1af4	N/A

server01	vda	disk	N/A	301G	0x1af4	N/A
server02	vda	disk	N/A	301G	0x1af4	N/A
server03	vda	disk	N/A	301G	0x1af4	N/A
server04	vda	disk	N/A	301G	0x1af4	N/A
spine01	vda	disk	N/A	6G	0x1af4	N/A
spine02	vda	disk	N/A	6G	0x1af4	N/A

You can filter the results of the command to view the disk information for a particular device. This example shows disk information for *leaf03* switch.

View Memory Information for a Switch

You can view the name, type, size, speed, vendor, and serial number for the memory installed in a single device or all devices. This example shows all of these characteristics for all devices.

cumulus@switch:~\$ netq show inventory memory

Matching inventory records:

Hostname Name Type Size Speed Vendor Serial No

dell-z9100-	05 DIMM0 BANK 0 DDR3 8192 MB 1600 MHz
	14391421
	5 DIMMO BANK 0 DDR3 8192 MB 1600 MHz InnoDisk
	n 0000000
·	I-05 ChannelA-DIMM0 DDR3 8192 MB 1600 MHz
017A	87416232
E	ANK 0
mlx-2700-1	1 ChannelA-DIMM0 DDR3 8192 MB 1600 MHz
017A	73215444
Е	ANK 0
mlx-2700-1	1 ChannelB-DIMM0 DDR3 8192 MB 1600 MHz
017A	73215444
Е	ANK 2
qct-ix1-08	N/A N/A 7907.45MB N/A N/A N/A
qct-ix7-04	DIMMO BANK 0 DDR3 8192 MB 1600 MHz
Transcend	00211415
st1-l1	DIMM0 BANK 0 DDR3 4096 MB 1333 MHz N/A N/A
st1-l2	DIMMO BANK 0 DDR3 4096 MB 1333 MHz N/A N/A
st1-l3	DIMMO BANK 0 DDR3 4096 MB 1600 MHz N/A N/A
st1-s1	A1_DIMM0 A1_BAN DDR3 8192 MB 1333 MHz
A1_Manufa	cturer0 A1_SerNum0
K	0
st1-s2	A1_DIMM0 A1_BAN DDR3 8192 MB 1333 MHz
A1_Manufa	cturer0 A1_SerNum0
K	0

You can filter the results of the command to view devices with a particular memory type or vendor. This example shows all of the devices with memory from *QEMU* .

Hostname	Name	Туре	Size S	Speed Ver	ndor	Serial No
eaf01	DIMM 0	RAM	1024 MB	Unknown	QEMU	Not
Specified						
eaf02	DIMM 0	RAM	1024 MB	Unknown	QEMU	Not
Specified						
eaf03	DIMM 0	RAM	1024 MB	Unknown	QEMU	Not
Specified						
eaf04	DIMM 0	RAM	1024 MB	Unknown	QEMU	Not
Specified						
oob-mgmt-	server DIMI	M 0 RAM	409	6 MB Unkr	nown	
QEMU	Not Spec	cified				
server01	DIMM 0	RAM	512 MB	Unknown	QEMU	Not
Specified						
server02	DIMM 0	RAM	512 MB	Unknown	QEMU	Not
Specified						
server03	DIMM 0	RAM	512 MB	Unknown	QEMU	Not
Specified						
server04	DIMM 0	RAM	512 MB	Unknown	QEMU	Not
Specified						
spine01	DIMM 0	RAM	1024 MB	Unknown	QEMU	Not
Specified						

Monitor Switch Hardware and Software

> You can filter the results to view memory information for a single switch, as shown here for leaf01.

Serial No
Not

View a Summary of Physical Inventory for the NetQ or NetQ Cloud Appliance

Using the opta option lets you view inventory information for the NetQ or NetQ Cloud Appliance(s) rather than all network nodes. This example give you a summary of the inventory on the device.

cumulus@spin	ie-1:mgmt-vr	f:~\$ netq s	how inver	ntory brie	f opta
Matching inve	ntory records	;.			
Hostname	Switch	OS	CPU	ASIC	Ports
10-20-14-158	VX	CL	x86_64	VX	N/A

View Memory for the NetQ or NetQ Cloud Appliance

You can be specific about which inventory item you want to view for an appliance. This example shows the memory information for a NetQ Appliance, letting you verify you have sufficient memory.

cumulus@netq-appliance:~\$ netq show inventory memory opta Matching inventory records:							
Hostname	Name	Туре	Size Speed Vendor	Serial No			
netq-app Specified	DIMM 0	RAM	64 GB Unknown QEMU	Not			

View Fan Health for All Switches

Fan, power supply unit, and temperature sensors are available to provide additional data about the NetQ Platform operation. To view the health of fans in your switches, use the **netq show sensors fan** command. If you name the fans in all of your switches consistently, you can view more information at once.

In this example, we look at the state of all fans with the name fan1.

cumulus@swi	itch:~\$ netq s	how sensors fan fan1				
Hostname Name		Description		ate	Speed	Max
Min Messa	ge	Last Changed				
exit01 f	an1 far	n tray 1, fan 1	ok	2500	290	000
		-				

2500		Fri Apr 19 16:01:1	7 2019				
exit02	fan1	fan tray 1, fan 1	ok	2500	29000		
2500		Fri Apr 19 16:01:3	3 2019				
leaf01	fan1	fan tray 1, fan 1	ok	2500	29000		
2500		Sun Apr 21 20:07:	12 2019				
leaf02	fan1	fan tray 1, fan 1	ok	2500	29000		
2500		Fri Apr 19 16:01:4	1 2019				
leaf03	fan1	fan tray 1, fan 1	ok	2500	29000		
2500 Fri Apr 19 16:01:44 2019							
leaf04	fan1	fan tray 1, fan 1	ok	2500	29000		
2500 Fri Apr 19 16:01:36 2019							
spine01	fan1	fan tray 1, fan 1	ok	2500	29000		
2500		Fri Apr 19 16:01:5	2 2019				
spine02	fan1	fan tray 1, fan 1	ok	2500	29000		
2500		Fri Apr 19 16:01:0	8 2019				



Use tab completion to determine the names of the fans in your switches:

cumulus@switch:~\$ netq show sensors fan <<pre>cress tab>></press tab>>

around: Go back in time to around ...

fan1: Fan Name

fan2: Fan Name

fan3: Fan Name

fan4: Fan Name

fan5: Fan Name

fan6: Fan Name

json: Provide output in JSON

psu1fan1: Fan Name

psu2fan1 : Fan Name

<ENTER>

To view the status for a particular switch, use the optional hostname parameter.

cumulus@switch:~\$ netq leaf01 show sensors fan fan1

Description State Speed Max Hostname Name

Last Changed Min Message

leaf01	fan1	fan tray 1, fan 1	ok	2500	29000	
2500		Sun Apr 21 20:07	:12 2019			

View PSU Health for All Switches

Fan, power supply unit, and temperature sensors are available to provide additional data about the NetQ Platform operation. To view the health of PSUs in your switches, use the **netq show sensors psu** command. If you name the PSUs in all of your switches consistently, you can view more information at once.

In this example, we look at the state of all PSUs with the name *psu2*.

cumulus@switch:~\$ netq show sensors psu psu2 Matching sensors records:							
Hostname	Name	State	Message	Last Changed			
exit01	psu2	ok		Fri Apr 19 16:01:17 2019			
exit02	psu2	ok		Fri Apr 19 16:01:33 2019			
leaf01	psu2	ok		Sun Apr 21 20:07:12 2019			
leaf02	psu2	ok		Fri Apr 19 16:01:41 2019			
leaf03	psu2	ok		Fri Apr 19 16:01:44 2019			
leaf04	psu2	ok		Fri Apr 19 16:01:36 2019			
spine01	psu2	ok		Fri Apr 19 16:01:52 2019			
spine02	psu2	ok		Fri Apr 19 16:01:08 2019			



TIP

Use Tab completion to determine the names of the PSUs in your switches. Use the optional hostname parameter to view the PSU state for a given switch.

View the Temperature in All switches

Fan, power supply unit, and temperature sensors are available to provide additional data about the NetQ Platform operation. To view the temperature sensor status, current temperature, and configured threshold values, use the **netq show sensors temp** command. If you name the temperature sensors in all of your switches consistently, you can view more information at once.

In this example, we look at the state of all temperature sensors with the name *psu1temp1*.

cumulus@switch:~\$ netq show sensors temp psu2temp1 Matching sensors records:								
_		Description	State	Temp	Critical			
Max Min Message Last Changed								
exit01	psu2temp1	psu2 temp sensor	ok	25	85	80		
5	Fri Apr 19 16:01:17 2019							
exit02	psu2temp1	psu2 temp sensor	ok	25	85	80		
5		Fri Apr 19 16:01:33 2019						

leaf01 5	psu2temp1	psu2 temp sensor Sun Apr 21 20:07:12 2019	ok	25	85	80
leaf02 5	psu2temp1	psu2 temp sensor Fri Apr 19 16:01:41 2019	ok	25	85	80
leaf03 5	psu2temp1	psu2 temp sensor Fri Apr 19 16:01:44 2019	ok	25	85	80
leaf04 5	psu2temp1	psu2 temp sensor Fri Apr 19 16:01:36 2019	ok	25	85	80
spine01	psu2temp1	psu2 temp sensor Fri Apr 19 16:01:52 2019	ok	25	85	80
spine02 5	psu2temp1	psu2 temp sensor Fri Apr 19 16:01:08 2019	ok	25	85	80



Use Tab completion to determine the names of the temperature sensors in your switches. Use the optional hostname parameter to view the temperature state, current temperature, and threshold values for a given switch.

View All Sensor Data

To view all fan data, all PSU data, or all temperature data from the sensors, you must view all of the sensor data. The more consistently you name your sensors, the easier it will be to view the full sensor data.

Hostname	Name	Description	State	Message	
Last Ch		Description.	State	message	
exit01	fan1	fan tray 1, fan 1	ok		Fri
Apr 19 16:0	1:17 2019				
exit01	fan2	fan tray 1, fan 2	ok		Fri
Apr 19 16:0	1:17 2019				
exit01	fan3	fan tray 2, fan 1	ok		Fri
Apr 19 16:0	1:17 2019				
exit01	fan4	fan tray 2, fan 2	ok		Fri
Apr 19 16:0	1:17 2019				
exit01	fan5	fan tray 3, fan 1	ok		Fri
Apr 19 16:0	1:17 2019				
exit01	fan6	fan tray 3, fan 2	ok		Fri
Apr 19 16:0	1:17 2019				
exit01	psu1fan1	psu1 fan	ok		Fri
Apr 19 16:0	1:17 2019				
exit01	psu1temp	1 psu1 temp sensor	ok		
Fri Apr 19	16:01:17 20	19			
exit01	psu2fan1	psu2 fan	ok		Fri

Apr 19 16:01:17 2019			
exit01 psu2tem	p1 psu2 temp sensor	ok	
Fri Apr 19 16:01:17 2	019		
exit01 temp1	board sensor near cp	u ok	
Fri Apr 19 16:01:17 2	019		
exit01 temp2	board sensor near vir	tual switch ok	
Fri Apr 19 16:01:1	7 2019		
exit01 temp3	board sensor at front	left corner ok	
Fri Apr 19 16:01:17	7 2019		
exit01 temp4	board sensor at front	right corner ok	
Fri Apr 19 16:01:1	7 2019		
exit01 temp5	board sensor near fa	n ok	
Fri Apr 19 16:01:17 20	19		
exit02 fan1	fan tray 1, fan 1	ok	Fri
Apr 19 16:01:33 2019			
exit02 fan2	fan tray 1, fan 2	ok	Fri
Apr 19 16:01:33 2019			
exit02 fan3	fan tray 2, fan 1	ok	Fri
Apr 19 16:01:33 2019			
exit02 fan4	fan tray 2, fan 2	ok	Fri
Apr 19 16:01:33 2019			
exit02 fan5	fan tray 3, fan 1	ok	Fri
Apr 19 16:01:33 2019			
exit02 fan6	fan tray 3, fan 2	ok	Fri
Apr 19 16:01:33 2019			
exit02 psu1fan1	psu1 fan	ok	Fri
Apr 19 16:01:33 2019			
exit02 psu1tem	p1 psu1 temp sensor	ok	

Fri Apr 19 16:01:33 2019

• • •

View All Sensor-related Events

You can view the events that are triggered by the sensors using the **netq show events** command. You can narrow the focus to only critical events using the severity **level** option.

cumulus@switch:~\$ netq show events type sensors

No matching events records found

cumulus@switch:~\$ netq show events level critical type sensors

No matching events records found

View Interface Statistics and Utilization

NetQ Agents collect performance statistics every 30 seconds for the physical interfaces on switches and hosts in your network. The NetQ Agent does not collect statistics for non-physical interfaces, such as bonds, bridges, and VXLANs. The NetQ Agent collects the following statistics:

- Statistics
 - **Transmit**: tx_bytes, tx_carrier, tx_colls, tx_drop, tx_errs, tx_packets
 - Receive: rx_bytes, rx_drop, rx_errs, rx_frame, rx_multicast, rx_packets
- Utilization
 - rx_util, tx_util
 - port speed

These can be viewed using the following NetQ CLI commands:

Where the various options are:

- hostname limits the output to a particular switch
- errors limits the output to only the transmit and receive errors found on the designated interfaces
- physical-port limits the output to a particular port
- around enables viewing of the data at a time in the past
- json outputs results in json format
- text-port limits output to a particular host and port; hostname is required with this option
- tx, rx limits output to the transmit or receive values, respectively

In this example, we view the interface statistics for all switches and all of their physical interfaces.

cumulus@s	witch:~\$ neta s	how interface-s	tats		
	roc_dev_stats re		tats		
Hostname	Interface	Duratio	n RX Bytes	RX Drop	RX
Errors	TX Bytes	TX Drop	TX Errors	Last Changed	TVX

edge01	eth0	30	2278 0	16
4007	0	0	Mon Jun 3 23:03:14 2019	
edge01	lo	30	864 0	0
864	0	0	Mon Jun 3 23:03:14 2019	
exit01	bridge	60	336 0	0
1176	0	0	Mon Jun 3 23:02:27 2019	
exit01	eth0	30	3424 0	0
6965	0	0	Mon Jun 3 23:02:58 2019	
exit01	mgmt	30	2682 0	0
7488	0	0	Mon Jun 3 23:02:58 2019	
exit01	swp44	30	2457 0	0
2457	0	0	Mon Jun 3 23:02:58 2019	
exit01	swp51	30	2462 0	0
1769	0	0	Mon Jun 3 23:02:58 2019	
exit01	swp52	30	2634 0	0
2629	0	0	Mon Jun 3 23:02:58 2019	
exit01	vlan4001	50	336 0	0
1176	0	0	Mon Jun 3 23:02:27 2019	
exit01	vrf1	60	1344 0	0
0	0	0	Mon Jun 3 23:02:27 2019	
exit01	vxlan4001	50	336 0	0
1368	0	0	Mon Jun 3 23:02:27 2019	
exit02	bridge	61	1008 0	0
392	0	0	Mon Jun 3 23:03:07 2019	
exit02	eth0	20	2711 0	0
4983	0	0	Mon Jun 3 23:03:07 2019	
exit02	mgmt	30	2162 0	0
5506	0	0	Mon Jun 3 23:03:07 2019	
exit02	swp44	20	3040 0	0

```
3824 0 0 Mon Jun 3 23:03:07 2019 ...
```

In this example, we view the interface statistics for switch port 29.

cumulus@sv	vitch:~\$ netq sl	how interface-sta	ts swp29	
Matching pr	oc_dev_stats re	ecords:		
Hostname	Interface	RX Bytes	RX Drop	RX Errors
TX Bytes	TX Drop	TX Errors	Last Updated	
spine01	swp29	12853778	0	0
13281292	0	0 We	d Sep 25 14:43:1	7 2019
spine02	swp29	11739987	0	0
13316634	0	0 We	d Sep 25 14:43:32	2 2019

In this example, we view the utilization for the leaf03 switch.

Matching po	ort_stats records	:			
Hostname	Interface	RX Bytes	RX Drop	RX Erro	ors
RX Util	TX Bytes	TX Drop	TX Errors	TX Util	Port
Speed	Last Changed				

View Interface Statistics and Utilization

leaf03	bond03		4447		0		0	0
5041	0	0		0		NA		Wed Sep 25
14:46:16								
2019								
leaf03	bond04		3811		0		0	0
4957	0	0		0		NA		Wed Sep 25
14:46:16								
2019								
	bridge	5	540		0		0	0
476		0		0		NA		Wed Sep 25
14:46:16								
2019								
leaf03	eth0	34	471		0		0	
0.00033102	10480		0		0		0.000	999451
1G	Wed Sep 25	14:46	:16					
2019								
•••								

In this example, we view the transmit utilization only.

cumulus@switch:~\$ netq show interface-utils tx

Matching port_stats records:

Hostname Interface TX Bytes TX Drop TX Errors

TX Util Port Speed Last Changed



Monitor Switch Software Information

The syntax for this command is:

netq [<hostname>] show agents

netq [<hostname>] show inventory brief [json]

netq [<hostname>] show inventory license [cumulus] [status ok|status missing]

[around <text-time>] [json]

netq [<hostname>] show inventory os [version <os-version>| name <os-name>]

[json]

netq [<hostname>] show events [level info|level error|level warning|level critical|

level debug] [type license | type os] [between <text-time> and <text-endtime>]

[ison]



(i) NOTE

The values for the name option are specific to your deployment. For example, if you have devices with only one type of OS, say Cumulus Linux, then that is the only option available for the os-name option value. If you have multiple OSs running, say you also have Ubuntu, then that would also be an option for you.



(i) NOTE

When entering a time value, you must include a numeric value and the unit of measure:

- w: week(s)
- d: day(s)
- h: hour(s)
- m: minute(s)
- s: second(s)
- now

For time ranges, the <text-time> is the most recent time and the <textendtime> is the oldest time. The values do not have to have the same unit of measure.

View OS Information for a Switch

You can view the name and version of the OS on a switch, and when it was last modified. This example shows the OS information for all devices.

Matching ir				
_	nventory	records:		
Hostname	Nai	me V	ersion	Last Changed
edge01	Ubur	ntu 16	.04	Fri Apr 19 16:01:18 2019
exit01	CL	3.7.5		Fri Apr 19 16:01:13 2019
exit02	CL	3.7.5		Fri Apr 19 16:01:38 2019

Monitor Switch Hardware and Software

Monitor Switch Software Information

leaf01	CL	3.7.5	Sun Apr 21 20:07:09 2019
leaf02	CL	3.7.5	Fri Apr 19 16:01:46 2019
leaf03	CL	3.7.5	Fri Apr 19 16:01:41 2019
leaf04	CL	3.7.5	Fri Apr 19 16:01:32 2019
server01	Ubuntu	16.04	Fri Apr 19 16:01:55 2019
server02	Ubuntu	16.04	Fri Apr 19 16:01:55 2019
server03	Ubuntu	16.04	Fri Apr 19 16:01:55 2019
server04	Ubuntu	16.04	Fri Apr 19 16:01:55 2019
spine01	CL	3.7.5	Fri Apr 19 16:01:49 2019
spine02	CL	3.7.5	Fri Apr 19 16:01:05 2019

You can filter the results of the command to view only devices with a particular operating system or version. This can be especially helpful when you suspect that a particular device has not been upgraded as expected. This example shows all devices with the Cumulus Linux version 3.7.5 installed.

cumulus@switch:~\$ netq show inventory os version 3.7.5							
Matching inventory records:							
Hostname	Name	Version	Last Changed				
exit01	CL	3.7.5	Fri Apr 19 16:01:13 2019				
exit02	CL	3.7.5	Fri Apr 19 16:01:38 2019				
leaf01	CL	3.7.5	Sun Apr 21 20:07:09 2019				
leaf02	CL	3.7.5	Fri Apr 19 16:01:46 2019				
leaf03	CL	3.7.5	Fri Apr 19 16:01:41 2019				
leaf04	CL	3.7.5	Fri Apr 19 16:01:32 2019				
spine01	CL	3.7.5	Fri Apr 19 16:01:49 2019				
spine02	CL	3.7.5	Fri Apr 19 16:01:05 2019				

This example shows changes that have been made to the OS on all devices between 16 and 21 days ago. Remember to use measurement units on the time values.

Matching inver	ntory records:				
Hostname	Name	Version	DB State	Last Changed	
 mlx-2410a1-05 18:30:53 2019	Cumulus Li	nux 3.7.3	Add	Tue Feb 12	
mlx-2700-11 18:30:45 2019	Cumulus Lin	ux 3.7.3	Add	Tue Feb 12	
mlx-2100-05 18:30:26 2019	Cumulus Lin	ux 3.7.3	Add	Tue Feb 12	
mlx-2100-05	Cumulus Lin	ux 3.7.3~153326317	4.bce9472	Add W	/ed

Feb 13 11:10:47	7 2019				
mlx-2700-11	Cumulus Linux	3.7.3~1533263174.bce9472	2	Add	Wed
Feb 13 11:10:38	8 2019				
mlx-2100-05	Cumulus Linux	3.7.3~1533263174.bce9472	2	Add	Wed
Feb 13 11:10:42	2 2019				
mlx-2700-11	Cumulus Linux	3.7.3~1533263174.bce9472	2	Add	Wed
Feb 13 11:10:51	1 2019				

View License Information for a Switch

You can view the name and current state of the license (whether it valid or not), and when it was last updated for one or more devices. If a license is no longer valid on a switch, it does not operate correctly. This example shows the license information for all devices.

cumulus@s	cumulus@switch:~\$ netq show inventory license								
Matching ir	Matching inventory records:								
Hostname	Name State	Last Changed							
edge01	Cumulus Linux N/A								
exit01	Cumulus Linux ok	Fri Apr 19 16:01:13 2019							
exit02	Cumulus Linux ok	Fri Apr 19 16:01:38 2019							
leaf01	Cumulus Linux ok	Sun Apr 21 20:07:09 2019							
leaf02	Cumulus Linux ok	Fri Apr 19 16:01:46 2019							
leaf03	Cumulus Linux ok	Fri Apr 19 16:01:41 2019							
leaf04	Cumulus Linux ok	Fri Apr 19 16:01:32 2019							
server01	Cumulus Linux N/A	Fri Apr 19 16:01:55 2019							
server02	Cumulus Linux N/A	Fri Apr 19 16:01:55 2019							

server03	Cumulus Linux N/A	Fri Apr 19 16:01:55 2019	
server04	Cumulus Linux N/A	Fri Apr 19 16:01:55 2019	
spine01	Cumulus Linux ok	Fri Apr 19 16:01:49 2019	
spine02	Cumulus Linux ok	Fri Apr 19 16:01:05 2019	

You can view the historical state of licenses using the around keyword. This example shows the license state for all devices about 7 days ago. Remember to use measurement units on the time values.

cumulus@s	cumulus@switch:~\$ netq show inventory license around 7d									
Matching ir	Matching inventory records:									
Hostname	Name State	Last Changed								
edge01	Cumulus Linux N/A	Tue Apr 2 14:01:18 2019								
exit01	Cumulus Linux ok	Tue Apr 2 14:01:13 2019								
exit02	Cumulus Linux ok	Tue Apr 2 14:01:38 2019								
leaf01	Cumulus Linux ok	Tue Apr 2 20:07:09 2019								
leaf02	Cumulus Linux ok	Tue Apr 2 14:01:46 2019								
leaf03	Cumulus Linux ok	Tue Apr 2 14:01:41 2019								
leaf04	Cumulus Linux ok	Tue Apr 2 14:01:32 2019								
server01	Cumulus Linux N/A	Tue Apr 2 14:01:55 2019								
server02	Cumulus Linux N/A	Tue Apr 2 14:01:55 2019								
server03	Cumulus Linux N/A	Tue Apr 2 14:01:55 2019								
server04	Cumulus Linux N/A	Tue Apr 2 14:01:55 2019								
spine01	Cumulus Linux ok	Tue Apr 2 14:01:49 2019								
spine02	Cumulus Linux ok	Tue Apr 2 14:01:05 2019								

You can filter the results to show license changes during a particular timeframe for a particular device. This example shows that there have been no changes to the license state on spine01 between now and 24 hours ago.

cumulus@switch:~\$ netq spine01 show events type license between now and 24h No matching events records found

View Summary of Operating System on a Switch

As with the hardware information, you can view a summary of the software information using the **brief** keyword. Specify a hostname to view the summary for a specific device.

cumulus@switch:~\$ netq show inventory brief								
Matching in	ventory reco	ords:						
Hostname	Switch	OS		CP	U	ASIC		Ports
edge01	N/A	Ubuntu	 I	x8(5_64	N/A		N/A
exit01	VX	CL	x86_	64	VX		N/A	
exit02	VX	CL	x86_	64	VX		N/A	
leaf01	VX	CL	x86_	64	VX		N/A	
leaf02	VX	CL	x86_	64	VX		N/A	
leaf03	VX	CL	x86_	64	VX		N/A	
leaf04	VX	CL	x86_	64	VX		N/A	
server01	N/A	Ubuntu	ı	x8	6_64	N/A		N/A
server02	N/A	Ubuntu	ı	x8	6_64	N/A		N/A
server03	N/A	Ubuntu	ı	x8	6_64	N/A		N/A
server04	N/A	Ubuntu	J	x8	6_64	N/A		N/A

spine01	VX	CL	x86_64 VX	N/A	
spine02	VX	CL	x86_64 VX	N/A	

Validate NetQ Agents are Running

You can confirm that NetQ Agents are running on switches and hosts (if installed) using the **netq show agents** command. Viewing the **Status** column of the output indicates whether the agent is up and current, labelled *Fresh*, or down and stale, labelled *Rotten*. Additional information is provided about the agent status, including whether it is time synchronized, how long it has been up, and the last time its state changed.

This example shows NetQ Agent state on all devices.

```
cumulus@switch:~$ netq show agents
Matching agents records:
                                     Sys Uptime
Hostname Status
                     NTP Sync Version
Agent Uptime Reinitialize Time Last Changed
edge01 Fresh yes 2.1.0-ub16.04u15~1555612152.6e34b56 2d:7h:
2m:12s
      2d:7h:2m:5s
                                          Sun Apr 21 16:00:50 2019
                          2d:7h:2m:5s
                       2.1.0-cl3u15~1555612272.6e34b56
                                                    2d:7h:1m:
exit01 Fresh
                  yes
30s
        2d:7h:1m:22s 2d:7h:1m:22s
                                        Sun Apr 21 16:00:52 2019
exit02
              yes 2.1.0-cl3u15~1555612272.6e34b56
                                                    2d:7h:1m:
        Fresh
36s
        2d:7h:1m:27s
                       2d:7h:1m:27s
                                        Sun Apr 21 16:01:19 2019
              yes 2.1.0-cl3u15~1555612272.6e34b56 2d:7h:1m:
leaf01
        Fresh
        2h:54m:12s
28s
                       2h:54m:12s
                                       Sun Apr 21 20:05:45 2019
leaf02 Fresh
              yes
                       2.1.0-cl3u15~1555612272.6e34b56
                                                    2d:7h:1m:
        2d:7h:1m:29s
38s
                        2d:7h:1m:29s
                                        Sun Apr 21 16:01:43 2019
```

leaf03	Fresh ye	es	2.1.0-cl3u15~1555612	272.6e34b56	2d:7h:1m:
37s	2d:7h:1m:28s		2d:7h:1m:28s	Sun Apr 21 16:	01:23 2019
leaf04	Fresh ye	es	2.1.0-cl3u15~1555612	272.6e34b56	2d:7h:1m:
39s	2d:7h:1m:31s		2d:7h:1m:31s	Sun Apr 21 16:	01:27 2019
server01	Fresh	yes	2.1.0-ub16.04u15~1	555612152.6e3	4b56 2d:6h:
59m:35s	2d:6h:59m:	:27s	2d:6h:59m:27s	Sun Apr	21 16:00:43
2019					
server02	Fresh	yes	2.1.0-ub16.04u15~1	555612152.6e3	4b56 2d:6h:
59m:34s	2d:6h:59m:	:26s	2d:6h:59m:26s	Sun Apr	21 16:00:46
2019					
server03	Fresh	yes	2.1.0-ub16.04u15~1	555612152.6e34	4b56 2d:6h:
59m:34s	2d:6h:59m:	:26s	2d:6h:59m:26s	Sun Apr	21 16:00:52
2019					
server04	Fresh	yes	2.1.0-ub16.04u15~1	555612152.6e34	4b56 2d:6h:
59m:34s	2d:6h:59m	:26s	2d:6h:59m:26s	Sun Apr	21 16:00:43
2019					
spine01	Fresh	yes	2.1.0-cl3u15~155561	2272.6e34b56	2d:7h:1m:
40s	2d:7h:1m:32s		2d:7h:1m:32s	Sun Apr 21 16:	01:33 2019
spine02	Fresh y	yes	2.1.0-cl3u15~155561	2272.6e34b56	2d:7h:1m:
34s	2d:7h:1m:26s		2d:7h:1m:26s	Sun Apr 21 16:	01:12 2019

You can narrow your focus in several ways:

- View the state of the NetQ Agent on a given device using the hostname keyword.
- View only the NetQ Agents that are fresh or rotten using the *fresh* or *rotten* keyword.
- View the state of NetQ Agents at an earlier time using the *around* keyword.

Monitor Software Services

Cumulus Linux and NetQ run a number of services to deliver the various features of these products. You can monitor their status using the **netq**show services command. The services related to system-level operation are described here. Monitoring of other services, such as those related to routing, are described with those topics. NetQ automatically monitors the following services:

- bgpd: BGP (Border Gateway Protocol) daemon
- **clagd**: MLAG (Multi-chassis Link Aggregation) daemon
- helpledmgrd: Switch LED manager daemon
- **Ildpd**: LLDP (Link Layer Discovery Protocol) daemon
- mstpd: MSTP (Multiple Spanning Tree Protocol) daemon
- neighmgrd: Neighbor Manager daemon for BGP and OSPF
- netq-agent: NetQ Agent service
- netqd: NetQ application daemon
- **ntp**: NTP service
- **ntpd**: NTP daemon
- ptmd: PTM (Prescriptive Topology Manager) daemon
- **pwmd**: PWM (Password Manager) daemon
- rsyslog: Rocket-fast system event logging processing service
- **smond**: System monitor daemon
- ssh: Secure Shell service for switches and servers
- status: License validation service
- syslog: System event logging service
- vrf: VRF (Virtual Route Forwarding) service
- zebra: GNU Zebra routing daemon

The CLI syntax for viewing the status of services is:

netq [<hostname>] show services [<service-name>] [vrf <vrf>] [active | monitored] [around <text-time>] [json] netq [<hostname>] show services [<service-name>] [vrf <vrf>] status (ok | warning | error|fail) [around <text-time>] [json] netq [<hostname>] show events [level info | level error | level warning | level critical | level debug] type services [between <text-time> and <text-endtime>] [ison]

View All Services on All Devices

This example shows all of the available services on each device and whether each is enabled, active, and monitored, along with how long the service has been running and the last time it was changed.



It is useful to have colored output for this show command. To configure colored output, run the netq config add color command.

cumulus@s	switch:~\$ netq	show services		
Hostname	Service	PID VRF	Enabled Active Monitore	d
Status	Uptime	Last Change	ed	
leaf01	bgpd	2872 default	yes yes yes ok	1d:6h:
43m:59s	Fri Feb 15	17:28:24 2019		
leaf01	clagd	n/a default	yes no yes n/a	1d:6h:

43m:35s	Fri Feb 15 17:28:48 2019	
leaf01	ledmgrd 1850 default yes yes no ok	1d:6h:
43m:59s	Fri Feb 15 17:28:24 2019	
leaf01	lldpd 2651 default yes yes ok	1d:6h:
43m:27s	Fri Feb 15 17:28:56 2019	
leaf01	mstpd 1746 default yes yes yes ok	1d:6h:
43m:35s	Fri Feb 15 17:28:48 2019	
leaf01	neighmgrd 1986 default yes yes no ok	1d:
6h:43m:59s	Fri Feb 15 17:28:24 2019	
leaf01	netq-agent 8654 mgmt yes yes ok	1d:
6h:43m:29s	Fri Feb 15 17:28:54 2019	
leaf01	netqd 8848 mgmt yes yes ok	1d:6h:
43m:29s	Fri Feb 15 17:28:54 2019	
leaf01	ntp 8478 mgmt yes yes ok	1d:6h:
43m:29s	Fri Feb 15 17:28:54 2019	
leaf01	ptmd 2743 default yes yes no ok	1d:6h:
43m:59s	Fri Feb 15 17:28:24 2019	
leaf01	pwmd 1852 default yes yes no ok	1d:6h:
43m:59s	Fri Feb 15 17:28:24 2019	
leaf01	smond 1826 default yes yes ok	1d:6h:
43m:27s	Fri Feb 15 17:28:56 2019	
leaf01	ssh 2106 default yes yes no ok	1d:6h:
43m:59s	Fri Feb 15 17:28:24 2019	
leaf01	syslog 8254 default yes yes no ok	1d:6h:
43m:59s	Fri Feb 15 17:28:24 2019	
leaf01	zebra 2856 default yes yes ok	1d:6h:
43m:59s	Fri Feb 15 17:28:24 2019	
leaf02	bgpd 2867 default yes yes ok	1d:6h:
43m:55s	Fri Feb 15 17:28:28 2019	

```
leaf02 clagd n/a default yes no yes n/a 1d:6h:
43m:31s Fri Feb 15 17:28:53 2019
leaf02 ledmgrd 1856 default yes yes no ok 1d:6h:
43m:55s Fri Feb 15 17:28:28 2019
leaf02 lldpd 2646 default yes yes yes ok 1d:6h:
43m:30s Fri Feb 15 17:28:53 2019
...
```

You can also view services information in JSON format:

```
cumulus@switch:~$ netq show services json
{
  "services":[
    {
      "status":"ok",
      "uptime":1550251734.0,
      "monitored":"yes",
      "service":"ntp",
      "lastChanged":1550251734.4790000916,
      "pid":"8478",
      "hostname":"leaf01",
      "enabled":"yes",
      "vrf":"mgmt",
      "active":"yes"
    },
      "status":"ok",
      "uptime":1550251704.0,
```

```
"monitored":"no",
  "service":"ssh",
  "lastChanged":1550251704.0929999352,
  "pid":"2106",
  "hostname":"leaf01",
  "enabled":"yes",
  "vrf":"default",
  "active":"yes"
},
  "status":"ok",
  "uptime":1550251736.0,
  "monitored":"yes",
  "service":"lldpd",
  "lastChanged":1550251736.5160000324,
  "pid":"2651",
  "hostname":"leaf01",
  "enabled":"yes",
  "vrf":"default",
  "active":"yes"
},
  "status":"ok",
  "uptime":1550251704.0,
  "monitored":"yes",
  "service":"bgpd",
  "lastChanged":1550251704.1040000916,
  "pid":"2872",
  "hostname":"leaf01",
```

```
"enabled":"yes",
  "vrf":"default",
  "active":"yes"
},
  "status":"ok",
  "uptime":1550251704.0,
  "monitored":"no",
  "service":"neighmgrd",
  "lastChanged":1550251704.0969998837,
  "pid":"1986",
  "hostname":"leaf01",
  "enabled":"yes",
  "vrf":"default",
  "active":"yes"
},
```

If you want to view the service information for a given device, simply use the **hostname** option when running the command.

View Information about a Given Service on All Devices

You can view the status of a given service at the current time, at a prior point in time, or view the changes that have occurred for the service during a specified timeframe.

This example shows how to view the status of the NTP service across the network. In this case, VRF is configured so the NTP service runs on both the default and management interface. You can perform the same command with the other services, such as bgpd, lldpd, and clagd.

Matching s	ervices record	ds:					
Hostname	Service	PID VRF	En	abled	Active	Monitored	d
Status	Uptime	Last Chang	ed 				
exit01	 ntp	 8478 mgmt	ves	ves	ves	ok	1d:6h:
	·	5 17:28:54 2019	,	,	<i>y</i>		
exit02	ntp	8497 mgmt	yes	yes	yes	ok	1d:6h:
52m:36s	Fri Feb 1	5 17:28:59 2019					
firewall01	ntp	n/a default	yes	yes	yes	ok	1d:6h:
53m:4s	Fri Feb 15	17:28:31 2019					
hostd-11	ntp	n/a default	yes	yes	yes	ok	1d:6h:
52m:46s	Fri Feb 1	5 17:28:49 2019					
hostd-21	ntp	n/a default	yes	yes	yes	ok	1d:6h:
52m:37s	Fri Feb 1	5 17:28:58 2019					
hosts-11	ntp	n/a default	yes	yes	yes	ok	1d:6h:
52m:28s	Fri Feb 1	5 17:29:07 2019					
hosts-13	ntp	n/a default	yes	yes	yes	ok	1d:6h:
52m:19s	Fri Feb 1	5 17:29:16 2019					
hosts-21	ntp	n/a default	yes	yes	yes	ok	1d:6h:
52m:14s	Fri Feb 1	5 17:29:21 2019					
hosts-23	ntp	n/a default	yes	yes	yes	ok	1d:6h:
		17:29:31 2019					
-	•	2148 default	yes	yes	yes	ok	1d:6h:
		5 17:27:52 2019					
	·	2148 default	yes	yes	yes	ok	1d:6h:
		5 17:27:57 2019					
spine01	ntp	8414 mgmt	yes	yes	s yes	ok	1d:6h:

53m:30s	Fri Feb	15 17:28:05 2019					
spine02	ntp	8419 mgmt	yes	yes	yes	ok	1d:6h:
53m:27s	Fri Feb	15 17:28:08 2019					
spine03	ntp	8443 mgmt	yes	yes	yes	ok	1d:6h:
53m:22s	Fri Feb	15 17:28:13 2019					
leaf01	ntp	8765 mgmt	yes	yes	yes	ok	1d:6h:
52m:52s	Fri Feb	15 17:28:43 2019					
leaf02	ntp	8737 mgmt	yes	yes	yes	ok	1d:6h:
52m:46s	Fri Feb	15 17:28:49 2019					
leaf11	ntp	9305 mgmt	yes	yes	yes	ok	1d:6h:
49m:22s	Fri Feb	15 17:32:13 2019					
leaf12	ntp	9339 mgmt	yes	yes	yes	ok	1d:6h:
49m:9s	Fri Feb	15 17:32:26 2019					
leaf21	ntp	9367 mgmt	yes	yes	yes	ok	1d:6h:
49m:5s	Fri Feb	15 17:32:30 2019					
leaf22	ntp	9403 mgmt	yes	yes	yes	ok	1d:6h:
52m:57s	Fri Feb	15 17:28:38 2019					

This example shows the status of the BGP daemon.

	witch:~\$ netq s	show services bgp	d				
Hostname	Service	PID VRF	Ena	abled .	Active	Monitored	
Status	Uptime	Last Changed					
exit01	bgpd	2872 default	yes	yes	yes	ok	1d:6h:
54m:37s	Fri Feb 15	17:28:24 2019					

exit02	bgpd	2867 default	yes	yes	yes	ok	1d:6h:
54m:33s	Fri Feb 15	17:28:28 2019					
firewall01	bgpd	21766 default	ye	s ye	s yes	ok	1d:6h:
54m:54s	Fri Feb 15	17:28:07 2019					
spine01	bgpd	2953 default	yes	yes	yes	ok	1d:6h:
55m:27s	Fri Feb 15	17:27:34 2019					
spine02	bgpd	2948 default	yes	yes	yes	ok	1d:6h:
55m:23s	Fri Feb 15	17:27:38 2019					
spine03	bgpd	2953 default	yes	yes	yes	ok	1d:6h:
55m:18s	Fri Feb 15	17:27:43 2019					
leaf01	bgpd	3221 default	yes	yes	yes	ok	1d:6h:
54m:48s	Fri Feb 15	17:28:13 2019					
leaf02	bgpd	3177 default	yes	yes	yes	ok	1d:6h:
54m:42s	Fri Feb 15	17:28:19 2019					
leaf11	bgpd	3521 default	yes	yes	yes	ok	1d:6h:
51m:18s	Fri Feb 15	17:31:43 2019					
leaf12	bgpd	3527 default	yes	yes	yes	ok	1d:6h:
51m:6s	Fri Feb 15	17:31:55 2019					
leaf21	bgpd	3512 default	yes	yes	yes	ok	1d:6h:
51m:1s	Fri Feb 15	17:32:00 2019					
leaf22	bgpd	3536 default	yes	yes	yes	ok	1d:6h:
54m:54s	Fri Feb 15	17:28:07 2019					

View Events Related to a Given Service

To view changes over a given time period, use the **netq show events** command. For more detailed information about events, refer to **Monitor Events**.

In this example, we want to view changes to the bgpd service in the last 48 hours.

Hastrana	Ma	acago Typo Coverity Message	Timestama
Hostname		ssage Type Severity Message 	Timestamp
leaf01	bgp	info BGP session with peer spin	e-1 swp3. 1d:6h:55m:37s
		3 vrf DataVrf1081 state changed f	ro
		m failed to Established	
leaf01	bgp	info BGP session with peer spin	e-2 swp4. 1d:6h:55m:37s
		3 vrf DataVrf1081 state changed f	ro
		m failed to Established	
leaf01	bgp	info BGP session with peer spin	e-3 swp5. 1d:6h:55m:37s
		3 vrf DataVrf1081 state changed f	ro
		m failed to Established	
leaf01	bgp	info BGP session with peer spin	e-1 swp3. 1d:6h:55m:37s
		2 vrf DataVrf1080 state changed f	ro
		m failed to Established	
leaf01	bgp	info BGP session with peer spin	e-3 swp5. 1d:6h:55m:37s
		2 vrf DataVrf1080 state changed f	ro
		m failed to Established	
leaf01	bgp	info BGP session with peer spin	e-2 swp4. 1d:6h:55m:37s
		2 vrf DataVrf1080 state changed f	ro
		m failed to Established	
leaf01	bgp	info BGP session with peer spin	e-3 swp5. 1d:6h:55m:37s
		4 vrf DataVrf1082 state changed f	ro
		m failed to Established	

Monitor Physical Layer Components

With NetQ, a network administrator can monitor OSI Layer 1 physical components on network devices, including interfaces, ports, links, and peers. NetQ provides the ability to:

- Manage physical inventory: view the performance and status of various components of a switch or host server
- Validate configurations: verify the configuration of network peers and ports

It helps answer questions such as:

- Are any individual or bonded links down?
- Are any links flapping?
- Is there a link mismatch anywhere in my network?
- Which interface ports are empty?
- Which transceivers are installed?
- What is the peer for a given port?

NetQ uses LLDP (Link Layer Discovery Protocol) to collect port information. NetQ can also identify peer ports connected to DACs (Direct Attached Cables) and AOCs (Active Optical Cables) without using LLDP, even if the link is not UP.

Monitor Physical Layer Inventory

Keeping track of the various physical layer components in your switches and servers ensures you have a fully functioning network and provides inventory management and audit capabilities. You can monitor ports, transceivers, and cabling deployed on a per port (interface), per vendor, per part number and so forth. NetQ enables you to view the

current status and the status an earlier point in time. From this information, you can, among other things:

- determine which ports are empty versus which ones have cables plugged in and thereby validate expected connectivity
- audit transceiver and cable components used by vendor, giving you insights for estimated replacement costs, repair costs, overall costs, and so forth to improve your maintenance and purchasing processes
- identify changes in your physical layer, and when they occurred

The netq show interfaces physical command is used to obtain the information from the devices. Its syntax is:

netq [<hostname>] show interfaces physical [<physical-port>] [empty|plugged] [peer] [vendor <module-vendor>|model <module-model>|module] [around <text-time>] [json]

netq [<hostname>] show events [level info|level error|level warning|level critical| level debug] type interfaces-physical [between <text-time> and <text-endtime>] [json]



(i) NOTE

When entering a time value, you must include a numeric value and the unit of measure:

- d: day(s)
- h: hour(s)
- m: minute(s)
- s: second(s)
- now

For time ranges, the <text-time> is the most recent time and the <textendtime> is the oldest time. The values are not required to have the same unit of measure.

View Detailed Cable Information for All Devices

You can view what cables are connected to each interface port for all devices, including the module type, vendor, part number and performance characteristics. You can also view the cable information for a given device by adding a hostname to the show command. This example shows cable information and status for all interface ports on all devices.

cumulus@switch:~\$ netq show interfaces physical

Matching cables records:

Hostname Interface Speed AutoNeg Module State

Vendor Part No Last Changed

edge01 eth0 up 1G on RJ45 n/a n/ a Fri Jun 7 00:42:52 2019 edge01 eth1 down 1G off RJ45 n/a n/ a Fri Jun 7 00:42:52 2019 edge01 eth2 down 1G off RJ45 n/a n/ a Fri Jun 7 00:42:52 2019 edge01 vagrant down 1G on RJ45 n/a n/ a Fri Jun 7 00:42:52 2019 exit01 eth0 up 1G off RJ45 n/a n/a Fri Jun 7 00:42:52 2019 exit01 swp1 down Unknown off RJ45 n/a n/ a Fri Jun 7 00:43:03 2019 exit01 swp44 up 1G off RJ45 n/a n/ a Fri Jun 7 00:51:28 2019 exit01 swp45 down Unknown off RJ45 n/a n/ a Fri Jun 7 00:43:03 2019 exit01 swp46 down Unknown off RJ45 n/a n/ a Fri Jun 7 00:43:03 2019 exit01 swp47 down Unknown off RJ45 n/ n/a a Fri Jun 7 00:43:03 2019 exit01 swp48 down Unknown off RJ45 n/a n/ a Fri Jun 7 00:43:03 2019 exit01 swp49 down Unknown off RJ45 n/a n/ a Fri Jun 7 00:43:03 2019 exit01 swp50 down Unknown off RJ45 n/a n/ a Fri Jun 7 00:42:53 2019 exit01 swp51 up 1G off RJ45 n/a n/ a Fri Jun 7 00:51:28 2019

exit01	swp52	up	1G	off	RJ45	n/a		n/	
а	Fri Jun 7 00:51:28 2	.019							
exit01	vagrant	down	Un	known	off	RJ45	n/a		n/
а	Fri Jun 7 00:43:03 2	.019							
exit02	eth0	up	1G	off	RJ45	n/a		n/a	
Fri Jun	7 00:42:51 2019								
exit02	swp1	down	Unk	nown	off	RJ45	n/a		n/
а	Fri Jun 7 00:43:01 2	.019							
exit02	swp44	up	1G	off	RJ45	n/a		n/	
а	Fri Jun 7 00:51:28 2	.019							
exit02	swp45	down	Unl	known	off	RJ45	n/a		n/
а	Fri Jun 7 00:43:01 2	.019							
exit02	swp46	down	Unl	known	off	RJ45	n/a		n/
а	Fri Jun 7 00:43:01 2	.019							
exit02	swp47	down	Unl	known	off	RJ45	n/a		n/
а	Fri Jun 7 00:43:01 2	.019							
exit02	swp48	down	Unl	known	off	RJ45	n/a		n/
а	Fri Jun 7 00:43:01 2	.019							
exit02	swp49	down	Unl	known	off	RJ45	n/a		n/
а	Fri Jun 7 00:43:01 2	.019							
exit02	swp50	down	Unl	known	off	RJ45	n/a		n/
а	Fri Jun 7 00:43:01 2	.019							
exit02	swp51	up	1G	off	RJ45	n/a		n/	
а	Fri Jun 7 00:51:28 2	.019							
exit02	swp52	up	1G	off	RJ45	n/a		n/	
а	Fri Jun 7 00:51:28 2	.019							
exit02	vagrant	down	Un	known	off	RJ45	n/a		n/
a	Fri Jun 7 00:43:01 2	.019							
leaf01	eth0	up	1G	off	RJ45	n/a		n/a	

Fri Jun 7 00:43:02 2019						
leaf01 swp1	up	1G	off	RJ45	n/a	n/a
Fri Jun 7 00:52:03 2019						
leaf01 swp2	up	1G	off	RJ45	n/a	n/a
Fri Jun 7 00:52:03 2019						

View Detailed Module Information for a Given Device

You can view detailed information about the transceiver modules on each interface port, including serial number, transceiver type, connector and attached cable length. You can also view the module information for a given device by adding a hostname to the **show** command. This example shows the detailed module information for the interface ports on *leaf02* switch.

Matching cal	·	leaf02 show interfaces are:	physical i	nodule	
Hostname	Interface	Module Ven	dor	Part No	Serial
No	Transceiver	Connector Lengt	h Last Cha	nged	
leaf02	swp1	RJ45 n/a	n/a	n/a	n/
	•	RJ45 n/a Thu Feb 7 22:49:37 20		n/a	n/
	n/a	-	19		n/
a n/a	n/a swp2	Thu Feb 7 22:49:37 20	19 MC2	2609130-003	.,,
a n/a leaf02	n/a swp2 5177 10	Thu Feb 7 22:49:37 20 SFP Mellanox	19 MC2	2609130-003	
a n/a leaf02 MT1507VS05	n/a swp2 5177 10	Thu Feb 7 22:49:37 20 SFP Mellanox	19 MC2	2609130-003 3m Thu F	

leaf02 swp47 QSFP+ CISCO AFBR-7IER05Z-CS1

AVE1823402U n/a n/a 5m Thu Feb 7 22:49:37 2019

leaf02 swp48 QSFP28 TE Connectivity 2231368-1

15250052 100G Base-CR4 or n/a 3m Thu Feb 7 22:49:37 2019

25G Base-CR CA-L

,40G Base-

CR4

leaf02 swp49 SFP OEM SFP-10GB-LR

ACSLR130408 10G Base-LR LC 10km, Thu Feb 7 22:49:37 2019

10000m

leaf02 swp50 SFP JDSU PLRXPLSCS4322N

CG03UF45M 10G Base-SR,Mult LC 80m, Thu Feb 7 22:49:37 2019

imode, 30m,

50um

(M5), Multim 300m

ode,

62.5um (M6),Shor

twave laser w/o

OFC (SN),interme

diate distance (

I)

leaf02 swp51 SFP Mellanox MC2609130-003

MT1507VS05177 1000Base-CX,Copp Copper pigtail 3m Thu Feb 7

22:49:37 2019

er Passive,Twin

Axial Pair (TW)

leaf02 swp52 SFP FINISAR CORP. FCLF8522P2BTL
PTN1VH2 1000Base-T RJ45 100m Thu Feb 7 22:49:37 2019

View Ports without Cables Connected for a Given Device

Checking for empty ports enables you to compare expected versus actual deployment. This can be very helpful during deployment or during upgrades. You can also view the cable information for a given device by adding a hostname to the **show** command. This example shows the ports that are empty on leaf01 switch.

cumulus@switch:~\$ netq leaf01 show interfaces physical empty Matching cables records are: Hostname Interface State Speed AutoNeg Module Vendor Part No Last Changed leaf01 swp49 down Unknown on empty n/a Thu n/a Feb 7 22:49:37 2019 leaf01 swp52 down Unknown on empty n/a n/a Thu Feb 7 22:49:37 2019

View Ports with Cables Connected for a Given Device

In a similar manner as checking for empty ports, you can check for ports that have cables connected, enabling you to compare expected versus actual deployment. You can also view the cable information for a given device by adding a hostname to the show command. If you add the around keyword, you can view which interface ports had cables connected at a previous time. This example shows the ports of *leaf01* switch that have attached cables.

cumulus@switch:~\$ netq leaf01 show interfaces physical plugged Matching cables records: Hostname Interface State Speed AutoNeg Module Vendor Part No Last Changed leaf01 eth0 up 1G on RJ45 n/a n/a Thu Feb 7 22:49:37 2019 leaf01 swp1 up 10G off SFP Amphenol 610640005 Thu Feb 7 22:49:37 2019 leaf01 swp2 up 10G off SFP Amphenol 610640005 Thu Feb 7 22:49:37 2019 leaf01 swp3 down 10G off SFP Mellanox MC3309130-001 Thu Feb 7 22:49:37 2019 leaf01 swp33 down 10G off SFP OEM SFP-H10GB-CU1M Thu Feb 7 22:49:37 2019 leaf01 swp34 down 10G off SFP Amphenol 571540007 Thu Feb 7 22:49:37 2019 leaf01 swp35 down 10G off SFP Amphenol 571540007 Thu Feb 7 22:49:37 2019 leaf01 swp36 down 10G off SFP OEM SFP-H10GB-CU1M Thu Feb 7 22:49:37 2019 swp37 down 10G off SFP OEM SFPleaf01 H10GB-CU1M Thu Feb 7 22:49:37 2019 swp38 down 10G off SFP OEM SFPleaf01 H10GB-CU1M Thu Feb 7 22:49:37 2019 leaf01 swp39 down 10G off SFP Amphenol 571540007 Thu Feb 7 22:49:37 2019 leaf01 swp40 down 10G off SFP Amphenol

```
571540007
          Thu Feb 7 22:49:37 2019
        swp49 up
leaf01
                          40G off QSFP+ Amphenol
624410001 Thu Feb 7 22:49:37 2019
leaf01
        swp5 down 10G off SFP Amphenol
571540007 Thu Feb 7 22:49:37 2019
                    down 40G off QSFP+ Amphenol
leaf01
     swp50
624410001 Thu Feb 7 22:49:37 2019
leaf01
    swp51
              down 40G off QSFP+ Amphenol
603020003 Thu Feb 7 22:49:37 2019
leaf01 swp52 up 40G off QSFP+ Amphenol
603020003 Thu Feb 7 22:49:37 2019
leaf01
        swp54
             down 40G off QSFP+ Amphenol
624410002 Thu Feb 7 22:49:37 2019
```

View Components from a Given Vendor

By filtering for a specific cable vendor, you can collect information such as how many ports use components from that vendor and when they were last updated. This information may be useful when you run a cost analysis of your network. This example shows all the ports that are using components by an *OEM* vendor.

leaf01	swp36	down	10G	off	SFP	OEM	SFP-
H10GB-CU	J1M Thu Feb 7 22	:49:37 20	19				
leaf01	swp37	down	10G	off	SFP	OEM	SFP-
H10GB-CU	J1M Thu Feb 7 22	:49:37 20	19				
leaf01	swp38	down	10G	off	SFP	OEM	SFP-
H10GB-CU	J1M Thu Feb 7 22	:49:37 20	19				

View All Devices Using a Given Component

You can view all of the devices with ports using a particular component. This could be helpful when you need to change out a particular component for possible failure issues, upgrades, or cost reasons. This example first determines which models (part numbers) exist on all of the devices and then those devices with a part number of QSFP-H40G-CU1M installed.

cumulus@switch:~\$ netq show interfaces physical model

2231368-1 : 2231368-1

624400001 : 624400001

QSFP-H40G-CU1M : QSFP-H40G-CU1M

QSFP-H40G-CU1MUS: QSFP-H40G-CU1MUS

n/a : n/a

cumulus@switch:~\$ netq show interfaces physical model QSFP-H40G-CU1M

Matching cables records:

Hostname Interface State Speed AutoNeg Module

Vendor Part No Last Changed

leaf01 swp50 up 1G off QSFP+ OEM QSFP-

```
H40G-CU1M Thu Feb 7 18:31:20 2019

leaf02 swp52 up 1G off QSFP+ OEM QSFP-
H40G-CU1M Thu Feb 7 18:31:20 2019
```

View Changes to Physical Components

Because components are often changed, NetQ enables you to determine what, if any, changes have been made to the physical components on your devices. This can be helpful during deployments or upgrades.

You can select how far back in time you want to go, or select a time range using the between keyword. Note that time values must include units to be valid. If no changes are found, a "No matching cable records found" message is displayed. This example illustrates each of these scenarios for all devices in the network.

```
cumulus@switch:~$ netq show events type interfaces-physical between now and
30d
Matching cables records:
Hostname Interface
                                 Speed AutoNeg Module
                           State
Vendor
            Part No Last Changed
leaf01
          swp1
                       up
                              1G
                                    off SFP
                                              AVAGO
AFBR-5715PZ-JU1 Thu Feb 7 18:34:20 2019
                              10G
leaf01
          swp2
                                    off SFP
                                               OEM
                       up
SFP-10GB-LR Thu Feb 7 18:34:20 2019
leaf01
      swp47
                           10G off SFP
                                              JDSU
                        up
PLRXPLSCS4322N Thu Feb 7 18:34:20 2019
leaf01
          swp48
                              40G
                                     off QSFP+
                                                Mellanox
                        up
```

MC2210130-002 Thu Feb 7 18:34:20 2019

leaf01 swp49 down 10G off empty n/a n/

a Thu Feb 7 18:34:20 2019

leaf01 swp50 up 1G off SFP FINISAR CORP.

FCLF8522P2BTL Thu Feb 7 18:34:20 2019

leaf01 swp51 up 1G off SFP FINISAR CORP.

FTLF1318P3BTL Thu Feb 7 18:34:20 2019

leaf01 swp52 down 1G off SFP CISCO-AGILENT

QFBR-5766LP Thu Feb 7 18:34:20 2019

leaf02 swp1 up 1G on RJ45 n/a n/

a Thu Feb 7 18:34:20 2019

leaf02 swp2 up 10G off SFP Mellanox

MC2609130-003 Thu Feb 7 18:34:20 2019

leaf02 swp47 up 10G off QSFP+ CISCO

AFBR-7IER05Z-CS1 Thu Feb 7 18:34:20 2019

leaf02 swp48 up 10G off QSFP+ Mellanox

MC2609130-003 Thu Feb 7 18:34:20 2019

leaf02 swp49 up 10G off SFP FIBERSTORE

SFP-10GLR-31 Thu Feb 7 18:34:20 2019

leaf02 swp50 up 1G off SFP OEM SFP-GLC-

T Thu Feb 7 18:34:20 2019

leaf02 swp51 up 10G off SFP Mellanox

MC2609130-003 Thu Feb 7 18:34:20 2019

leaf02 swp52 up 1G off SFP FINISAR CORP.

FCLF8522P2BTL Thu Feb 7 18:34:20 2019

leaf03 swp1 up 10G off SFP Mellanox

MC2609130-003 Thu Feb 7 18:34:20 2019

leaf03 swp2 up 10G off SFP Mellanox

MC3309130-001 Thu Feb 7 18:34:20 2019

leaf03 swp47 10G off SFP CISCO-AVAGO up AFBR-7IER05Z-CS1 Thu Feb 7 18:34:20 2019 leaf03 swp48 up 10G off SFP Mellanox MC3309130-001 Thu Feb 7 18:34:20 2019 leaf03 swp49 down 1G off SFP FINISAR CORP. FCLF8520P2BTL Thu Feb 7 18:34:20 2019 leaf03 swp50 up 1G off SFP FINISAR CORP. FCLF8522P2BTL Thu Feb 7 18:34:20 2019 leaf03 swp51 up 10G off QSFP+ Mellanox MC2609130-003 Thu Feb 7 18:34:20 2019 oob-mgmt-server swp1 up 1G off RJ45 n/a n/ Thu Feb 7 18:34:20 2019 oob-mgmt-server swp2 up 1G off RJ45 n/a n/ Thu Feb 7 18:34:20 2019 a cumulus@switch:~\$ netq show events interfaces-physical between 6d and 16d Matching cables records: Hostname Interface State Speed AutoNeg Module Vendor Part No Last Changed leaf01 swp1 up 1G off SFP **AVAGO** AFBR-5715PZ-JU1 Thu Feb 7 18:34:20 2019 leaf01 swp2 up 10G off SFP OEM SFP-10GB-LR Thu Feb 7 18:34:20 2019 leaf01 swp47 up 10G off SFP JDSU PLRXPLSCS4322N Thu Feb 7 18:34:20 2019 leaf01 swp48 up 40G off QSFP+ Mellanox

```
MC2210130-002 Thu Feb 7 18:34:20 2019
leaf01
                                   10G
           swp49
                           down
                                           off empty
                                                                    n/
                                                        n/a
       Thu Feb 7 18:34:20 2019
а
leaf01
           swp50
                                  1G
                                        off
                                             SFP
                                                    FINISAR CORP.
                           up
FCLF8522P2BTL Thu Feb 7 18:34:20 2019
leaf01
           swp51
                                  1G
                                        off
                                             SFP
                                                    FINISAR CORP.
                           up
FTLF1318P3BTL Thu Feb 7 18:34:20 2019
leaf01
                           down
                                          off SFP
                                                      CISCO-AGILENT
           swp52
                                   1G
QFBR-5766LP
              Thu Feb 7 18:34:20 2019
```

cumulus@switch:~\$ netq show events type interfaces-physical between 0s and 5h No matching cables records found

Validate Physical Layer Configuration

Beyond knowing what physical components are deployed, it is valuable to know that they are configured and operating correctly. NetQ enables you to confirm that peer connections are present, discover any misconfigured ports, peers, or unsupported modules, and monitor for link flaps.

NetQ checks peer connections using LLDP. For DACs and AOCs, NetQ determines the peers using their serial numbers in the port EEPROMs, even if the link is not UP.

Confirm Peer Connections

You can validate peer connections for all devices in your network or for a specific device or port. This example shows the peer hosts and their status for leaf03 switch.

	ables records:					
Hostname	Interface	Peer I	Hostname	Peer Inter	face	State
Message						
		a a b ma a ma	t quitab au			
leaf03	swp1	oob-mgm	t-switch sv	vp7		
up						
leaf03	swp2			down	Peer po	rt
unknown						
leaf03	swp47	leaf04	swp47	L	ıp	
leaf03	swp48	leaf04	swp48	L	ір	
leaf03	swp49	leaf04	swp49	L	ір	
leaf03	swp50	leaf04	swp50	L	ір	
leaf03	swp51	exit01	swp51	U	ıp	

This example shows the peer data for a specific interface port.

	witch:~\$ netq le	af01 show inte	erfaces ph	ysical swp47	
Matching ca	ables records:				
Hostname	Interface	Peer H	ostname	Peer Interface	State
Message					
leaf01	swp47	leaf02	swp47	up	

Monitor Physical Layer Components

Validate Physical Layer Configuration

Discover Misconfigurations

You can verify that the following configurations are the same on both sides of a peer

interface:

Admin state

Operational state

Link speed

Auto-negotiation setting

The netq check interfaces command is used to determine if any of the interfaces have any continuity errors. This command only checks the physical interfaces; it does not check bridges, bonds or other software constructs. You can check all interfaces at once. It enables you to compare the current status of the interfaces, as well as their status at

an earlier point in time. The command syntax is:

netq check interfaces [around <text-time>] [json]

 (\checkmark) TIP

If NetQ cannot determine a peer for a given device, the port is marked as unverified.

If you find a misconfiguration, use the netq show interfaces physical command for clues about the cause.

Example: Find Mismatched Operational States

Monitor Physical Layer Components

Validate Physical Layer Configuration

In this example, we check all of the interfaces for misconfigurations and we find that one interface port has an error. We look for clues about the cause and see that the Operational states do not match on the connection between leaf 03 and leaf04: leaf03 is up, but leaf04 is down. If the misconfiguration was due to a mismatch in the administrative state, the message would have been *Admin state mismatch (up, down)* or *Admin state mismatch (down, up)*.

cumulus@s	switch:~\$ netq ch	neck interface	es		
	odes: 18, Failed N				
Checked Po	orts: 741, Failed F	Ports: 1, Unve	erified Ports	s: 414	
cumulus@s	switch:~\$ netq sh	now interface	es physical p	peer	
Matching c	ables records:				
Hostname	Interface	Peer l	Hostname	Peer Interface	Message
•••					
leaf03	swp1	oob-mgm	t-switch		
swp7					
leaf03	swp2			Peer port	
unknown					
leaf03	swp47	leaf04	swp47		
leaf03	swp48	leaf04	swp48	State mis	smatch (up,
down)					
leaf03	swp49	leaf04	swp49		
leaf03	swp50	leaf04	swp50		
leaf03	swp52			Port cage	
empty					
•••					

Example: Find Mismatched Peers

This example uses the *and* keyword to check the connections between two peers. An error is seen, so we check the physical peer information and discover that the incorrect

peer has been specified. After fixing it, we run the check again, and see that there are no longer any interface errors.

cumulus@switch:~\$ netq check interfaces

Checked Nodes: 1, Failed Nodes: 1

Checked Ports: 1, Failed Ports: 1, Unverified Ports: 0

cumulus@switch:~\$ netq show interfaces physical peer

Matching cables records:

Hostname Interface Peer Hostname Peer Interface Message

leaf01 swp50 leaf04 swp49 Incorrect peer

specified. Real peer

is leaf04 swp50

cumulus@switch:~\$ netq check interfaces

Checked Nodes: 1, Failed Nodes: 0

Checked Ports: 1, Failed Ports: 0, Unverified Ports: 0

Example: Find Mismatched Link Speeds

This example checks for for configuration mismatches and finds a link speed mismatch on server03. The link speed on swp49 is 40G and the peer port swp50 is unspecified.

cumulus@switch:~\$ netq check interfaces

Checked Nodes: 10, Failed Nodes: 1

Checked Ports: 125, Failed Ports: 2, Unverified Ports: 35

Hostname	Interface	Peer Hosti	name Peer l	Interface	Message
				-	
server03	swp49	server03	swp50	Speed	mismatch
(40G, Unkno	wn)				
server03	swp50	server03	swp49	Speed	mismatch
(Unknown, 4	0G)				

Example: Find Mismatched Auto-negotiation Settings

This example checks for configuration mismatches and finds auto-negotation setting mismatches between the servers and leafs. Auto-negotiation is *off* on the leafs, but *on* on the servers.

	odes: 15, Failed I orts: 118, Failed I		ified Ports	s· 9 <i>4</i>	
Hostname	Interface			Peer Interface	Message
leaf01	swp1	server01	eth1	Autoneg	mismatch
(off, on) leaf02 (off, on)	swp2	server02	eth2	Autoneg	mismatch
leaf03 (off, on)	swp1	server03	eth1	Autoneg	mismatch
leaf04	swp2	server04	eth2	Autoneg	

Monitor Physical Layer Components

Validate Physical Layer Configuration

server01	eth1	leaf01	swp1	Autoneg mismatch
(on, off)				
server02	eth2	leaf02	swp2	Autoneg mismatch
(on, off)				
server03	eth1	leaf03	swp1	Autoneg mismatch
(on, off)				
server04	eth2	leaf04	swp2	Autoneg mismatch
(on, off)				

Identify Flapping Links

You can also determine whether a link is flapping using the **netq check interfaces** command. If a link is flapping, NetQ indicates this in a message:

Monitor Data Link Layer Devices and Protocols

With NetQ, a network administrator can monitor OSI Layer 2 devices and protocols, including switches, bridges, link control, and physical media access. Keeping track of the various data link layer devices in your network ensures consistent and error-free communications between devices. NetQ provides the ability to:

- Monitor and validate device and protocol configurations
- View available communication paths between devices

It helps answer questions such as:

- Is a VLAN misconfigured?
- Is there an MTU mismatch in my network?
- Is MLAG configured correctly?
- Is there an STP loop?
- Can device A reach device B using MAC addresses?

Monitor LLDP Operation

LLDP is used by network devices for advertising their identity, capabilities, and neighbors on a LAN. You can view this information for one or more devices. You can also view the information at an earlier point in time or view changes that have occurred to the information during a specified timeframe. NetQ enables you to view LLDP information for your devices using the **netq show lldp** command. The syntax for this command is:

netq [<hostname>] show lldp [<remote-physical-interface>] [around <text-time>]
[json]

netq [<hostname>] show events [level info|level error|level warning|level critical| level debug] type lldp [between <text-time> and <text-endtime>] [json]

View LLDP Information for All Devices

This example shows the interface and peer information that is advertised for each device.

cumulus@switch:~\$ netq show lldp						
_	dp records: Interface	Peer Ho	stname	Peer Interface	Last	
exit01	swp1	edge01	swp5	Thu Feb	7 18:31:53	
2019 exit01 2019	swp2	edge02	swp5	Thu Feb	7 18:31:53	
	swp3	spine01	swp9	Thu Feb	7 18:31:53	
exit01 2019	swp4	spine02	swp9	Thu Feb	7 18:31:53	
	swp5	spine03	swp9	Thu Feb	7 18:31:53	
	swp6)19	firewall01	mac:00	0:02:00:00:00:11	Thu Feb 7	

exit01	swp7	firewall02	swp3	Thu Feb 7 18:31:53
2019				
exit02	swp1	edge01	swp6	Thu Feb 7 18:31:49
2019				
exit02	swp2	edge02	swp6	Thu Feb 7 18:31:49
2019				
exit02	swp3	spine01	swp10	Thu Feb 7 18:31:49
2019				
exit02	swp4	spine02	swp10	Thu Feb 7 18:31:49
2019				
exit02	swp5	spine03	swp10	Thu Feb 7 18:31:49
2019				
exit02	swp6	firewall01	mac:00:02:00:0	0:00:12 Thu Feb 7
18:31:49 20	019			
exit02	swp7	firewall02	swp4	Thu Feb 7 18:31:49
2019				
firewall01	swp1	edge01	swp14	Thu Feb 7
18:31:26 20	019			
firewall01	swp2	edge02	swp14	Thu Feb 7
18:31:26 20	019			
firewall01	swp3	exit01	swp6	Thu Feb 7 18:31:26
2019				
firewall01	swp4	exit02	swp6	Thu Feb 7 18:31:26
2019				
firewall02	swp1	edge01	swp15	Thu Feb 7
18:31:31 20	019			
firewall02	swp2	edge02	swp15	Thu Feb 7
18:31:31 20)19			
firewall02	swp3	exit01	swp7	Thu Feb 7 18:31:31

2019				
firewall02	swp4	exit02	swp7	Thu Feb 7 18:31:31
2019				
server11	swp1	leaf01	swp7	Thu Feb 7 18:31:43
2019				
server11	swp2	leaf02	swp7	Thu Feb 7 18:31:43
2019				
server11	swp3	edge01	swp16	Thu Feb 7
18:31:43 201	19			
server11	swp4	edge02	swp16	Thu Feb 7
18:31:43 201	19			
server12	swp1	leaf01	swp8	Thu Feb 7 18:31:47
2019				
server12	swp2	leaf02	swp8	Thu Feb 7 18:31:47
2019				

Monitor Interface Health

Interface (link) health can be monitored using the netq show

interfaces command. You can view status of the links, whether they are operating over a VRF interface, the MTU of the link, and so forth. Using the hostname option enables you to view only the interfaces for a given device. View changes to interfaces using the netq show events command.

The syntax for these commands is:

netq [<hostname>] show interfaces [type bond|type bridge|type eth|type loopback|type macvlan|type swp|type vlan|type vrf|type vxlan] [state <remote-interface-state>] [around <text-time>] [json]

netq <hostname> show interfaces [type bond|type bridge|type eth|type loopback|type macvlan|type swp|type vlan|type vrf|type vxlan] [state <remote-interface-state>] [around <text-time>] [count] [json] netq [<hostname>] show events [level info | level error | level warning | level critical | level debug] type interfaces [between <text-time> and <text-endtime>] [json]

View Status for All Interfaces

Viewing the status of all interfaces at once can be helpful when you are trying to compare configuration or status of a set of links, or generally when changes have been made.

This example shows all interfaces network-wide.

```
cumulus@switch:~$ netq show interfaces
Matching link records:
Hostname
            Interface
                           Type
                                     State VRF
                                                     Details
    Last Changed
                                  up default
exit01
          bridge
                        bridge
                                                 , Root bridge:
exit01, Mon Apr 29 20:57:59 2019
                                       Root port:, Members: vxlan4001,
                                       bridge,
exit01 eth0
                       eth
                                       mgmt
                                                 MTU: 1500
                                up
  Mon Apr 29 20:57:59 2019
                                 up default
          lo
                                                 MTU: 65536
exit01
                      loopback
    Mon Apr 29 20:57:58 2019
```

exit01 mgmt vrf up table: 1001, MTU:

65536, Mon Apr 29 20:57:58 2019

Members: mgmt, eth0,

exit01 swp1 swp down default VLANs:, PVID: 0

MTU: 1500 Mon Apr 29 20:57:59 2019

exit01 swp44 swp up vrf1 VLANs:,

Mon Apr 29 20:57:58 2019

PVID: 0 MTU: 1500 LLDP:

internet:sw

p1

exit01 swp45 swp down default VLANs: , PVID: 0

MTU: 1500 Mon Apr 29 20:57:59 2019

exit01 swp46 swp down default VLANs:, PVID: 0

MTU: 1500 Mon Apr 29 20:57:59 2019

exit01 swp47 swp down default VLANs:, PVID: 0

MTU: 1500 Mon Apr 29 20:57:59 2019

• • •

leaf01 bond01 bond up default Slave:swp1 LLDP:

server01:eth1 Mon Apr 29 20:57:59 2019

leaf01 bond02 bond up default Slave:swp2 LLDP:

server02:eth1 Mon Apr 29 20:57:59 2019

leaf01 bridge bridge up default , Root bridge:

leaf01, Mon Apr 29 20:57:59 2019

Root port:, Members: vxlan4001,

bond02, vni24, vni13, bond01,

bridge, peerlink,

leaf01 eth0 eth up mgmt MTU:

1500 Mon Apr 29 20:58:00 2019 loopback up default leaf01 lo MTU: 65536 Mon Apr 29 20:57:59 2019 vrf up leaf01 mgmt table: 1001, MTU: Mon Apr 29 20:57:59 2019 65536, Members: mgmt, eth0, leaf01 peerlink bond up default Slave:swp50 LLDP: leaf02:swp49 LLDP Mon Apr 29 20:58:00 2019 : leaf02:swp50

View Interface Status for a Given Device

If you are interested in only a the interfaces on a specific device, you can view only those.

This example shows all interfaces on the *spine01* device.

cumulus@switch:~\$ netq spine01 show interfaces Matching link records: Interface Type State VRF Hostname Details Last Changed spine01 eth0 eth up mgmt MTU: 1500 Mon Apr 29 21:12:47 2019 spine01 lo loopback up default MTU: 65536 Mon Apr 29 21:12:47 2019 table: 1001, MTU: spine01 mgmt vrf up

65536, Mon Apr 29 21:12:46 2019

Members: mgmt, eth0,

spine01 swp1 swp up default VLANs:,

Mon Apr 29 21:12:47 2019

PVID: 0 MTU: 9216 LLDP:

leaf01:swp5

1

spine01 swp2 swp up default VLANs:,

Mon Apr 29 21:12:47 2019

PVID: 0 MTU: 9216 LLDP:

leaf02:swp5

1

spine01 swp29 swp up default VLANs:,

Mon Apr 29 21:12:47 2019

PVID: 0 MTU: 9216 LLDP:

exit02:swp5

1

spine01 swp3 swp up default VLANs:,

Mon Apr 29 21:12:46 2019

PVID: 0 MTU: 9216 LLDP:

leaf03:swp5

1

spine01 swp30 swp up default VLANs:,

Mon Apr 29 21:12:47 2019

PVID: 0 MTU: 9216 LLDP:

exit01:swp5

1

spine01 swp31 swp up default VLANs:,

Mon Apr 29 21:12:46 2019

PVID: 0 MTU: 9216 LLDP: spine02:swp 31 spine01 swp32 up default VLANs:, swp Mon Apr 29 21:12:46 2019 PVID: 0 MTU: 9216 LLDP: spine02:swp 32 spine01 swp4 swp up default VLANs:, Mon Apr 29 21:12:47 2019 PVID: 0 MTU: 9216 LLDP: leaf04:swp5 1

View All Interfaces of a Given Type

It can be can be useful to see the status of a particular type of interface.

This example shows all bond interfaces that are down, and then those that are up.

leaf01 bond01 up default Slave:swp1 LLDP:

server01:eth1 Mon Apr 29 21:19:07 2019

leaf01 bond02 bond up default Slave:swp2 LLDP:

server02:eth1 Mon Apr 29 21:19:07 2019

leaf01 peerlink bond up default Slave:swp50 LLDP:

leaf02:swp49 LLDP Mon Apr 29 21:19:07 2019

: leaf02:swp50

leaf02 bond01 bond up default Slave:swp1 LLDP:

server01:eth2 Mon Apr 29 21:19:07 2019

leaf02 bond02 bond up default Slave:swp2 LLDP:

server02:eth2 Mon Apr 29 21:19:07 2019

leaf02 peerlink bond up default Slave:swp50 LLDP:

leaf01:swp49 LLDP Mon Apr 29 21:19:07 2019

: leaf01:swp50

leaf03 bond03 bond up default Slave:swp1 LLDP:

server03:eth1 Mon Apr 29 21:19:07 2019

leaf03 bond04 bond up default Slave:swp2 LLDP:

server04:eth1 Mon Apr 29 21:19:07 2019

leaf03 peerlink bond up default Slave:swp50 LLDP:

leaf04:swp49 LLDP Mon Apr 29 21:19:07 2019

: leaf04:swp50

leaf04 bond03 bond up default Slave:swp1 LLDP:

server03:eth2 Mon Apr 29 21:19:07 2019

leaf04 bond04 bond up default Slave:swp2 LLDP:

server04:eth2 Mon Apr 29 21:19:07 2019

leaf04 peerlink bond up default Slave:swp50 LLDP:

leaf03:swp49 LLDP Mon Apr 29 21:19:07 2019

: leaf03:swp50

server01 bond0 bond up default Slave:bond0

LLDP: lea	af02:swp1	Mon Apr 29 21:19:07	2019		
server02	bond0	bond	up	default	Slave:bond0
LLDP: lea	af02:swp2	Mon Apr 29 21:19:07	2019		
server03	bond0	bond	up	default	Slave:bond0
LLDP: lea	af04:swp1	Mon Apr 29 21:19:07	2019		
server04	bond0	bond	up	default	Slave:bond0
LLDP: lea	af04:swp2	Mon Apr 29 21:19:07	2019		

View the Total Number of Interfaces

For a quick view of the amount of interfaces currently operating on a device, use the hostname and count options together.

This example shows the count of interfaces on the *leaf03* switch.

cumulus@switch:~\$ netq leaf03 show interfaces count Count of matching link records: 28

View the Total Number of a Given Interface Type

It can be useful to see how many interfaces of a particular type you have on a device.

This example shows the count of swp interfaces are on the *leaf03* switch.

cumulus@switch:~\$ netq leaf03 show interfaces type swp count

Count of matching link records: 11

Monitor Data Link Layer Devices and Protocols

Monitor Interface Health

View Changes to Interfaces

If you suspect that an interface is not working as expected, seeing a drop in performance or a large number of dropped messages for example, you can view changes that have been made to interfaces network-wide.

This example shows info level events for all interfaces in your network:

30d			
Matching ev	ents record	s:	
Hostname	Messag	е Туре	Severity Message
Timestamp			
server03	link	info	HostName server03 changed state fro 3d:
12h:8m:28s			
		m d	lown to up Interface:eth2
server03	link	info	HostName server03 changed state fro 3d:
12h:8m:28s			
		m d	lown to up Interface:eth1
server01	link	info	HostName server01 changed state fro 3d:
12h:8m:30s			
		m d	lown to up Interface:eth2
server01	link	info	HostName server01 changed state fro 3d:
12h:8m:30s			
		m d	lown to up Interface:eth1
server02	link	info	HostName server02 changed state fro 3d:
12h:8m:34s			
		m d	lown to up Interface:eth2

Check for MTU Inconsistencies

The maximum transmission unit (MTU) determines the largest size packet or frame that can be transmitted across a given communication link. When the MTU is not configured to the same value on both ends of the link, communication problems can occur. With

NetQ, you can verify that the MTU is correctly specified for each link using the netq check mtu command.

This example shows that four switches have inconsistently specified link MTUs. Now the network administrator or operator can reconfigure the switches and eliminate the communication issues associated with this misconfiguration.

cumulus@s	witch:~\$ netq ch	eck mtu			
Checked No	odes: 15, Checked	d Links: 215, Failed N	odes: 4, Fa	iled Links: 7	
MTU misma	atch found on fol	lowing links			
Hostname	Interface	MTU Peer	Peer l	interface P	eer
MTU Error					
spine01	swp30	 9216 exit01	swp51	1500	MTU
Mismatch			'		
exit01	swp51	1500 spine01	swp30	9216	MTU
Mismatch					
spine01	swp29	9216 exit02	swp51	1500	MTU
Mismatch					
exit02			-	Rotten Agent	
exit01	swp52	1500 spine02	swp30	9216	MTU
Mismatch					
spine02	swp30	9216 exit01	swp52	1500	MTU
Mismatch					
spine02	swp29	9216 exit02	swp52	1500	MTU
Mismatch					

Monitor VLAN Configurations

A VLAN (Virtual Local Area Network) enables devices on one or more LANs to communicate as if they were on the same network, without being physically connected. The VLAN enables network administrators to partition a network for functional or security requirements without changing physical infrastructure. With NetQ, you can view the operation of VLANs for one or all devices. You can also view the information at an earlier point in time or view changes that have occurred to the information during a specified timeframe. NetQ enables you to view basic VLAN information for your devices using the netq show vlan command. Additional show commands enable you to view VLAN information associated with interfaces and MAC addresses. The syntax for these commands is:

netq [<hostname>] show interfaces [type vlan] [state <remote-interface-state>]
[around <text-time>] [json]
netq <hostname> show interfaces [type vlan] [state <remote-interface-state>]
[around <text-time>] [count] [json]
netq [<hostname>] show events [level info | level error | level warning | level
critical | level debug] type vlan [between <text-time> and <text-endtime>] [json]
netq show macs [<mac>] [vlan <1-4096>] [origin] [around <text-time>] [json]
netq <hostname> show macs [<mac>] [vlan <1-4096>] [origin | count] [around
<text-time>] [json]
netq <hostname> show macs egress-port <egress-port> [<mac>] [vlan <1-4096>]
[origin] [around <text-time>] [json]
netq [<hostname>] show vlan [<1-4096>] [around <text-time>] [json]



(i) NOTE

When entering a time value, you must include a numeric value and the unit of measure:

- d: day(s)
- h: hour(s)
- m: minute(s)
- s: second(s)
- now

For time ranges, the <text-time> is the most recent time and the <textendtime> is the oldest time. The values do not have to have the same unit of measure.

View VLAN Information for All Devices

This example shows the VLANs configured across your network.

umulus@s	switch:~\$ netq sh	now vlan	
Matching v	lan records:		
Hostname	VLANs	SVIs	Last Changed
exit01	4001	4001	Thu Feb 7 18:31:38 2019
exit02	4001	4001	Thu Feb 7 18:31:38 2019
leaf01	1,13,24,4001	13 24 4001	Thu Feb 7 18:31:38 2019
leaf02	1,13,24,4001	13 24 4001	Thu Feb 7 18:31:38 2019

Monitor Data Link Layer Devices and Protocols

Monitor VLAN Configurations

leaf03	1,13,24,4001	13 24 4001	Thu Feb 7 18:31:38 2019	
leaf04	1,13,24,4001	13 24 4001	Thu Feb 7 18:31:38 2019	

View VLAN Interface Information

You can view the current or past state of the interfaces associated with VLANs using the **netq show interfaces** command. This provides the status of the interface, its specified MTU, whether it is running over a VRF, and the last time it was changed.

9	link records:				
lostname	e Interface	Туре	St	ate VRF	Details
Last	Changed				
					-
xit01	vlan4001	vlan	up	vrf1	MTU:1500
Fri Fe	b 8 00:24:28 2019				
xit02	vlan4001	vlan	up	vrf1	MTU:1500
Fri Fe	b 8 00:24:28 2019				
eaf01	peerlink.4094	vlan	up	defau	lt MTU:9000
Fri	Feb 8 00:24:28 20	19			
eaf01	vlan13	vlan	up	vrf1	MTU:1500
Fri Feb	8 00:24:28 2019				
eaf01	vlan24	vlan	up	vrf1	MTU:1500
Fri Feb	8 00:24:28 2019				
eaf01	vlan4001	vlan	up	vrf1	MTU:1500
Fri Fe	b 8 00:24:28 2019				
eaf02	peerlink.4094	vlan	up	defau	lt MTU:9000
Fri	Feb 8 00:24:28 20	19			

leaf02	vlan13	vlan	up	vrf1	MTU:1500
Fri Feb	8 00:24:28 2019				
leaf02	vlan24	vlan	up	vrf1	MTU:1500
Fri Feb	8 00:24:28 2019				
leaf02	vlan4001	vlan	up	vrf1	MTU:1500
Fri Fe	b 8 00:24:28 2019				
leaf03	peerlink.4094	vlan	up	defa	ult MTU:9000
Fri	i Feb 8 00:24:28 20)19			
leaf03	vlan13	vlan	up	vrf1	MTU:1500
Fri Feb	8 00:24:28 2019				
leaf03	vlan24	vlan	up	vrf1	MTU:1500
Fri Feb	8 00:24:28 2019				
leaf03	vlan4001	vlan	up	vrf1	MTU:1500
Fri Fe	b 8 00:24:28 2019				
leaf04	peerlink.4094	vlan	up	defa	ult MTU:9000
Fri	i Feb 8 00:24:28 20)19			
leaf04	vlan13	vlan	up	vrf1	MTU:1500
Fri Feb	8 00:24:28 2019				
leaf04	vlan24	vlan	up	vrf1	MTU:1500
Fri Feb	8 00:24:28 2019				
leaf04	vlan4001	vlan	up	vrf1	MTU:1500
Fri Fe	b 8 00:24:28 2019				

View MAC Addresses Associated with a VLAN

You can determine the MAC addresses associated with a given VLAN using the netq show macs vlan command. The command also provides the hostnames of the devices, the egress port for the interface, whether the MAC address originated from the given

device, whether it learns the MAC address from the peer (remote=yes), and the last time the configuration was changed.

This example shows the MAC addresses associated with *VLAN13*.

	ching mac records:				
Orig	jin MAC Address	VLA	N Hostnam	e Egress Port	Remote Last
Cha	nged				
no		13	leaf01	bond01:server01 n	o Fri Feb 8
00:2	24:28 2019				
no	00:03:00:11:11:01	13	leaf02	bond01:server01 n	o Fri Feb 8
00:2	4:28 2019				
no	00:03:00:11:11:01	13	leaf03	vni13:leaf01 yes	Fri Feb 8
00:2	4:28 2019				
no	00:03:00:11:11:01	13	leaf04	vni13:leaf01 yes	Fri Feb 8
00:2	4:28 2019				
no	00:03:00:33:33:01	13	leaf01	vni13:10.0.0.134 yes	Fri Feb 8
00:2	4:28 2019				
no	00:03:00:33:33:01	13	leaf02	vni13:10.0.0.134 yes	Fri Feb 8
00:2	4:28 2019				
no	00:03:00:33:33:01	13	leaf03	bond03:server03 no	o Fri Feb 8
00:2	4:28 2019				
no	00:03:00:33:33:01	13	leaf04	bond03:server03 n	o Fri Feb 8
00:2	4:28 2019				
no	02:03:00:11:11:01	13	leaf01	bond01:server01 ne	o Fri Feb 8
00:2	4:28 2019				
	02:03:00:11:11:01	12	loof02	bond01:server01 ne	o Fri Feb 8

00:24:28 2019	
no 02:03:00:11:11:01 13 leaf03	vni13:leaf01 yes Fri Feb 8
00:24:28 2019	
no 02:03:00:11:11:01 13 leaf04	vni13:leaf01 yes Fri Feb 8
00:24:28 2019	
no 02:03:00:11:11:02 13 leaf01	bond01:server01 no Fri Feb 8
00:24:28 2019	
no 02:03:00:11:11:02 13 leaf02	bond01:server01 no Fri Feb 8
00:24:28 2019	
no 02:03:00:11:11:02 13 leaf03	vni13:leaf01 yes Fri Feb 8
00:24:28 2019	
no 02:03:00:11:11:02 13 leaf04	vni13:leaf01 yes Fri Feb 8
00:24:28 2019	
no 02:03:00:33:33:01 13 leaf01	vni13:10.0.0.134 yes Fri Feb 8
00:24:28 2019	
no 02:03:00:33:33:01 13 leaf02	vni13:10.0.0.134 yes Fri Feb 8
00:24:28 2019	
no 02:03:00:33:33:01 13 leaf03	bond03:server03 no Fri Feb 8
00:24:28 2019	
no 02:03:00:33:33:01 13 leaf04	bond03:server03 no Fri Feb 8
00:24:28 2019	
no 02:03:00:33:33:02 13 leaf01	vni13:10.0.0.134 yes Fri Feb 8
00:24:28 2019	
no 02:03:00:33:33:02 13 leaf02	vni13:10.0.0.134 yes Fri Feb 8
00:24:28 2019	
no 02:03:00:33:33:02 13 leaf03	bond03:server03 no Fri Feb 8
00:24:28 2019	
no 02:03:00:33:33:02 13 leaf04	bond03:server03 no Fri Feb 8
00:24:28 2019	

yes 44:38:39:00:00:03 13	leaf01	bridge	no	Fri Feb 8 00:24:28
2019				
yes 44:38:39:00:00:15 13	leaf02	bridge	no	Fri Feb 8 00:24:28
2019				
yes 44:38:39:00:00:23 13	leaf03	bridge	no	Fri Feb 8 00:24:28
2019				
yes 44:38:39:00:00:5c 13	leaf04	bridge	no	Fri Feb 8 00:24:28
2019				
yes 44:39:39:ff:00:13 13	leaf01	bridge	no	Fri Feb 8 00:24:28
2019				
yes 44:39:39:ff:00:13 13	leaf02	bridge	no	Fri Feb 8 00:24:28
2019				
yes 44:39:39:ff:00:13 13	leaf03	bridge	no	Fri Feb 8 00:24:28
2019				
yes 44:39:39:ff:00:13 13	leaf04	bridge	no	Fri Feb 8 00:24:28
2019				

View MAC Addresses Associated with an Egress Port

You can filter that information down to just the MAC addresses that are associated with a given VLAN that use a particular egress port. This example shows MAC addresses associated with the *leaf03* switch and *VLAN 13* that use the *bridge* port.

```
cumulus@switch:~$ netq leaf03 show macs egress-port bridge vlan 13

Matching mac records:

Origin MAC Address VLAN Hostname Egress Port Remote Last

Changed

-----
yes 44:38:39:00:00:23 13 leaf03 bridge no Fri Feb 8 00:24:28
```

2019 yes 44:39:39:ff:00:13 13 leaf03 bridge no Fri Feb 8 00:24:28 2019

View the MAC Addresses Associated with VRR Configurations

You can view all of the MAC addresses associated with your VRR (virtual router reflector) interface configuration using the netq show

interfaces type macvlan command. This is useful for determining if the specified MAC address inside a VLAN is the same or different across your VRR configuration.

Matching li	nk records:				
Hostname	Interface	Туре	State	VRF	Details
Last C	hanged				
leaf01	vlan13-v0	macvlan	up	vrf1	MAC: 44:39:39:ff:
00:13,	Fri Feb 8 00:28	:09 2019			
			Мо	de: Privat	e
leaf01	vlan24-v0	macvlan	up	vrf1	MAC: 44:39:39:ff:
00:24,	Fri Feb 8 00:28	:09 2019			
			Mod	de: Privat	е
leaf02	vlan13-v0	macvlan	up	vrf1	MAC: 44:39:39:ff:
00:13,	Fri Feb 8 00:28	:09 2019			
			Mod	de: Privat	е
leaf02	vlan24-v0	macvlan	up	vrf1	MAC: 44:39:39:ff:
00:24,	Fri Feb 8 00:28	:09 2019			
			Мо	de: Privat	е

leaf03	vlan13-v0	macvlan	up	vrf1	MAC: 44:39:39:ff:
00:13,	Fri Feb 8 00:28	8:09 2019			
			M	ode: Priva	te
leaf03	vlan24-v0	macvlan	up	vrf1	MAC: 44:39:39:ff:
00:24,	Fri Feb 8 00:28	8:09 2019			
			M	ode: Priva	te
leaf04	vlan13-v0	macvlan	up	vrf1	MAC: 44:39:39:ff:
00:13,	Fri Feb 8 00:28	8:09 2019			
			M	ode: Priva	te
leaf04	vlan24-v0	macvlan	up	vrf1	MAC: 44:39:39:ff:
00:24,	Fri Feb 8 00:28	8:09 2019			
			M	ode: Priva	te

Monitor MLAG Configurations

Multi-Chassis Link Aggregation (MLAG) is used to enable a server or switch with a two-port bond (such as a link aggregation group/LAG, EtherChannel, port group or trunk) to connect those ports to different switches and operate as if they are connected to a single, logical switch. This provides greater redundancy and greater system throughput. Dual-connected devices can create LACP bonds that contain links to each physical switch. Therefore, active-active links from the dual-connected devices are supported even though they are connected to two different physical switches.



TIP

MLAG or CLAG? The Cumulus Linux implementation of MLAG is referred to by other vendors as CLAG, MC-LAG or VPC. You will even see references to CLAG in Cumulus Linux and Cumulus NetQ, including the management daemon, named clagd, and other options in the code, such as clag-id, which exist for historical purposes. The Cumulus Linux implementation is truly a multi-chassis link aggregation protocol, so we call it MLAG.

For instructions on configuring MLAG, refer to the MLAG topic in the Cumulus Linux User Guide.

With NetQ, you can view the configuration and operation of devices using MLAG using the **netq show clag** command. You can view the current configuration and the configuration at a prior point in time, as well as view any changes that have been made within a timeframe. The syntax for the show command is:

netq [<hostname>] show clag [around <text-time>] [json]
netq [<hostname>] show events [level info|level error|level warning|level critical|
level debug] type clag [between <text-time> and <text-endtime>] [json]

View MLAG Configuration and Status for all Devices

This example shows the configuration and status of MLAG for all devices. In this case, three MLAG pairs are seen between leaf11 and leaf12 (which happens to be down), edge01(P) and edge02, and leaf21(P) and leaf22.

Matching c	lag records:								
Hostname	Peer	SysMac	State	Bac	kup #	Bor	nd #D	ual Last	
Changed									
			S						
 leaf11	44	:38:39:ff:ff:01	down	n/a (0	Tł	nu Fe	b 7 18:30):49
2019									
leaf12	44	:38:39:ff:ff:01	down	down	8 (0	Thu F	eb 7	
18:30:53 20	19								
edge01(P)	edge02	00:01:01:1	0:00:01 u	up	up	25	25	Thu Feb	7
18:31:02 20	19								
edge02	edge01(P)	00:01:01:1	0:00:01 ι	up	up	25	25	Thu Feb	7
18:31:15 20	19								
leaf21(P)	leaf22	44:38:39:ff:ff:	02 up	up	8	8	Thu F	eb 7	
18:31:20 20	19								
leaf22	leaf21(P)	44:38:39:ff:ff:	02 un	un	8	8	Thu F	eb 7	

You can go back in time to see when this first MLAG pair went down. These results indicate that the pair became disconnected some time in the last five minutes.

```
cumulus@switch:~$ netq show clag around 5m

Matching clag records:

Hostname Peer SysMac State Backup #Bond #Dual Last

Changed
```

edge01(P)	edge02	00:01:01:10:00:01	up	up	25	25	Thu Feb 7
18:31:30 20°	19						
edge02	edge01(P)	00:01:01:10:00:01	up	up	25	25	Thu Feb 7
18:31:30 20°	19						
leaf11(P)	leaf12	44:38:39:ff:ff:01 up	up	8	8	Thu	Feb 7
18:31:30 20°	19						
leaf12	leaf11(P)	44:38:39:ff:ff:01 up	up	8	8	Thu	Feb 7
18:31:30 20 ⁻	19						
leaf21(P)	leaf22	44:38:39:ff:ff:02 up	up	8	8	Thu	Feb 7
18:31:30 20 ⁻	19						
leaf22	leaf21(P)	44:38:39:ff:ff:02 up	up	8	8	Thu	Feb 7
18:31:30 20 ⁻	19						

View MLAG Configuration and Status for Given Devices

This example shows that leaf22 is up and MLAG properly configured with a peer connection to leaf21 through 8 bonds, all of which are dual bonded.

cumulus@s	witch:~\$ neto	ղ leaf22 show cla	g					
Matching C	LAG session	records are:						
Hostname	Peer	SysMac	State	Back	kup	#Bc	ond #Dual Last	
Changed								
			S					
leaf22	leaf21(P)	44:38:39:ff:ff:0	2 up	up	8	8	Thu Feb 7	
18:31:30 20	19							

When you're directly on the switch, you can run clagctl to get the state:

```
cumulus@switch:~$ sudo clagctl
The peer is alive
Peer Priority, ID, and Role: 4096 00:02:00:00:00:4e primary
Our Priority, ID, and Role: 8192 44:38:39:00:a5:38 secondary
Peer Interface and IP: peerlink-3.4094 169.254.0.9
VxLAN Anycast IP: 36.0.0.20
Backup IP: 27.0.0.20 (active)
System MAC: 44:38:39:ff:ff:01
CLAG Interfaces
Our Interface Peer Interface CLAG Id Conflicts Proto-Down Reason
vx-38 - -
vx-33 - -
hostbond4 hostbond4 1 -
hostbond5 hostbond5 2 -
vx-37 vx-37 - -
vx-36 vx-36 -
vx-35 vx-35 -
vx-34 vx-34
```

Monitor Time Synchronization Status for Devices

It is important that the switches and hosts remain in time synchronization with the NetQ Platform to ensure collected data is properly captured and processed. You can use the **netq show ntp** command to view the time synchronization status for all devices or filter for devices that are either in synchronization or out of synchronization, currently or at a time in the past. The syntax for the show command is:

netq [<hostname>] show ntp [out-of-sync|in-sync] [around <text-time>] [json]
netq [<hostname>] show events [level info|level error|level warning|level critical|
level debug] type ntp [between <text-time> and <text-endtime>] [json]

This example shows the time synchronization status for all devices in the network.

```
cumulus@switch:~$ netg show ntp
Matching ntp records:
Hostname
             NTP Sync Current Server Stratum NTP App
edge01
          yes services01.it.c 3 ntpg
exit01
          yes time.tritn.com 2
                                   ntpq
exit02
                time.tritn.com 2
          yes
                                   ntpq
                               ntpq
internet
                          16
          no
          yes services01.it.c 2
leaf01
                                  ntpq
leaf02
               services01.it.c 2
                                  ntpq
          yes
leaf03
                107.181.191.189 2
          yes
                                     ntpq
leaf04
                grom.polpo.org 2
           yes
                                     ntpq
oob-mgmt-server yes
                      linode227395.st 2
                                          ntpq
                 192.168.0.254 3
server01
            yes
                                     ntpq
                 192.168.0.254 3
server02
            yes
                                     ntpq
                 192.168.0.254 3
server03
            yes
                                     ntpq
```

```
      server04
      yes
      192.168.0.254
      3
      ntpq

      spine01
      yes
      107.181.191.189
      2
      ntpq

      spine02
      yes
      t2.time.bf1.yah
      2
      ntpq
```

This example shows all devices in the network that are out of time synchronization, and consequently might need to be investigated.

```
cumulus@switch:~$ netq show ntp out-of-sync

Matching ntp records:

Hostname NTP Sync Current Server Stratum NTP App

------
internet no - 16 ntpq
```

This example shows the time synchronization status for *leaf01*.

```
cumulus@switch:~$ netq leaf01 show ntp

Matching ntp records:

Hostname NTP Sync Current Server Stratum NTP App

------
leaf01 yes kilimanjaro 2 ntpq
```

Monitor Spanning Tree Protocol Configuration

The Spanning Tree Protocol (STP) is used in Ethernet-based networks to prevent communication loops when you have redundant paths on a bridge or switch. Loops cause excessive broadcast messages greatly impacting the network performance. With NetQ, you can view the STP topology on a bridge or switch to ensure no loops have been created using the netq show stp topology command. You can also view the topology information for a prior point in time to see if any changes were made around then. The syntax for the show command is:

netq <hostname> show stp topology [around <text-time>] [json]

This example shows the STP topology as viewed from the *spine1* switch.

```
cumulus@switch:~$ netq spine1 show stp topology

Root(spine1) -- spine1:sw_clag200 -- leaf2:EdgeIntf(sng_hst2) -- hsleaf21
-- leaf2:EdgeIntf(dual_host2) -- hdleaf2
-- leaf2:EdgeIntf(dual_host1) -- hdleaf1
-- leaf2:ClagIsl(peer-bond1) -- leaf1
-- leaf1:EdgeIntf(sng_hst2) -- hsleaf11
-- leaf1:EdgeIntf(dual_host2) -- hdleaf2
-- leaf1:EdgeIntf(dual_host1) -- hdleaf1
-- leaf1:ClagIsl(peer-bond1) -- leaf2
-- spine1:ClagIsl(peer-bond1) -- spine2
-- spine1:sw_clag300 -- edge1:EdgeIntf(sng_hst2) -- hsedge11
-- edge1:EdgeIntf(dual_host1) -- hdedge2
-- edge1:EdgeIntf(dual_host1) -- hdedge1
-- edge1:ClagIsl(peer-bond1) -- edge2
```

```
-- edge2:EdgeIntf(sng_hst2) -- hsedge21
                  -- edge2:EdgeIntf(dual_host2) -- hdedge2
                   -- edge2:EdgeIntf(dual_host1) -- hdedge1
                  -- edge2:ClagIsl(peer-bond1) -- edge1
Root(spine2) -- spine2:sw_clag200 -- leaf2:EdgeIntf(sng_hst2) -- hsleaf21
                   -- leaf2:EdgeIntf(dual_host2) -- hdleaf2
                  -- leaf2:EdgeIntf(dual_host1) -- hdleaf1
                   -- leaf2:ClagIsl(peer-bond1) -- leaf1
                  -- leaf1:EdgeIntf(sng_hst2) -- hsleaf11
                  -- leaf1:EdgeIntf(dual_host2) -- hdleaf2
                  -- leaf1:EdgeIntf(dual_host1) -- hdleaf1
                  -- leaf1:ClagIsl(peer-bond1) -- leaf2
       -- spine2:ClagIsl(peer-bond1) -- spine1
       -- spine2:sw_clag300 -- edge2:EdgeIntf(sng_hst2) -- hsedge21
                  -- edge2:EdgeIntf(dual_host2) -- hdedge2
                   -- edge2:EdgeIntf(dual_host1) -- hdedge1
                   -- edge2:ClagIsl(peer-bond1) -- edge1
                  -- edge1:EdgeIntf(sng_hst2) -- hsedge11
                   -- edge1:EdgeIntf(dual_host2) -- hdedge2
                   -- edge1:EdgeIntf(dual_host1) -- hdedge1
                   -- edge1:ClagIsl(peer-bond1) -- edge2
```

Validate Paths between Devices

If you have VLANs configured, you can view the available paths between two devices on the VLAN currently and at a time in the past using their MAC addresses. You can view the output in one of three formats (*json*, *pretty*, and *detail*). JSON output provides the output in a JSON file format for ease of importing to other applications or software. Pretty output lines up the paths in a pseudo-graphical manner to help visualize multiple

paths. Detail output is useful for traces with higher hop counts where the pretty output wraps lines, making it harder to interpret the results. The detail output displays a table with a row for each path.

To view the paths:

- 1. Identify the MAC address and VLAN ID for the destination device
- 2. Identify the IP address or hostname for the source device
- 3. Use the **netq trace** command to see the available paths between those devices.

The trace command syntax is:

netq trace <mac> [vlan <1-4096>] from (<src-hostname>|<ip-src>) [vrf <vrf>] [around <text-time>] [json|detail|pretty] [debug]



The syntax requires the destination device address first, mac, and then the source device address or hostname. Additionally, the vlan keyword-value pair is required for layer 2 traces even though the syntax indicates it is optional.

The tracing function only knows about addresses that have already been learned. If you find that a path is invalid or incomplete, you may need to ping the identified device so that its address becomes known.

View Paths between Two Switches with Pretty Output

This example shows the available paths between a top of rack switch, *tor-1*, and a server, *server11*. The request is to go through VLAN *1001* from the VRF *vrf1*. The results

include a summary of the trace, including the total number of paths available, those with errors and warnings, and the MTU of the paths. In this case, the results are displayed in pseudo-graphical output.

```
cumulus@switch:~$ netq trace 00:02:00:00:00:02 vlan 1001 from leaf01 vrf vrf1
pretty

Number of Paths: 4

Number of Paths with Errors: 0

Number of Paths with Warnings: 0

Path MTU: 9152

leaf01 vni: 34 uplink-2 -- downlink-5 spine02 downlink-2 -- uplink-2 vni: 34 leaf12
hostbond4 -- swp2 server11

uplink-2 -- downlink-5 spine02 downlink-1 -- uplink-2 vni: 34 leaf11
hostbond4 -- swp1 server11

leaf01 vni: 34 uplink-1 -- downlink-5 spine01 downlink-2 -- uplink-1 vni: 34 leaf12
hostbond4 -- swp2 server11

uplink-1 -- downlink-5 spine01 downlink-1 -- uplink-1 vni: 34 leaf11
hostbond4 -- swp1 server11
```

Alternately, you can use the IP address of the source device, as shown in this example.

```
cumulus@redis-1:~$ netq trace 00:02:00:00:02 vlan 1001 from 10.0.0.8 vrf vrf1 pretty
```

Number of Paths: 4

Number of Paths with Errors: 0

Number of Paths with Warnings: 0

Path MTU: 9152

```
server11 swp1 -- swp5 <vlan1000> tor-1 <vlan1001> vni: 34 uplink-2 -- downlink-5
spine02 downlink-2 -- uplink-2 vni: 34 <vlan1001> leaf12 hostbond4 -- swp2
server11
                                uplink-2 -- downlink-5 spine02 downlink-1 --
uplink-2 vni: 34 <vlan1001> leaf11 hostbond4 -- swp1 server11
     swp1 -- swp5 <vlan1000> tor-1 <vlan1001> vni: 34 uplink-1 -- downlink-5
spine01 downlink-2 -- uplink-1 vni: 34 <vlan1001> leaf12 hostbond4 -- swp2
server11
                                uplink-1 -- downlink-5 spine01 downlink-1 --
uplink-1 vni: 34 <vlan1001> leaf11 hostbond4 -- swp1 server11
```

View Paths between Two Switches with Detailed Output

This example provides the same path information as the pretty output, but displays the information in a tabular output.

```
cumulus@switch:~$ netq trace 00:02:00:00:00:02 vlan 1001 from 10.0.0.8 vrf vrf1
detail
Number of Paths: 4
Number of Paths with Errors: 0
Number of Paths with Warnings: 0
Path MTU: 9152
Id Hop Hostname InPort InVlan InTunnel InRtrIf InVRF
OutRtrIf OutVRF OutTunnel OutPort OutVlan
1 1
server11
swp1
         1000
```

```
2 leaf01 swp5 1000 vlan1000 vrf1
vlan1001 vrf1 vni: 34 uplink-2
 3 spine02 downlink-5
                    downlink-5 default
downlink-2 default
                    downlink-2
 4 leaf12 uplink-2 vni: 34 vlan1001
vrf1
                 hostbond4 1001
 5 server11 swp2
2 1
server11
swp1 1000
2 leaf01 swp5 1000 vlan1000 vrf1
vlan1001 vrf1 vni: 34 uplink-2
 3 spine02 downlink-5 default
downlink-1 default downlink-1
 4 leaf11 uplink-2 vni: 34 vlan1001
                  hostbond4 1001
vrf1
 5 server11 swp1
3 1
server11
swp1 1000
2 leaf01 swp5 1000 vlan1000 vrf1
vlan1001 vrf1 vni: 34 uplink-1
 3 spine01 downlink-5 default
downlink-2 default downlink-2
4 leaf12 uplink-1 vni: 34 vlan1001
```

vrf1 hostbond4 1001
5 server11 swp2
4 1
server11
swp1 1000
2 leaf01 swp5 1000 vlan1000 vrf1
vlan1001 vrf1 vni: 34 uplink-1
3 spine01 downlink-5 downlink-5 default
downlink-1 default downlink-1
4 leaf11 uplink-1 vni: 34 vlan1001
vrf1 hostbond4 1001
5 server11 swp1

Monitor Network Layer Protocols

With NetQ, a network administrator can monitor OSI Layer 3 network protocols running on Linux-based hosts, including IP (Internet Protocol), BGP (Border Gateway Protocol) and OSPF (Open Shortest Path First). NetQ provides the ability to:

- Validate protocol configurations
- Validate layer 3 communication paths

It helps answer questions such as:

- Who are the IP neighbors for a switch?
- How many IPv4 and IPv6 addresses am I using?
- When did changes occur to my IP configuration?
- Is BGP working as expected?
- Is OSPF working as expected?
- Can device A reach device B using IP addresses?

Monitor IP Configuration

NetQ enables you to view the current status and the status an earlier point in time. From this information, you can:

- determine IP addresses of one or more interfaces
- determine IP neighbors for one or more devices
- determine IP routes owned by a device
- identify changes to the IP configuration

The **netq show ip** command is used to obtain the address, neighbor, and route information from the devices. Its syntax is:

```
netq <hostname> show ip addresses [<remote-interface>] [<ipv4>|<ipv4/|
prefixlen>] [vrf <vrf>] [around <text-time>] [count] [json]
netq [<hostname>] show ip addresses [<remote-interface>] [<ipv4>|<ipv4/|
prefixlen>] [vrf <vrf>] [around <text-time>] [json]
netq <hostname> show ip neighbors [<remote-interface>] [<ipv4>|<ipv4> vrf
<vrf>|vrf <vrf>] [<mac>] [around <text-time>] [json]
netq [<hostname>] show ip neighbors [<remote-interface>] [<ipv4>|<ipv4> vrf
<vrf>|vrf <vrf>] [<mac>] [around <text-time>] [count] [json]
netq <hostname> show ip routes [<ipv4>|<ipv4/prefixlen>] [vrf <vrf>] [origin]
[around <text-time>] [count] [json]
netq [<hostname>] show ip routes [<ipv4>|<ipv4/prefixlen>] [vrf <vrf>] [origin]
[around <text-time>] [json]
```

netq <hostname> show ipv6 addresses [<remote-interface>] [<ipv6>|<ipv6/
prefixlen>] [vrf <vrf>] [around <text-time>] [count] [json]
netq [<hostname>] show ipv6 addresses [<remote-interface>] [<ipv6>|<ipv6/
prefixlen>] [vrf <vrf>] [around <text-time>] [json]
netq <hostname> show ipv6 neighbors [<remote-interface>] [<ipv6>|<ipv6> vrf
<vrf>|vrf <vrf>] [<mac>] [around <text-time>] [count] [json]
netq [<hostname>] show ipv6 neighbors [<remote-interface>] [<ipv6>|<ipv6> vrf
<vrf>|vrf <vrf>] [<mac>] [around <text-time>] [json]
netq <hostname> show ipv6 routes [<ipv6>|<ipv6/prefixlen>] [vrf <vrf>] [origin]
[around <text-time>] [count] [json]
netq [<hostname>] show ipv6 routes [<ipv6>|<ipv6/prefixlen>] [vrf <vrf>] [origin]
[around <text-time>] [json]

(i) NOTE

When entering a time value, you must include a numeric value and the unit of measure:

- d: day(s)
- h: hour(s)
- m: minute(s)
- s: second(s)
- now

For time ranges, the <text-time> is the most recent time and the <textendtime> is the oldest time. The values do not have to have the same unit of measure.

View IP Address Information

You can view the IPv4 and IPv6 address information for all of your devices, including the interface and VRF for each device. Additionally, you can:

- view the information at an earlier point in time
- filter against a particular device, interface or VRF assignment
- obtain a count of all of the addresses.

Each of these provides information for troubleshooting potential configuration and communication issues at the layer 3 level.

Example: View IPv4 address information for all devices

Address	Hostname	Interface	VRF	Last Changed
10.0.0.11/32	leaf01	lo	default	Thu Feb 7 18:30:53
2019				
10.0.0.12/32	leaf02	lo	default	Thu Feb 7 18:30:53
2019				
10.0.0.13/32	leaf03	lo	default	Thu Feb 7 18:30:53
2019				
10.0.0.14/32	leaf04	lo	default	Thu Feb 7 18:30:53
2019				
10.0.0.21/32	spine01	lo	default	Thu Feb 7 18:30:53
2019				
10.0.0.22/32	spine02	lo	default	Thu Feb 7 18:30:53
2019				
10.0.0.254/32	oob-mgmt	-server eth0	de	efault Thu Feb 7
18:30:53 2019				
172.16.1.1/24	leaf01	br0	default	Thu Feb 7 18:30:53
2019				
172.16.1.101/24	server01	eth1	defau	llt Thu Feb 7
18:30:53 2019				
172.16.2.1/24	leaf02	br0	default	Thu Feb 7 18:30:53
2019				
172.16.2.101/24	server02	eth2	defau	llt Thu Feb 7
18:30:53 2019				
172.16.3.1/24	leaf03	br0	default	Thu Feb 7 18:30:53
2019				
172.16.3.101/24	server03	eth1	defau	ılt Thu Feb 7

18:30:53 2019				
172.16.4.1/24	leaf04	br0	default	Thu Feb 7 18:30:53
2019				
172.16.4.101/24	server04	eth2	default	Thu Feb 7
18:30:53 2019				
172.17.0.1/16	oob-mgmt-	server docker0	def	Fault Thu Feb 7
18:30:53 2019				
192.168.0.11/24	leaf01	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.12/24	leaf02	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.13/24	leaf03	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.14/24	leaf04	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.21/24	spine01	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.22/24	spine02	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.254/24	oob-mgm	nt-server eth1	def	fault Thu Feb 7
18:30:53 2019				
192.168.0.31/24	server01	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.32/24	server02	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.33/24	server03	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.34/24	server04	eth0	default	Thu Feb 7
18:30:53 2019				

Example: View IPv6 address information for all devices

Matching add	dress records:			
Address	Hostname	Interface	VRF L	ast Changed
e80::203:ff:fe	e11:1101/64 server01	eth1	default	Thu Feb 7
18:30:53 2019	9			
fe80::203:ff:fe	e22:2202/64 server02	eth2	default	Thu Feb 7
18:30:53 2019	9			
fe80::203:ff:fe	e33:3301/64 server03	eth1	default	Thu Feb 7
18:30:53 2019	9			
fe80::203:ff:fe	e44:4402/64 server04	eth2	default	Thu Feb 7
18:30:53 2019	9			
fe80::4638:39	ff:fe00:18/6 leaf02	br0	default	Thu Feb 7
18:30:53 2019	e			
fe80::4638:39	ff:fe00:1b/6 leaf03	swp52	default	Thu Feb 7
18:30:53 2019	9			
fe80::4638:39	ff:fe00:1c/6 spine02	swp3	default	Thu Feb 7
18:30:53 2019	9			
fe80::4638:39	ff:fe00:23/6 leaf03	br0	default	Thu Feb 7
18:30:53 2019	9			
fe80::4638:39	ff:fe00:24/6 leaf01	swp52	default	Thu Feb 7
18:30:53 2019	9			
fe80::4638:39	ff:fe00:25/6 spine02	swp1	default	Thu Feb 7
18:30:53 2019	9			
fe80::4638:39	ff:fe00:28/6 leaf02	swp51	default	Thu Feb 7
18:30:53 2019	9			
fe80::4638:39	ff:fe00:29/6 spine01	swp2	default	Thu Feb 7

18:30:53 2019			
fe80::4638:39ff:fe00:2c/6 leaf04	br0	default	Thu Feb 7
18:30:53 2019			
fe80::4638:39ff:fe00:3/64 leaf01	br0	default	Thu Feb 7
18:30:53 2019			
fe80::4638:39ff:fe00:3b/6 leaf04	swp51	default	Thu Feb 7
18:30:53 2019			
fe80::4638:39ff:fe00:3c/6 spine01	swp4	default	Thu Feb 7
18:30:53 2019			
fe80::4638:39ff:fe00:46/6 leaf04	swp52	default	Thu Feb 7
18:30:53 2019			
fe80::4638:39ff:fe00:47/6 spine02	swp4	default	Thu Feb 7
18:30:53 2019			
fe80::4638:39ff:fe00:4f/6 leaf03	swp51	default	Thu Feb 7
18:30:53 2019			
fe80::4638:39ff:fe00:50/6 spine01	swp3	default	Thu Feb 7
18:30:53 2019			
fe80::4638:39ff:fe00:53/6 leaf01	swp51	default	Thu Feb 7
18:30:53 2019			
fe80::4638:39ff:fe00:54/6 spine01	swp1	default	Thu Feb 7
18:30:53 2019			
fe80::4638:39ff:fe00:57/6 oob-mgm	it-server eth1	defa	ault Thu Feb
7 18:30:53 2019			
fe80::4638:39ff:fe00:5d/6 leaf02	swp52	default	Thu Feb 7
18:30:53 2019			
fe80::4638:39ff:fe00:5e/6 spine02	swp2	default	Thu Feb 7
18:30:53 2019			
fe80::5054:ff:fe77:c277/6 oob-mgm	t-server eth0	defa	ault Thu Feb
7 18:30:53 2019			

18:30:53 2019 fe80::a200:ff:fe00:12/64 leaf02 eth0 default Thu Feb 7 18:30:53 2019
18:30:53 2019
fe80::a200:ff:fe00:13/64 leaf03 eth0 default Thu Feb 7
18:30:53 2019
fe80::a200:ff:fe00:14/64 leaf04 eth0 default Thu Feb 7
18:30:53 2019
fe80::a200:ff:fe00:21/64 spine01 eth0 default Thu Feb 7
18:30:53 2019
fe80::a200:ff:fe00:22/64 spine02 eth0 default Thu Feb 7
18:30:53 2019
fe80::a200:ff:fe00:31/64 server01 eth0 default Thu Feb 7
18:30:53 2019
fe80::a200:ff:fe00:32/64 server02 eth0 default Thu Feb 7
18:30:53 2019
fe80::a200:ff:fe00:33/64 server03 eth0 default Thu Feb 7
18:30:53 2019
fe80::a200:ff:fe00:34/64 server04 eth0 default Thu Feb 7
18:30:53 2019

Example: Filter IP Address Information for a Specific Interface

This example shows the IPv4 address information for the eth0 interface on all devices.

cumulus@switch:~\$ netq show ip addresses eth0

Matching address records:

Address Hostname Interface VRF Last Changed

10.0.0.254/32	oob-mgmt-s	server eth0	defau	ılt Thu Feb 7
18:30:53 2019				
192.168.0.11/24	leaf01	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.12/24	leaf02	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.13/24	leaf03	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.14/24	leaf04	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.21/24	spine01	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.22/24	spine02	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.31/24	server01	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.32/24	server02	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.33/24	server03	eth0	default	Thu Feb 7
18:30:53 2019				
192.168.0.34/24	server04	eth0	default	Thu Feb 7
18:30:53 2019				

Example: Filter IP Address Information for a Specific Device

This example shows the IPv6 address information for the leaf01 switch.

Matching address	s records:			
Address	Hostname	Interface	VRF La	st Changed
2001:c15c:d06:f0	od::16/12 leaf01	lo	default	Fri Feb 8
00:35:07 2019				
8				
2001:cafe:babe:0	22::/128 leaf01	DataVrf1080	DataVı	rf1080 Fri Feb
8 00:35:07 2019				
2001:cafe:babe:1	22::/128 leaf01	DataVrf1081	DataVı	rf1081 Fri Feb
8 00:35:07 2019				
2001:cafe:babe:2	22::/128 leaf01	DataVrf1082	DataVı	rf1082 Fri Feb
8 00:35:07 2019				
2001:fee1:600d:1	0::1/64 leaf01	VlanA-1.102	DataVrf	1082 Fri Feb
8 00:35:07 2019				
2001:fee1:600d:1	1::1/64 leaf01	VlanA-1.103	default	Fri Feb 8
00:35:07 2019				
2001:fee1:600d:1	2::1/64 leaf01	VlanA-1.104	default	Fri Feb 8
00:35:07 2019				
2001:fee1:600d:1	3::1/64 leaf01	VlanA-1.105	default	Fri Feb 8
00:35:07 2019				
	4::1/64 leaf01	VlanA-1.106	default	Fri Feb 8
00:35:07 2019				
	::1/64 leaf01	VlanA-1.100	DataVrf1	1080 Fri Feb 8
00:35:07 2019				
	:1/64 leaf01	VlanA-1.101	DataVrf1	081 Fri Feb 8
00:35:07 2019				
2001:fee1:d00d:1	::1/64 leaf01	vlan1001-v0	vrf1	Fri Feb 8

2001:fee1:d00d:1::2/64	leaf01	vlan1001	vrf1	Fri Feb 8
00:35:07 2019				
2001:fee1:d00d:2::1/64	leaf01	vlan1002-v0	vrf1	Fri Feb 8
00:35:07 2019				

Example: View Changes to IP Address Information

This example shows the IPv4 address information that changed for all devices around 1 day ago.

Matching address					
Address	Hostname	Interface	VRF	Last Changed	
192.168.0.15/24	leaf01	eth0	mgmt	Thu Feb 7	
22:49:26 2019					
27.0.0.22/32	leaf01	lo	default	Thu Feb 7 22:49:26	
2019					
3.0.3.129/26	leaf01	VlanA-1.100	DataVrf	f1080 Thu Feb 7	
22:49:26 2019					
3.0.3.193/26	leaf01	VlanA-1.101	DataVrf	f1081 Thu Feb 7	
22:49:26 2019					
3.0.4.1/26	leaf01	VlanA-1.102	DataVrf1	082 Thu Feb 7	
22:49:26 2019					
3.0.4.129/26	leaf01	VlanA-1.104	default	Thu Feb 7	
22:49:26 2019					
3.0.4.193/26	leaf01	VlanA-1.105	default	Thu Feb. 7	

3.0.4.65/26	leaf01	VlanA-1.103	default Thu Feb 7
22:49:26 2019			
3.0.5.1/26	leaf01	VlanA-1.106	default Thu Feb 7
22:49:26 2019			
30.0.0.22/32	leaf01	DataVrf1080	DataVrf1080 Thu Feb 7
22:49:26 2019			
30.0.1.22/32	leaf01	DataVrf1081	DataVrf1081 Thu Feb 7
22:49:26 2019			
30.0.2.22/32	leaf01	DataVrf1082	DataVrf1082 Thu Feb 7
22:49:26 2019			
45.0.0.13/26	leaf01	NetQBond-1	mgmt Thu Feb 7
22:49:26 2019			
6.0.0.1/26	leaf01	vlan1000-v0	vrf1 Thu Feb 7 22:49:26
2019			
6.0.0.129/26	leaf01	vlan1002-v0	vrf1 Thu Feb 7
22:49:26 2019			

Example: Obtain a Count of IP Addresses Used on a Node

This example shows the number of IPv4 and IPv6 addresses on the node leaf01. Note that you must specify a hostname to use the count option.

cumulus@switch:~\$ netq leaf01 show ip addresses count

Count of matching address records: 33

cumulus@switch:~\$ netq leaf01 show ipv6 addresses count

Count of matching address records: 42

View IP Neighbor Information

You can view the IPv4 and IPv6 neighbor information for all of your devices, including the interface port, MAC address, VRF assignment, and whether it learns the MAC address from the peer (remote=yes). Additionally, you can:

- view the information at an earlier point in time
- filter against a particular device, interface, address or VRF assignment
- obtain a count of all of the addresses

Each of these provides information for troubleshooting potential configuration and communication issues at the layer 3 level.

Example: View IPv4 Neighbor Information for All Devices

cumulus@switch:~\$ netq show ip neighbors Matching neighbor records:						
		Interfoce	MAC Address VRF			
		Interrace	WAC Address VRF			
Remote Last Ch	anged					
10.255.5.1	oob-mgmt-	server eth0	52:54:00:0f:79:30			
default no	Thu Feb 7 22	:49:26 2019				
169.254.0.1	leaf01	swp51	44:38:39:00:00:54 default			
no Thu Feb 7	22:49:26 2019					
169.254.0.1	leaf01	swp52	44:38:39:00:00:25 default			
no Thu Feb 7	22:49:26 2019					
169.254.0.1	leaf02	swp51	44:38:39:00:00:29 default			
no Thu Feb 7	22:49:26 2019					
169.254.0.1	leaf02	swp52	44:38:39:00:00:5e default			
no Thu Feb 7	22:49:26 2019					

169.254.0.1 leaf03	swp51	44:38:39:00:00:50 default
no Thu Feb 7 22:49:26 2019	3WP31	44.30.33.00.00.30 acidali
169.254.0.1 leaf03	swp52	44:38:39:00:00:1c default
no Thu Feb 7 22:49:26 2019	3ννρ32	44.30.33.00.00.1C default
169.254.0.1 leaf04	swp51	44:38:39:00:00:3c default
	Swball	44.56.59.00.00.5C default
no Thu Feb 7 22:49:26 2019	aven E 2	44.20.20.00.00.47 default
169.254.0.1 leaf04	swp52	44:38:39:00:00:47 default
no Thu Feb 7 22:49:26 2019	aa 1	44.20.20.00.00.52 defects
169.254.0.1 spine01	swpı	44:38:39:00:00:53 default
no Thu Feb 7 22:49:26 2019		44.20.20.00.00.20 . [.]
169.254.0.1 spine01	swp2	44:38:39:00:00:28 default
no Thu Feb 7 22:49:26 2019		
169.254.0.1 spine01	swp3	44:38:39:00:00:4f default
no Thu Feb 7 22:49:26 2019		
169.254.0.1 spine01	swp4	44:38:39:00:00:3b default
no Thu Feb 7 22:49:26 2019		
169.254.0.1 spine02	swp1	44:38:39:00:00:24 default
no Thu Feb 7 22:49:26 2019		
169.254.0.1 spine02	swp2	44:38:39:00:00:5d default
no Thu Feb 7 22:49:26 2019		
169.254.0.1 spine02	swp3	44:38:39:00:00:1b default
no Thu Feb 7 22:49:26 2019		
169.254.0.1 spine02	swp4	44:38:39:00:00:46 default
no Thu Feb 7 22:49:26 2019		
192.168.0.11 oob-mgmt	-server eth1	a0:00:00:00:00:11
default no Thu Feb 7 22:	49:26 2019	
192.168.0.12 oob-mgmt	-server eth1	a0:00:00:00:00:12
default no Thu Feb 7 22:	49:26 2019	
192.168.0.13 oob-mgmt	-server eth1	a0:00:00:00:00:13

default no	Thu Feb 7 22:49:26 2019	
192.168.0.14	oob-mgmt-server eth1	a0:00:00:00:00:14
default no	Thu Feb 7 22:49:26 2019	
192.168.0.21	oob-mgmt-server eth1	a0:00:00:00:00:21
default no	Thu Feb 7 22:49:26 2019	
192.168.0.22	oob-mgmt-server eth1	a0:00:00:00:00:22
default no	Thu Feb 7 22:49:26 2019	
192.168.0.253	oob-mgmt-server eth1	a0:00:00:00:50
default no	Thu Feb 7 22:49:26 2019	
192.168.0.254	leaf01 eth0	44:38:39:00:00:57 default
no Thu Feb 7	22:49:26 2019	
192.168.0.254	leaf02 eth0	44:38:39:00:00:57 default
no Thu Feb 7	22:49:26 2019	
•••		

Example: View IPv6 Neighbor Information for a Given Device

This example shows the IPv6 neighbors for leaf02 switch.

cumulus@switch	n\$ netq leaf02 sh	ow ipv6 neighb	ors			
Matching neighbor records:						
IP Address	Hostname	Interface	MAC Address VRF			
Remote Last Cha	anged					
fe80::203:ff:fe22	:2202 leaf02	br0	00:03:00:22:22:02			
default no Thu Feb 7 22:49:26 2019						
fe80::4638:39ff:f	e00:29 leaf02	swp51	44:38:39:00:00:29			

default no Thu Feb 7 22:49	:26 2019	
fe80::4638:39ff:fe00:4 leaf02	eth0	44:38:39:00:00:04
default no Thu Feb 7 22:49	:26 2019	
fe80::4638:39ff:fe00:5e leaf02	swp52	44:38:39:00:00:5e
default no Thu Feb 7 22:49	:26 2019	
fe80::a200:ff:fe00:31 leaf02	eth0	a0:00:00:00:00:31
default no Thu Feb 7 22:49	:26 2019	
fe80::a200:ff:fe00:32 leaf02	eth0	a0:00:00:00:00:32
default no Thu Feb 7 22:49	:26 2019	
fe80::a200:ff:fe00:33 leaf02	eth0	a0:00:00:00:00:33
default no Thu Feb 7 22:49	:26 2019	
fe80::a200:ff:fe00:34 leaf02	eth0	a0:00:00:00:34
default no Thu Feb 7 22:49	:26 2019	

View IP Routes Information

You can view the IPv4 and IPv6 routes for all of your devices, including the IP address (with or without mask), the destination (by hostname) of the route, next hops available, VRF assignment, and whether a host is the owner of the route or MAC address.

Additionally, you can:

- view the information at an earlier point in time
- filter against a particular address or VRF assignment
- obtain a count of all of the routes

Each of these provides information for troubleshooting potential configuration and communication issues at the layer 3 level.

Example: View IP Routes for All Devices

This example shows the IPv4 and IPv6 routes for all devices in the network.

Matching rout	es records:				
Origin VRF	Prefix	Hostname		Nexthops	Last
Changed 					
yes default		server04	lo		Thu Feb 7
22:49:26 2019					
yes default	::/0	server03	lo		Thu Feb 7
22:49:26 2019					
yes default	::/0	server01	lo		Thu Feb 7
22:49:26 2019					
yes default	::/0	server02	lo		Thu Feb 7
22:49:26 2019					
	tch:~\$ netq show ip	routes			
		routes			
cumulus@swit	es records:	routes Hostname		Nexthops	Last
cumulus@swit Matching rout	es records:			Nexthops	Last
cumulus@swit Matching rout Origin VRF	es records: Prefix			Nexthops	Last
cumulus@swit Matching rout Origin VRF Changed	es records: Prefix	Hostname		· 	Last
cumulus@swit Matching rout Origin VRF Changed	es records: Prefix 080 3.0.3.128/26	Hostname		· 	Last
cumulus@swith Matching routh Origin VRF Changedyes DataVrf1 Fri Feb 8 00:4	es records: Prefix 080 3.0.3.128/26	Hostname leaf0	1	VlanA-1.100	Last
cumulus@swith Matching routh Origin VRF Changedyes DataVrf1 Fri Feb 8 00:4	es records: Prefix 080 3.0.3.128/26 6:17 2019 080 3.0.3.129/32	Hostname leaf0	1	VlanA-1.100	Last
cumulus@swith Matching routh Origin VRF Changedyes DataVrf1 Fri Feb 8 00:40 yes DataVrf1 Fri Feb 8 00:40	es records: Prefix 080 3.0.3.128/26 6:17 2019 080 3.0.3.129/32	Hostname leaf0	 1	VlanA-1.100 VlanA-1.100	Last
cumulus@swith Matching routh Origin VRF Changedyes DataVrf1 Fri Feb 8 00:40 yes DataVrf1 Fri Feb 8 00:40	es records: Prefix 080 3.0.3.128/26 6:17 2019 080 3.0.3.129/32 6:17 2019 080 30.0.0.22/32	Hostname leaf0	 1	VlanA-1.100 VlanA-1.100	Last

yes DataVrf108	3.0.3.193/32	leaf01	VlanA-1.101	
Fri Feb 8 00:46:	17 2019			
yes DataVrf108	30.0.1.22/32	leaf01	DataVrf1081	
Fri Feb 8 00:46:	17 2019			
yes DataVrf108	3.0.4.0/26	leaf01	VlanA-1.102	
Fri Feb 8 00:46:	17 2019			
yes DataVrf108	3.0.4.1/32	leaf01	VlanA-1.102	
Fri Feb 8 00:46:	17 2019			
yes DataVrf108	30.0.2.22/32	leaf01	DataVrf1082	
Fri Feb 8 00:46:	17 2019			
yes default	27.0.0.22/32	leaf01	lo	Fri Feb
8 00:46:17 2019				
yes default	3.0.4.128/26	leaf01	VlanA-1.104	Fri
Feb 8 00:46:17 2	2019			
yes default	3.0.4.129/32	leaf01	VlanA-1.104	Fri
Feb 8 00:46:17 2	2019			
yes default	3.0.4.192/26	leaf01	VlanA-1.105	Fri
Feb 8 00:46:17 2	2019			
yes default	3.0.4.193/32	leaf01	VlanA-1.105	Fri
Feb 8 00:46:17 2	2019			

Example: View IP Routes for a Given IP Address

This example shows the routes available for an IP address of 10.0.0.12.

cumulus@switch:~\$ netq show ip routes 10.0.0.12 Matching routes records:

Origin VRF	Prefix	Hostname	Nexthops	Last
Changed				
no default	10.0.0.12/32	leaf03	10.0.0.21: swp51, 1	0.0.0.22:
swp52 Fri Feb	8 00:46:17 2019			
no default	10.0.0.12/32	leaf01	10.0.0.21: swp51, 1	0.0.0.22:
swp52 Fri Feb	8 00:46:17 2019			
no default	10.0.0.12/32	leaf04	10.0.0.21: swp51, 1	0.0.0.22:
swp52 Fri Feb	8 00:46:17 2019			
no default	10.0.0.12/32	spine02	10.0.0.12: swp2	
Fri Feb 8 00:46	:17 2019			
no default	10.0.0.12/32	spine01	10.0.0.12: swp2	
Fri Feb 8 00:46	:17 2019			
yes default	10.0.0.12/32	leaf02	lo	Fri Feb
8 00:46:17 2019	9			

Example: View IP Routes Owned by a Given Device

This example shows the IPv4 routes that are owned by spine01 switch.

cumulus@swit	ch:~\$ netq spine01	show ip routes o	rigin	
Origin VRF Changed	Prefix	Hostname	Nexthops	Last
yes default	10.0.0.21/32	spine01	lo	Fri Feb

8 00:46:17 2019								
yes default	192.168.0.0/24	spine01	eth0	Fri				
Feb 8 00:46:17 2019								
yes default	192.168.0.21/32	spine01	eth0	Fri				
Feb 8 00:46:17 2019								

Example: View IP Routes for a Given Device at a Prior Time

This example show the IPv4 routes for spine01 switch about 24 hours ago.

Matching rout		I I a atra a ras a	Ni asatia asa	14				
	Prefix	Hostname	Nextnops	Last				
Changed								
	10.0.0.11/32	spine01	169.254.0.1: swp1					
Fri Feb 8 00:46:17 2019								
no default	10.0.0.12/32	spine01	169.254.0.1: swp2					
Fri Feb 8 00:46:17 2019								
no default	10.0.0.13/32	spine01	169.254.0.1: swp3					
Fri Feb 8 00:46:17 2019								
no default	10.0.0.14/32	spine01	169.254.0.1: swp4					
Fri Feb 8 00:46:17 2019								
no default	172.16.1.0/24	spine01	169.254.0.1: swp1					
Fri Feb 8 00:46:17 2019								
no default	172.16.2.0/24	spine01	169.254.0.1: swp2					

no	default	172.16.3.0/24	spine01	169.254.0.1: swp3			
Fri Feb 8 00:46:17 2019							
no	default	172.16.4.0/24	spine01	169.254.0.1: swp4			
Fri Feb 8 00:46:17 2019							
yes	default	10.0.0.21/32	spine01	lo	Fri Feb		
8 00:46:17 2019							
yes	default	192.168.0.0/24	spine01	eth0	Fri		
Feb 8 00:46:17 2019							
yes	default	192.168.0.21/32	spine01	eth0	Fri		
Feb 8 00:46:17 2019							

Example: View the Number of IP Routes on a Node

This example shows the total number of IP routes for all devices on a node.

cumulus@switch:~\$ netq leaf01 show ip routes count

Count of matching routes records: 125

cumulus@switch:~\$ netq leaf01 show ipv6 routes count

Count of matching routes records: 5

Monitor BGP Configuration

If you have BGP running on your switches and hosts, you can monitor its operation using the NetQ CLI. For each device, you can view its associated neighbors, ASN

(autonomous system number), peer ASN, receive IP or EVPN address prefixes, and VRF assignment. Additionally, you can:

- view the information at an earlier point in time
- filter against a particular device, ASN, or VRF assignment
- validate it is operating correctly across the network

The netq show bgp command is used to obtain the BGP configuration information from the devices. The netq check bgp command is used to validate the configuration. The syntax of these commands is:

netq [<hostname>] show bgp [<bgp-session>|asn <number-asn>] [vrf <vrf>]
[around <text-time>] [json]
netq [<hostname>] show events [level info|level error|level warning|level critical|
level debug] type bgp [between <text-time> and <text-endtime>] [json]
netq check bgp [vrf <vrf>] [around <text-time>] [json]

(i) NOTE

When entering a time value, you must include a numeric value and the unit of measure:

- d: day(s)
- h: hour(s)
- m: minute(s)
- s: second(s)
- now

For time ranges, the <text-time> is the most recent time and the <textendtime> is the oldest time. The values do not have to have the same unit of measure.

View BGP Configuration Information

NetQ enables you to view the BGP configuration of a single device or across all of your devices at once. You can filter the results based on an ASN, BGP session (IP address or interface name), or VRF assignment. You can view the configuration in the past and view changes made to the configuration within a given timeframe.

Example: View BGP Configuration Information Across Network

This example shows the BGP configuration across all of your switches. In this scenario, BGP routing is configured between two spines and four leafs. Each leaf switch has a unique ASN and the spine switches share an ASN. The PfxRx column indicates that these devices have IPv4 address prefixes. The second and third values in this column indicate IPv6 and EVPN address prefixes when configured. This configuration was changed just over one day ago.

Matching	bgp records:					
Hostname	e Neighbor	VRF	ASN	Pee	er ASN	PfxRx Last
Changed						
exit-1	swp3(spine-1)	default	655537	65	5435 2	29/25/434
Thu Feb 7	7 18:19:50 2019					
exit-1	swp3.2(spine-1)	DataVrf108	0 655	537	655435	5 15/13/0
Thu Feb 7	7 18:19:50 2019					
exit-1	swp3.3(spine-1)	DataVrf108	1 655	537	65543	5 14/13/0
Thu Feb 7	7 18:19:50 2019					
exit-1	swp3.4(spine-1)	DataVrf108	2 655	537	65543	5 16/13/0
Thu Feb 7	7 18:19:50 2019					
exit-1	swp4(spine-2)	default	655537	65	5435	29/25/434
Thu Feb 7	7 18:19:50 2019					
exit-1	swp4.2(spine-2)	DataVrf108	0 655	5537	65543	5 16/13/0
Thu Feb 7	7 18:19:50 2019					
exit-1	swp4.3(spine-2)	DataVrf108	1 655	537	65543	5 14/13/0
Thu Feb 7	7 18:19:50 2019					
exit-1	swp4.4(spine-2)	DataVrf108	2 655	537	65543	5 16/13/0
Thu Feb 7	7 18:19:50 2019					
exit-1	swp5(spine-3)	default	655537	65	5435	30/25/434
Thu Feb 7	7 18:19:50 2019					
exit-1	swp5.2(spine-3)	DataVrf108	0 655	537	65543	5 15/13/0
	7 18:19:50 2019					
	swp5.3(spine-3)	DataVrf108	1 655	5537	65543	5 14/13/0
Thu Feb 7	7 18:19:50 2019					
exit-1	swp5.4(spine-3)	DataVrf108	2 655	537	65543	5 16/13/0

Thu Feb 7 18:19:50 2019	
exit-1 swp7	default 655537 - NotEstd Thu Feb 7
18:31:44 2019	
exit-1 swp7.2	DataVrf1080 655537 - NotEstd Thu
Feb 7 18:31:44 2019	
exit-1 swp7.3	DataVrf1081 655537 - NotEstd Thu
Feb 7 18:31:44 2019	
exit-1 swp7.4	DataVrf1082 655537 - NotEstd Thu
Feb 7 18:31:44 2019	
exit-2 swp3(spine-1)	default 655538 655435 28/24/434
Thu Feb 7 18:19:50 2019	
exit-2 swp3.2(spine-1)	DataVrf1080 655538 655435 14/12/0
Thu Feb 7 18:19:50 2019	
exit-2 swp3.3(spine-1)	DataVrf1081 655538 655435 15/12/0
Thu Feb 7 18:19:50 2019	
exit-2 swp3.4(spine-1)	DataVrf1082 655538 655435 15/12/0
Thu Feb 7 18:19:50 2019	
exit-2 swp4(spine-2)	default 655538 655435 28/24/434
Thu Feb 7 18:19:50 2019	
exit-2 swp4.2(spine-2)	DataVrf1080 655538 655435 14/12/0
Thu Feb 7 18:19:50 2019	
exit-2 swp4.3(spine-2)	DataVrf1081 655538 655435 15/12/0
Thu Feb 7 18:19:50 2019	
exit-2 swp4.4(spine-2)	DataVrf1082 655538 655435 15/12/0
Thu Feb 7 18:19:50 2019	
exit-2 swp5(spine-3)	default 655538 655435 27/24/434
Thu Feb 7 18:19:50 2019	
exit-2 swp5.2(spine-3)	DataVrf1080 655538 655435 15/12/0
Thu Feb 7 18:19:50 2019	

exit-2	swp5.3(spine-3)	DataVrf1081 655538 655435 15/12/0
Thu Feb 7	7 18:19:50 2019	
exit-2	swp5.4(spine-3)	DataVrf1082 655538 655435 15/12/0
Thu Feb 7	7 18:19:50 2019	
exit-2	swp7	default 655538 - NotEstd Thu Feb 7
18:31:49 2	2019	
exit-2	swp7.2	DataVrf1080 655538 - NotEstd Thu
Feb 7 18:3	31:49 2019	
exit-2	swp7.3	DataVrf1081 655538 - NotEstd Thu
Feb 7 18:3	31:49 2019	
exit-2	swp7.4	DataVrf1082 655538 - NotEstd Thu
Feb 7 18:	31:49 2019	
spine-1	swp10(exit-2)	default 655435 655538 10/5/0 Thu
Feb 718:	19:50 2019	
spine-1	swp10.2(exit-2)	DataVrf1080 655435 655538 10/5/0
Thu Feb 7	7 18:19:50 2019	
spine-1	swp10.3(exit-2)	DataVrf1081 655435 655538 10/5/0
Thu Feb 7	7 18:19:50 2019	
spine-1	swp10.4(exit-2)	DataVrf1082 655435 655538 10/5/0
Thu Feb 7	7 18:19:50 2019	
spine-1	swp3(leaf-11)	default 655435 655559 19/6/94
Thu Feb 7	7 18:19:50 2019	
spine-1	swp3.2(leaf-11)	DataVrf1080 655435 655559 14/2/0
Thu Feb 7	7 18:19:50 2019	
spine-1	swp3.3(leaf-11)	DataVrf1081 655435 655559 14/2/0
Thu Feb 7	7 18:19:50 2019	
spine-1	swp3.4(leaf-11)	DataVrf1082 655435 655559 14/2/0
Thu Feb 7	7 18:19:50 2019	
spine-1	swp4(leaf-12)	default 655435 655560 19/6/64

Thu Feb 7 18:19:50 2019	
spine-1 swp4.2(leaf-12)	DataVrf1080 655435 655560 14/2/0
Thu Feb 7 18:19:50 2019	
spine-1 swp4.3(leaf-12)	DataVrf1081 655435 655560 14/2/0
Thu Feb 7 18:19:50 2019	
spine-1 swp4.4(leaf-12)	DataVrf1082 655435 655560 14/2/0
Thu Feb 7 18:19:50 2019	
spine-1 swp5(leaf-21)	default 655435 655561 19/6/50
Thu Feb 7 18:19:50 2019	
spine-1 swp5.2(leaf-21)	DataVrf1080 655435 655561 14/2/0
Thu Feb 7 18:19:50 2019	
spine-1 swp5.3(leaf-21)	DataVrf1081 655435 655561 14/2/0
Thu Feb 7 18:19:50 2019	
spine-1 swp5.4(leaf-21)	DataVrf1082 655435 655561 14/2/0
Thu Feb 7 18:19:50 2019	
spine-1 swp6(leaf-22)	default 655435 655562 19/6/62
Thu Feb 7 18:19:50 2019	
spine-1 swp6.2(leaf-22)	DataVrf1080 655435 655562 14/2/0
Thu Feb 7 18:19:50 2019	
spine-1 swp6.3(leaf-22)	DataVrf1081 655435 655562 14/2/0
Thu Feb 7 18:19:50 2019	
spine-1 swp6.4(leaf-22)	DataVrf1082 655435 655562 14/2/0
Thu Feb 7 18:19:50 2019	
spine-1 swp7(leaf-1)	default 655435 655557 17/5/54 Thu
Feb 7 18:19:50 2019	
spine-1 swp7.2(leaf-1)	DataVrf1080 655435 655557 14/2/0
Thu Feb 7 18:19:50 2019	
spine-1 swp7.3(leaf-1)	DataVrf1081 655435 655557 14/2/0
Thu Feb 7 18:19:50 2019	

spine-1	swp7.4(leaf-1)	DataVrf1082	655435	655557	14/2/0
Thu Feb 71	18:19:50 2019				
spine-1	swp8(leaf-2)	default 65	5435 65	5558 17/	/5/54 Thu
Feb 7 18:19):50 2019				
spine-1	swp8.2(leaf-2)	DataVrf1080	655435	655558	14/2/0
Thu Feb 71	18:19:50 2019				
spine-1	swp8.3(leaf-2)	DataVrf1081	655435	655558	14/2/0
Thu Feb 71	18:19:50 2019				
spine-1	swp8.4(leaf-2)	DataVrf1082	655435	655558	14/2/0
Thu Feb 71	18:19:50 2019				
spine-1	swp9(exit-1)	default 65	5435 65	5537 19/	/5/0 Thu
Feb 7 18:19	9:50 2019				
spine-1	swp9.2(exit-1)	DataVrf1080	655435	655537	19/5/0
Thu Feb 71	18:19:50 2019				
spine-1	swp9.3(exit-1)	DataVrf1081	655435	655537	19/5/0
Thu Feb 71	18:19:50 2019				
spine-1	swp9.4(exit-1)	DataVrf1082	655435	655537	19/5/0
Thu Feb 71	18:19:50 2019				
spine-2	swp10(exit-2)	default 65	5435 65	55538 10)/5/0 Thu
Feb 7 18:19	9:50 2019				
spine-2	swp10.3(exit-2)	DataVrf1081	655435	655538	10/5/0
Thu Feb 71	18:19:50 2019				
spine-2	swp10.4(exit-2)	DataVrf1082	655435	655538	10/5/0
Thu Feb 71	18:19:50 2019				
spine-2	swp3.2(leaf-11)	DataVrf1080	655435	655559	14/2/0
Thu Feb 71	18:19:50 2019				

Example: View BGP Configuration Information for a Given Device

This example shows the BGP configuration information for the spine02 switch. The switch is peered with swp1 on leaf01, swp2 on leaf02, and so on. Spine02 has an ASN of 65020 and each of the leafs have unique ASNs.

	bgp records:						
Hostname	e Neighbor	VRF	ASN	Pee	r ASN	PfxRx	Last
Changed							
spine02	swp3(spine01)	default	655	5557 6	555435	42/27	//324
Thu Feb 7	7 18:19:50 2019						
spine02	swp3.2(spine01)	DataVrf1	080	655557	655	435	
31/18/0	Thu Feb 7 18:19:50 2019						
spine02	swp3.3(spine01)	DataVrf1	081	655557	655	435	
31/18/0	Thu Feb 7 18:19:50 2019						
spine02	swp3.4(spine01)	DataVrf1	082	655557	655	435	
29/18/0	Thu Feb 7 18:19:50 2019						
spine02	swp5(spine03)	default	655	5557 6	555435	42/27	//324
Thu Feb 7	7 18:19:50 2019						
spine02	swp5.2(spine03)	DataVrf1	080	655557	655	435	
31/18/0	Thu Feb 7 18:19:50 2019						
spine02	swp5.3(spine03)	DataVrf1	081	655557	655	435	
31/18/0	Thu Feb 7 18:19:50 2019						
spine02	swp5.4(spine03)	DataVrf1	082	655557	655	435	
29/18/0	Thu Feb 7 18:19:50 2019						

Example: View BGP Configuration Information for a Given ASN

This example shows the BGP configuration information for ASN of *655557*. This ASN is associated with spine02 and so the results show the BGP neighbors for that switch.

Matching	bgp records:						
Hostnam	e Neighbor	VRF	ASN	Pe	er ASN	PfxRx	Last
Changed							
	swp3(spine01)	default	655	5557	655435	42/27	7/324
Thu Feb	7 18:19:50 2019						
spine02	swp3.2(spine01)	DataVrf10	080	65555	7 655	435	
31/18/0	Thu Feb 7 18:19:50 2019						
spine02	swp3.3(spine01)	DataVrf10	081	65555	7 655	435	
31/18/0	Thu Feb 7 18:19:50 2019						
spine02	swp3.4(spine01)	DataVrf10)82	65555	7 655	435	
29/18/0	Thu Feb 7 18:19:50 2019						
spine02	swp5(spine03)	default	655	5557	655435	42/27	7/324
Thu Feb	7 18:19:50 2019						
spine02	swp5.2(spine03)	DataVrf10	080	65555	7 655	435	
31/18/0	Thu Feb 7 18:19:50 2019						
spine02	swp5.3(spine03)	DataVrf10	081	65555	7 655	435	
31/18/0	Thu Feb 7 18:19:50 2019						
spine02	swp5.4(spine03)	DataVrf10	082	65555	7 655	435	
29/18/0	Thu Feb 7 18:19:50 2019						

Example: View BGP Configuration Information for a Prior Time

This example shows the BGP configuration information as it was 12 hours earlier.

Hostname	e Neighbor	VRF	ASN	Pee	r ASN I	PfxRx Last
Changed	J					
exit01	swp3(spine01)	default	6555	537 65	5435	29/25/434
Thu Feb 7	7 18:19:50 2019					
exit01	swp3.2(spine01)	DataVrf10	080	655537	65543	5 15/13/0
Thu Feb 7	7 18:19:50 2019					
exit01	swp3.3(spine01)	DataVrf10	081	655537	65543	5 14/13/0
Γhu Feb 7	7 18:19:50 2019					
exit01	swp3.4(spine01)	DataVrf10	082	655537	65543	5 16/13/0
Γhu Feb 7	7 18:19:50 2019					
exit01	swp4(spine02)	default	6555	537 65	5435	29/25/434
Thu Feb 7	7 18:19:50 2019					
exit01	swp4.2(spine02)	DataVrf10	080	655537	65543	5 16/13/0
Γhu Feb 7	7 18:19:50 2019					
exit01	swp4.3(spine02)	DataVrf10	081	655537	65543	5 14/13/0
Γhu Feb 7	7 18:19:50 2019					
exit01	swp4.4(spine02)	DataVrf10	082	655537	65543	5 16/13/0
Thu Feb 7	7 18:19:50 2019					
exit01	swp5(spine03)	default	6555	537 65	5435	30/25/434
Thu Feb 7	7 18:19:50 2019					
exit01	swp5.2(spine03)	DataVrf1	080	655537	65543	5 15/13/0

Thu Feb 7 18:19:50 2019	
exit01 swp5.3(spine03) DataVrf1	081 655537 655435 14/13/0
Thu Feb 7 18:19:50 2019	
exit01 swp5.4(spine03) DataVrf1	082 655537 655435 16/13/0
Thu Feb 7 18:19:50 2019	
exit01 swp6(firewall01) default	655537 655539 73/69/-
Thu Feb 7 18:26:30 2019	
exit01 swp6.2(firewall01) DataVrf1	080 655537 655539 73/69/-
Thu Feb 7 18:26:30 2019	
exit01 swp6.3(firewall01) DataVrf1	081 655537 655539 73/69/-
Thu Feb 7 18:26:30 2019	
exit01 swp6.4(firewall01) DataVrf1	082 655537 655539 73/69/-
Thu Feb 7 18:26:30 2019	
exit01 swp7 default 69	55537 - NotEstd Thu Feb 7
18:31:44 2019	
exit01 swp7.2 DataVrf1080	655537 - NotEstd Thu
Feb 7 18:31:44 2019	
exit01 swp7.3 DataVrf1081	655537 - NotEstd Thu
Feb 7 18:31:44 2019	
exit01 swp7.4 DataVrf1082	655537 - NotEstd Thu
Feb 7 18:31:44 2019	
exit02 swp3(spine01) default	655538 655435 28/24/434
Thu Feb 7 18:19:50 2019	
exit02 swp3.2(spine01) DataVrf1	080 655538 655435 14/12/0
Thu Feb 7 18:19:50 2019	
exit02 swp3.3(spine01) DataVrf1	081 655538 655435 15/12/0
Thu Feb 7 18:19:50 2019	
exit02 swp3.4(spine01) DataVrf1	082 655538 655435 15/12/0
Thu Feb 7 18:19:50 2019	

exit02 swp4(spine02)	default 655538 655435 28/24/434
Thu Feb 7 18:19:50 2019	
exit02 swp4.2(spine02)	DataVrf1080 655538 655435 14/12/0
Thu Feb 7 18:19:50 2019	
exit02 swp4.3(spine02)	DataVrf1081 655538 655435 15/12/0
Thu Feb 7 18:19:50 2019	
exit02 swp4.4(spine02)	DataVrf1082 655538 655435 15/12/0
Thu Feb 7 18:19:50 2019	
exit02 swp5(spine03)	default 655538 655435 27/24/434
Thu Feb 7 18:19:50 2019	
exit02 swp5.2(spine03)	DataVrf1080 655538 655435 15/12/0
Thu Feb 7 18:19:50 2019	
exit02 swp5.3(spine03)	DataVrf1081 655538 655435 15/12/0
Thu Feb 7 18:19:50 2019	
exit02 swp5.4(spine03)	DataVrf1082 655538 655435 15/12/0
Thu Feb 7 18:19:50 2019	
exit02 swp6(firewall01)	default 655538 655539 7/5/- Thu
Feb 7 18:26:30 2019	
exit02 swp6.2(firewall01)	DataVrf1080 655538 655539 7/5/-
Thu Feb 7 18:26:30 2019	
exit02 swp6.3(firewall01)	DataVrf1081 655538 655539 7/5/-
Thu Feb 7 18:26:30 2019	
exit02 swp6.4(firewall01)	DataVrf1082 655538 655539 7/5/-
Thu Feb 7 18:26:30 2019	
exit02 swp7	default 655538 - NotEstd Thu Feb 7
18:31:49 2019	
exit02 swp7.2	DataVrf1080 655538 - NotEstd Thu
Feb 7 18:31:49 2019	
ovi+02 cup7.2	DataVrf1081 655538 - NotEstd Thu
exit02 swp7.3	Dataviiiooi 055550 Notesta iiid

Feb 7 18:3	31:49 2019				
exit02	swp7.4	DataVrf1082	655538	- Not	tEstd Thu
Feb 7 18:3	31:49 2019				
firewall01	swp3(exit01)	default	655539	655537	29/27/-
Thu Feb 7	7 18:26:30 2019				
firewall01	swp3.2(exit01)	default	655539	655537	15/15/-
Thu Feb 7	7 18:26:30 2019				
firewall01	swp3.3(exit01)	default	655539	655537	15/15/-
Thu Feb 7	7 18:26:30 2019				
firewall01	swp3.4(exit01)	default	655539	655537	15/15/-
Thu Feb 7	7 18:26:30 2019				
firewall01	swp4(exit02)	default	655539	655538	29/27/-
Thu Feb 7	7 18:26:30 2019				
firewall01	swp4.2(exit02)	default	655539	655538	15/15/-
Thu Feb 7	7 18:26:30 2019				
firewall01	swp4.3(exit02)	default	655539	655538	15/15/-
Thu Feb 7	7 18:26:30 2019				
firewall01	swp4.4(exit02)	default	655539	655538	15/15/-
Thu Feb 7	7 18:26:30 2019				
spine01	swp10(exit02)	default	655435	655538	10/5/0
Thu Feb 7	7 18:19:50 2019				
spine01	swp10.2(exit02)	DataVrf1	080 6554	35 6555	538 10/5/0
Thu Feb 7	7 18:19:50 2019				
spine01	swp10.3(exit02)	DataVrf1	081 6554	35 6555	538 10/5/0
Thu Feb 7	7 18:19:50 2019				
spine01	swp10.4(exit02)	DataVrf1	082 6554	35 6555	338 10/5/0
Thu Feb 7	7 18:19:50 2019				
spine01	swp7(leaf01)	default	655435	655557	17/5/54
Thu Feb 7	7 18:19:50 2019				

spine01 swp7.2(leaf01)	DataVrf1080 655435 655557 14/2/0
Thu Feb 7 18:19:50 2019	
spine01 swp7.3(leaf01)	DataVrf1081 655435 655557 14/2/0
Thu Feb 7 18:19:50 2019	
spine01 swp7.4(leaf01)	DataVrf1082 655435 655557 14/2/0
Thu Feb 7 18:19:50 2019	
spine01 swp8(leaf02)	default 655435 655558 17/5/54
Thu Feb 7 18:19:50 2019	
spine01 swp8.2(leaf02)	DataVrf1080 655435 655558 14/2/0
Thu Feb 7 18:19:50 2019	
spine01 swp8.3(leaf02)	DataVrf1081 655435 655558 14/2/0
Thu Feb 7 18:19:50 2019	
spine01 swp8.4(leaf02)	DataVrf1082 655435 655558 14/2/0
Thu Feb 7 18:19:50 2019	
spine01 swp9(exit01)	default 655435 655537 19/5/0 Thu
Feb 7 18:19:50 2019	
spine01 swp9.2(exit01)	DataVrf1080 655435 655537 19/5/0
Thu Feb 7 18:19:50 2019	
spine01 swp9.3(exit01)	DataVrf1081 655435 655537 19/5/0
Thu Feb 7 18:19:50 2019	
spine01 swp9.4(exit01)	DataVrf1082 655435 655537 19/5/0
Thu Feb 7 18:19:50 2019	
spine02 swp10(exit02)	default 655435 655538 10/5/0
Thu Feb 7 18:19:50 2019	
spine02 swp10.3(exit02)	DataVrf1081 655435 655538 10/5/0
Thu Feb 7 18:19:50 2019	
spine02 swp10.4(exit02)	DataVrf1082 655435 655538 10/5/0
Thu Feb 7 18:19:50 2019	
spine02 swp7(leaf01)	default 655435 655557 17/5/62

Thu Feb 7 18:19:50 2019				
spine02 swp7.2(leaf01)	DataVrf1080	655435	655557	14/2/0
Thu Feb 7 18:19:50 2019				
spine02 swp7.3(leaf01)	DataVrf1081	655435	655557	14/2/0
Thu Feb 7 18:19:50 2019				
spine02 swp7.4(leaf01)	DataVrf1082	655435	655557	14/2/0
Thu Feb 7 18:19:50 2019				
spine02 swp8(leaf02)	default 65	5435 65	5558 17	/5/62
Thu Feb 7 18:19:50 2019				
spine02 swp8.2(leaf02)	DataVrf1080	655435	655558	14/2/0
Thu Feb 7 18:19:50 2019				
spine02 swp8.3(leaf02)	DataVrf1081	655435	655558	14/2/0
Thu Feb 7 18:19:50 2019				
spine02 swp8.4(leaf02)	DataVrf1082	655435	655558	14/2/0
Thu Feb 7 18:19:50 2019				
spine02 swp9(exit01)	default 65	5435 65	5537 19	/5/0 Thu
Feb 7 18:19:50 2019				
spine02 swp9.2(exit01)	DataVrf1080	655435	655537	19/5/0
Thu Feb 7 18:19:50 2019				
spine02 swp9.4(exit01)	DataVrf1082	655435	655537	19/5/0
Thu Feb 7 18:19:50 2019				
spine02 swp10.2(exit02)	DataVrf1080	655435	655538	10/5/0
Thu Feb 7 18:19:50 2019				
spine02 swp9.3(exit01)	DataVrf1081	655435	655537	19/5/0
Thu Feb 7 18:19:50 2019				
leaf01 swp3(spine01)	default 65	5557 65	5435 42	/27/324
Thu Feb 7 18:19:50 2019				
leaf01 swp3.2(spine01)	DataVrf1080	655557	655435	31/18/0
leaf01 swp3.2(spine01) Thu Feb 7 18:19:50 2019	DataVrf1080	655557	655435	31/18/0

leaf01	swp3.3(spine01)	DataVrf1081 655557 655435 31/18/0
Thu Feb 7	18:19:50 2019	
leaf01	swp3.4(spine01)	DataVrf1082 655557 655435 29/18/0
Thu Feb 7	18:19:50 2019	
leaf01	swp4(spine02)	default 655557 655435 42/27/324
Thu Feb 7	18:19:50 2019	
leaf01	swp4.2(spine02)	DataVrf1080 655557 655435 31/18/0
Thu Feb 7	18:19:50 2019	
leaf01	swp4.3(spine02)	DataVrf1081 655557 655435 31/18/0
Thu Feb 7	18:19:50 2019	
leaf01	swp4.4(spine02)	DataVrf1082 655557 655435 29/18/0
Thu Feb 7	18:19:50 2019	
leaf01	swp5(spine03)	default 655557 655435 42/27/324
Thu Feb 7	18:19:50 2019	
leaf01	swp5.2(spine03)	DataVrf1080 655557 655435 31/18/0
Thu Feb 7	18:19:50 2019	
leaf01	swp5.3(spine03)	DataVrf1081 655557 655435 31/18/0
Thu Feb 7	18:19:50 2019	
leaf01	swp5.4(spine03)	DataVrf1082 655557 655435 29/18/0
Thu Feb 7	18:19:50 2019	
leaf02	swp3(spine01)	default 655558 655435 42/27/372
Thu Feb 7	18:19:50 2019	
leaf02	swp3.2(spine01)	DataVrf1080 655558 655435 31/18/0
Thu Feb 7	18:19:50 2019	
leaf02	swp3.3(spine01)	DataVrf1081 655558 655435 31/18/0
Thu Feb 7	18:19:50 2019	
leaf02	swp3.4(spine01)	DataVrf1082 655558 655435 31/18/0
Thu Feb 7	18:19:50 2019	
leaf02	swp4(spine02)	default 655558 655435 42/27/372

Thu Feb 7	18:19:50 2019				
leaf02	swp4.2(spine02)	DataVrf1080	655558	655435	31/18/0
Thu Feb 7	18:19:50 2019				
leaf02	swp4.3(spine02)	DataVrf1081	655558	655435	31/18/0
Thu Feb 7	18:19:50 2019				
leaf02	swp4.4(spine02)	DataVrf1082	655558	655435	31/18/0
Thu Feb 7	18:19:50 2019				
leaf02	swp5(spine03)	default 65	5558 65	5435 42	/27/372
Thu Feb 7	18:19:50 2019				
leaf02	swp5.2(spine03)	DataVrf1080	655558	655435	31/18/0
Thu Feb 7	18:19:50 2019				
leaf02	swp5.3(spine03)	DataVrf1081	655558	655435	31/18/0
Thu Feb 7	18:19:50 2019				
leaf02	swp5.4(spine03)	DataVrf1082	655558	655435	31/18/0
Thu Feb 7	18:19:50 2019				
•••					

Example: View BGP Configuration Changes

This example shows that BGP configuration changes were made about five days ago on this network.

h:10m:11s
h:10m:11s
h:10m:11s

Validate BGP Operation

A single command enables you to validate that all configured route peering is established across the network. The command checks for duplicate router IDs and sessions that are in an unestablished state. Either of these conditions trigger a configuration check failure. When a failure is found, the reason is identified in the output along with the time the issue occurred.

This example shows a check on the BGP operations that found no failed sessions.

cumulus@switch:~\$ netq check bgp

Total Nodes: 15, Failed Nodes: 0, Total Sessions: 16, Failed Sessions: 0

This example shows 24 failed BGP sessions with a variety of reasons.

cumulus@switch:~\$ netq check bgp

Total Nodes: 25, Failed Nodes: 3, Total Sessions: 220, Failed Sessions: 24,

Hostname VRF Peer Name Peer Hostname

Reason Last Changed

exit-1 DataVrf1080 swp6.2 firewall-1 BGP session with peer

firewall-1 swp6.2: AFI/ 1d:7h:56m:9s

SAFI evpn not activated on peer

exit-1 DataVrf1080 swp7.2 firewall-2 BGP session with peer

firewall-2 (swp7.2 vrf 1d:7h:49m:31s

DataVrf1080) failed,

reason: Peer not configured

exit-1 DataVrf1081 swp6.3 firewall-1 BGP session with peer

firewall-1 swp6.3: AFI/ 1d:7h:56m:9s

SAFI evpn not activated on peer

exit-1 DataVrf1081 swp7.3 firewall-2 BGP session with peer

firewall-2 (swp7.3 vrf 1d:7h:49m:31s

DataVrf1081) failed,

reason: Peer not configured

exit-1 DataVrf1082 swp6.4 firewall-1 BGP session with peer

firewall-1 swp6.4: AFI/ 1d:7h:56m:9s

SAFI evpn not activated on peer

exit-1 DataVrf1082 swp7.4 firewall-2 BGP session with peer

firewall-2 (swp7.4 vrf 1d:7h:49m:31s

DataVrf1082) failed,

reason: Peer not configured

exit-1 default swp6 firewall-1 BGP session with peer

firewall-1 swp6: AFI/SA 1d:7h:56m:9s

FI evpn not activated on peer

exit-1 default swp7 firewall-2 BGP session with peer

firewall-2 (swp7 vrf de 1d:7h:49m:31s

fault) failed, reason: Peer not configured

exit-2 DataVrf1080 swp6.2 firewall-1 BGP session with peer

firewall-1 swp6.2: AFI/ 1d:7h:56m:9s

SAFI evpn not activated on peer

exit-2 DataVrf1080 swp7.2 firewall-2 BGP session with peer

firewall-2 (swp7.2 vrf 1d:7h:49m:26s

DataVrf1080) failed,

reason: Peer not configured

exit-2 DataVrf1081 swp6.3 firewall-1 BGP session with peer

firewall-1 swp6.3: AFI/ 1d:7h:56m:9s

SAFI evpn not activated on peer

. . .

Monitor OSPF Configuration

If you have OSPF running on your switches and hosts, you can monitor its operation using the NetQ CLI. For each device, you can view its associated interfaces, areas, peers, state, and type of OSPF running (numbered or unnumbered). Additionally, you can:

- view the information at an earlier point in time
- filter against a particular device, interface, or area
- validate it is operating correctly across the network

The **netq show ospf**command is used to obtain the OSPF configuration information from the devices. The **netq check ospf** command is used to validate the configuration. The syntax of these commands is:

netq [<hostname>] show ospf [<remote-interface>] [area <area-id>] [around <texttime>] [json]

netq [<hostname>] show events [level info|level error|level warning|level critical| level debug] type ospf [between <text-time> and <text-endtime>] [json] netq check ospf [around <text-time>] [json]

(i) NOTE

When entering a time value, you must include a numeric value and the unit of measure:

- d: day(s)
- h: hour(s)
- m: minute(s)
- s: second(s)
- now

For time ranges, the <text-time> is the most recent time and the <textendtime> is the oldest time. The values do not have to have the same unit of measure.

View OSPF Configuration Information

NetQ enables you to view the OSPF configuration of a single device or across all of your devices at once. You can filter the results based on a device, interface, or area. You can view the configuration in the past and view changes made to the configuration within a given timeframe.

Example: View OSPF Configuration Information Across the Network

This example shows all devices included in OSPF unnumbered routing, the assigned areas, state, peer and interface, and the last time this information was changed.

cumulus@switch:~\$ netq show ospf

Hostname	Interface	Area	Туре	State	Peer Hostname
Peer Interfa	ace Last Ch	anged			
	swp51			Full	spine01
swp1		14:42:16 20		- "	
leaf01	swp52			Full	spine02
·	Thu Feb 7				
leaf02	'			Full	spine01
-	Thu Feb 7				
	swp52			Full	spine02
swp2		14:42:16 20			
leaf03	swp51			Full	spine01
swp3	Thu Feb 7	14:42:16 20	19		
leaf03	swp52	0.0.0.0	Unnumbered	Full	spine02
swp3	Thu Feb 7	14:42:16 20	19		
leaf04	swp51	0.0.0.0	Unnumbered	Full	spine01
swp4	Thu Feb 7	14:42:16 20	19		
leaf04	swp52	0.0.0.0	Unnumbered	Full	spine02
swp4	Thu Feb 7	14:42:16 20	19		
spine01	swp1	0.0.0.0	Unnumbered	Full	leaf01
swp51	Thu Feb 7	7 14:42:16 20	019		
spine01	swp2	0.0.0.0	Unnumbered	Full	leaf02
swp51	Thu Feb 7	7 14:42:16 20	019		
spine01	swp3	0.0.0.0	Unnumbered	Full	leaf03
swp51	Thu Feb 7	7 14:42:16 20	019		
spine01	swp4	0.0.0.0	Unnumbered	Full	leaf04
swp51	Thu Feb 7	7 14:42:16 20	019		

spine02	swp1	0.0.0.0	Unnumbered	Full	leaf01
swp52	Thu Feb 7	7 14:42:16 2	019		
spine02	swp2	0.0.0.0	Unnumbered	Full	leaf02
swp52	Thu Feb 7	7 14:42:16 2	019		
spine02	swp3	0.0.0.0	Unnumbered	Full	leaf03
swp52	Thu Feb 7	7 14:42:16 2			
spine02	swp4	0.0.0.0	Unnumbered	Full	leaf04
swp52	Thu Feb 7	7 14:42:16 2	019		

Example: View OSPF Configuration Information for a Given Device

This example show the OSPF configuration information for *leaf01*.

cumulus@s\	witch:~\$ netq lea	f01 show o	spf		
Matching os	spf records:				
Hostname	Interface	Area	Туре	State	Peer Hostname
Peer Interfa	ce Last Cha	anged			
leaf01	swp51	0.0.0.0	Unnumbered	Full	spine01
swp1	Thu Feb 7	14:42:16 20	19		
leaf01	swp52	0.0.0.0	Unnumbered	Full	spine02
swp1	Thu Feb 7	14:42:16 20	19		

Example: View OSPF Configuration Information for a Given Interface

This example shows the OSPF configuration for all devices with the *swp51* interface.

cumulus@	switch:~\$ netq sh	ow ospf sw	p51		
Matching	ospf records:				
Hostname	Interface	Area	Туре	State	Peer Hostname
Peer Inter	face Last Ch	anged			
leaf01	swp51	0.0.0.0	Unnumbered	Full	spine01
swp1	Thu Feb 7	14:42:16 20	019		
leaf02	swp51	0.0.0.0	Unnumbered	Full	spine01
swp2	Thu Feb 7	14:42:16 20	019		
leaf03	swp51	0.0.0.0	Unnumbered	Full	spine01
swp3	Thu Feb 7	14:42:16 20	019		
leaf04	swp51	0.0.0.0	Unnumbered	Full	spine01
swp4	Thu Feb 7	14:42:16 20	019		
-					

Example: View OSPF Configuration Information at a Prior Time

This example shows the OSPF configuration for all leaf switches about five minutes ago.

swp1	Thu Feb 714	4:42:16 20	019		
leaf01	swp52	0.0.0.0	Unnumbered	Full	spine02
swp1	Thu Feb 714	4:42:16 20	019		
leaf02	swp51	0.0.0.0	Unnumbered	Full	spine01
swp2	Thu Feb 714	4:42:16 20	019		
leaf02	swp52	0.0.0.0	Unnumbered	Full	spine02
swp2	Thu Feb 714	4:42:16 20	019		
leaf03	swp51	0.0.0.0	Unnumbered	Full	spine01
swp3	Thu Feb 714	4:42:16 20	019		
leaf03	swp52	0.0.0.0	Unnumbered	Full	spine02
swp3	Thu Feb 714	4:42:16 20	019		
leaf04	swp51	0.0.0.0	Unnumbered	Full	spine01
swp4	Thu Feb 714	4:42:16 20	019		
leaf04	swp52	0.0.0.0	Unnumbered	Full	spine02
swp4	Thu Feb 714	4:42:16 20	019		

Validate OSPF Operation

A single command, **netq check ospf**, enables you to validate that all configured route peering is established across the network. The command checks for:

- router ID conflicts, such as duplicate IDs
- links that are down, or have mismatched MTUs
- mismatched session parameters (hello timer, dead timer, area ids, and network type)

When peer information is not available, the command verifies whether OSPF is configured on the peer and if so, whether the service is disabled, shutdown, or not functioning.

All of these conditions trigger a configuration check failure. When a failure is found, the reason is identified in the output along with the time the issue occurred.

This example shows a check on the OSPF operations that found no failed sessions.

cumulus@switch:~\$ netq check ospf

Total Sessions: 16, Failed Sessions: 0

This example shows a check on the OSPF operations that found two failed sessions. The results indicate the reason for the failure is a mismatched MTU for two links.

cumulus@switch:~\$ netq check ospf

Total Nodes: 21, Failed Nodes: 2, Total Sessions: 40, Failed Sessions: 2,

Hostname Interface PeerID Peer IP

Reason Last Changed

spine03 swp6 0.0.0.23 27.0.0.23 mtu mismatch,

mtu mismatch Thu Feb 7 14:42:16 2019

leaf22 swp5 0.0.0.17 27.0.0.17 mtu mismatch,

mtu mismatch Thu Feb 7 14:42:16 2019

View Paths between Devices

You can view the available paths between two devices on the network currently and at a time in the past using their IPv4 or IPv6 addresses. You can view the output in one of three formats (*json*, *pretty*, and *detail*). JSON output provides the output in a JSON file

format for ease of importing to other applications or software. Pretty output lines up the paths in a pseudo-graphical manner to help visualize multiple paths. Detail output is the default when not specified, and is useful for traces with higher hop counts where the pretty output wraps lines, making it harder to interpret the results. The detail output displays a table with a row per hop and a set of rows per path.

To view the paths, first identify the addresses for the source and destination devices using the netq show ip addresses command (see syntax above), and then use the netq trace command to see the available paths between those devices. The trace command syntax is:

netq trace <ip> from (<src-hostname>|<ip-src>) [vrf <vrf>] [around <text-time>] [json|detail|pretty] [debug]



(i) NOTE

The syntax requires the destination device address first, <*ip*>, and then the source device address or hostname.

The tracing function only knows about addresses that have already been learned. If you find that a path is invalid or incomplete, you may need to ping the identified device so that its address becomes known.

View Paths between Two Switches with Pretty Output

This example first determines the IP addresses of the leaf01 and leaf03 switches, then shows the available paths between them. The results include a summary of the trace,

including the total number of paths available, those with errors and warnings, and the MTU of the paths. In this case, the results are displayed in pseudo-graphical output.

	s records:			
Address	Hostname	Interface	VRF	Last Changed
10.0.0.11/32	leaf01	lo	default	Fri Feb 8 01:35:49
2019				
10.0.0.11/32	leaf01	swp51	default	Fri Feb 8 01:35:49
2019				
10.0.0.11/32	leaf01	swp52	default	Fri Feb 8 01:35:49
2019				
172.16.1.1/24	leaf01	br0	default	Fri Feb 8 01:35:49
2019				
192.168.0.11/24	leaf01	eth0	default	Fri Feb 8 01:35:49
2019				
cumulus@switch	:~\$ netq leaf0	3 show ip addr	esses	
Matching addres	s records:			
Address	Hostname	Interface	VRF	Last Changed
10.0.0.13/32	leaf03	lo	default	Thu Feb 7 18:31:29
2019				
10.0.0.13/32	leaf03	swp51	default	Thu Feb 7 18:31:29
2019				
10 0 0 12/22	leaf03	swp52	default	Thu Feb 7 18:31:29
10.0.0.13/32 2019				

2019

192.168.0.13/24 leaf03 eth0 default Thu Feb 7

18:31:29 2019

cumulus@switch:~\$ netq trace 10.0.0.13 from 10.0.0.11 pretty

Number of Paths: 2

Number of Paths with Errors: 0

Number of Paths with Warnings: 0

Path MTU: 1500

leaf01 swp52 -- swp1 spine02 swp3 -- swp52 leaf03 <lo>

swp51 -- swp1 spine01 swp3 -- swp51 leaf03 <lo>

View Paths between Two Switches with Detailed Output

This example provides the same path information as the pretty output, but displays the information in a tabular output. In this case there, no VLAN is configured, so the related fields are left blank.

cumulus@switch:~\$ netq trace 10.0.0.13 from 10.0.0.11 detail

Number of Paths: 2

Number of Paths with Errors: 0

Number of Paths with Warnings: 0

Path MTU: 1500

Id Hop Hostname InPort InVlan InTunnel InRtrIf InVRF

OutRtrIf OutVRF OutTunnel OutPort OutVlan

1 1 leaf01				swp52	default
	swp52				
2 spine02	swp1		swp1	default	swp3
default	9	wp3			
3 leaf03	swp52		swp52	default	lo
2 1 leaf01				swp51	default
	swp51				
2 spine01	swp1		swp1	default	swp3
default	9	wp3			
3 leaf03	swp51		swp51	default	lo

Monitor Virtual Network Overlays

With NetQ, a network administrator can monitor virtual network components in the data center, including VXLAN, EVPN, and LNV software constructs. NetQ provides the ability to:

- Manage virtual constructs: view the performance and status of VXLANs, EVPN,
 and LNV
- Validate overlay communication paths

It helps answer questions such as:

- Is my overlay configured and operating correctly?
- Is my control plane configured correctly?
- Can device A reach device B?

(i) NOTE

Lightweight network virtualization (LNV) was deprecated in Cumulus Linux 3.7.4 and will be removed in Cumulus Linux 4.0.0. Cumulus NetQ will continue to support and return LNV data as long as you are running a supported version of Cumulus Linux earlier than 4.0.0. For information on the support timeline, read this knowledge base article.

Monitor Virtual Extensible LANs

Virtual Extensible LANs (VXLANs) provide a way to create a virtual network on top of layer 2 and layer 3 technologies. It is intended for organizations, such as data centers, that require larger scale without additional infrastructure and more flexibility than is available with existing infrastructure equipment. With NetQ, you can monitor the current and historical configuration and status of your VXLANs using the following command:

netq [<hostname>] show vxlan [vni <text-vni>] [around <text-time>] [json]
netq show interfaces type vxlan [state <remote-interface-state>] [around <text-time>] [json]

netq <hostname> show interfaces type vxlan [state <remote-interface-state>]
[around <text-time>] [count] [json]

netq [<hostname>] show events [level info|level error|level warning|level critical| level debug] type vxlan [between <text-time> and <text-endtime>] [json]

(i) NOTE

When entering a time value, you must include a numeric value and the unit of measure:

- d: day(s)
- h: hour(s)
- m: minute(s)
- s: second(s)
- now

For time ranges, the <text-time> is the most recent time and the <textendtime> is the oldest time. The values do not have to have the same unit of measure.

View All VXLANs in Your Network

You can view a list of configured VXLANs for all devices, including the VNI (VXLAN network identifier), protocol, address of associated VTEPs (VXLAN tunnel endpoint), replication list, and the last time it was changed. You can also view VXLAN information for a given device by adding a hostname to the show command. You can filter the results by VNI.

This example shows all configured VXLANs across the network. In this network, there are three VNIs (13, 24, and 104001) associated with three VLANs (13, 24, 4001), EVPN is the virtual protocol deployed, and the configuration was last changed around 23 hours ago.

cumulus@switch:~\$ netq show vxlan Matching vxlan records: Hostname VNI Protoc VTEP IP VLAN Replication List Last Changed ol exit01 104001 EVPN 10.0.0.41 4001 Fri Feb 8 01:35:49 2019 exit02 104001 EVPN 10.0.0.42 Fri Feb 8 4001 01:35:49 2019 leaf01 13 EVPN 10.0.0.112 13 10.0.0.134(leaf04, leaf03) Fri Feb 8 01:35:49 2019 leaf01 24 EVPN 10.0.0.112 24 10.0.0.134(leaf04, leaf03) Fri Feb 8 01:35:49 2019 4001 leaf01 104001 EVPN 10.0.0.112 Fri Feb 8 01:35:49 2019 leaf02 13 EVPN 10.0.0.112 13 10.0.0.134(leaf04, leaf03) Fri Feb 8 01:35:49 2019 leaf02 24 EVPN 10.0.0.112 24 10.0.0.134(leaf04, leaf03) Fri Feb 8 01:35:49 2019 leaf02 104001 EVPN 10.0.0.112 4001 Fri Feb 8 01:35:49 2019 13 EVPN 10.0.0.134 13 10.0.0.112(leaf02, leaf01) leaf03 Fri Feb 8 01:35:49 2019 24 EVPN 10.0.0.134 24 10.0.0.112(leaf02, leaf01) leaf03 Fri Feb 8 01:35:49 2019 leaf03 104001 EVPN 10.0.0.134 4001 Fri Feb 8 01:35:49 2019 leaf04 13 EVPN 10.0.0.134 13 10.0.0.112(leaf02, leaf01) Fri

This example shows the events and configuration changes that have occurred on the VXLANs in your network in the last 24 hours. In this case, the EVPN configuration was added to each of the devices in the last 24 hours.

	vxlan recor	·	w everies cy	Je valali L	etween now and	Z-TI I	
Hostname			oc VTEP IP	VLAN	Replication List		DB
State Las	st Changed				1		
	ol						
exit02	104001	EVPN	10.0.0.42	4001		Add	Fri
Feb 8 01:	35:49 2019						
exit02	104001	EVPN	10.0.0.42	4001		Add	Fri
Feb 8 01:	35:49 2019						
exit02	104001	EVPN	10.0.0.42	4001		Add	Fri
Feb 8 01:	35:49 2019						
exit02	104001	EVPN	10.0.0.42	4001		Add	Fri
Feb 8 01:	35:49 2019						
exit02	104001	EVPN	10.0.0.42	4001		Add	Fri
Feb 8 01:	35:49 2019						
exit02	104001	EVPN	10.0.0.42	4001		Add	Fri

Feb 8 01:35:49 2019						
exit02	104001	EVPN	10.0.0.42	4001	Add	Fri
Feb 8 01:35:49 2019						
exit01	104001	EVPN	10.0.0.41	4001	Add	Fri
Feb 8 01:3	5:49 2019					
exit01	104001	EVPN	10.0.0.41	4001	Add	Fri
Feb 8 01:35:49 2019						
exit01	104001	EVPN	10.0.0.41	4001	Add	Fri
Feb 8 01:35:49 2019						
exit01	104001	EVPN	10.0.0.41	4001	Add	Fri
Feb 8 01:35:49 2019						
exit01	104001	EVPN	10.0.0.41	4001	Add	Fri
Feb 8 01:35:49 2019						
exit01	104001	EVPN	10.0.0.41	4001	Add	Fri
Feb 8 01:35:49 2019						
exit01	104001	EVPN	10.0.0.41	4001	Add	Fri
Feb 8 01:35:49 2019						
exit01	104001	EVPN	10.0.0.41	4001	Add	Fri
Feb 8 01:35:49 2019						
leaf04	104001	EVPN	10.0.0.134	4001	Add	Fri
Feb 8 01:35:49 2019						
leaf04	104001	EVPN	10.0.0.134	4001	Add	Fri
Feb 8 01:35:49 2019						
leaf04	104001	EVPN	10.0.0.134	4001	Add	Fri
Feb 8 01:35:49 2019						
leaf04	104001	EVPN	10.0.0.134	4001	Add	Fri
Feb 8 01:35:49 2019						
leaf04	104001	EVPN	10.0.0.134	4001	Add	Fri
Feb 8 01:35:49 2019						

leaf04	1040	01 EVP	N 10.0.0.134	4	001	Add	Fri
Feb 8 01:3	5:49 20	19					
leaf04	1040	01 EVP	N 10.0.0.134	4	001	Add	Fri
Feb 8 01:3	5:49 20	19					
leaf04	13	EVPN	10.0.0.134	13	10.0.0.112()	Add	Fri
Feb 8 01:3	5:49 20	19					
leaf04	13	EVPN	10.0.0.134	13	10.0.0.112()	Add	Fri
Feb 8 01:3	5:49 20	19					
leaf04	13	EVPN	10.0.0.134	13	10.0.0.112()	Add	Fri
Feb 8 01:3	5:49 20	19					
leaf04	13	EVPN	10.0.0.134	13	10.0.0.112()	Add	Fri
Feb 8 01:3	5:49 20	19					
leaf04	13	EVPN	10.0.0.134	13	10.0.0.112()	Add	Fri
Feb 8 01:3	5:49 20	19					
leaf04	13	EVPN	10.0.0.134	13	10.0.0.112()	Add	Fri
Feb 8 01:3	5:49 20	19					
leaf04	13	EVPN	10.0.0.134	13	10.0.0.112()	Add	Fri
Feb 8 01:3	5:49 20	19					
•••							

Consequently, if you looked for the VXLAN configuration and status for last week, you would find either another configuration or no configuration. This example shows that no VXLAN configuration was present.

cumulus@switch:~\$ netq show vxlan around 7d No matching vxlan records found

You can filter the list of VXLANs to view only those associated with a particular VNI. The VNI option lets you specify single VNI (100), a range of VNIs (10-100), or provide a comma-separated list (10,11,12). This example shows the configured VXLANs for *VNI 24*.

cumulus@switch:~\$ netq show vxlan vni 24 Matching vxlan records:								
Hostname	e V	NI Pr	otoc VTEP IP	V	LAN Replication List	Last		
Changed								
	0							
leaf01	24	EVPN	10.0.0.112	24	10.0.0.134(leaf04, leaf03)	Fri		
Feb 8 01:	35:49 20)19						
leaf02	24	EVPN	10.0.0.112	24	10.0.0.134(leaf04, leaf03)	Fri		
Feb 8 01:	35:49 20)19						
leaf03	24	EVPN	10.0.0.134	24	10.0.0.112(leaf02, leaf01)	Fri		
Feb 8 01:35:49 2019								
leaf04	24	EVPN	10.0.0.134	24	10.0.0.112(leaf02, leaf01)	Fri		
Feb 8 01:35:49 2019								

View the Interfaces Associated with VXLANs

You can view detailed information about the VXLAN interfaces using the **netq show interface** command. You can also view this information for a given device by adding a hostname to the **show** command. This example shows the detailed VXLAN interface information for the *leaf02* switch.

cumulus@switch:~\$ netq leaf02 show interfaces type vxlan Matching link records:

Hostname	Interface	Туре	9	State	VRF	
Details	Last	Changed				
leaf02	vni13	vxlan	up	defa	ult	VNI: 13, PVID: 13,
Master: br	idge, Fri Feb 80°	1:35:49 2019				
				VTEP	: 10.0.0	0.112, MTU: 9000
leaf02	vni24	vxlan	up	defa	ult	VNI: 24, PVID: 24,
Master: br	idge, Fri Feb 80°	1:35:49 2019				
				VTEP	: 10.0.0	0.112, MTU: 9000
leaf02	vxlan4001	vxlan	up	d	efault	VNI: 104001, PVID:
4001,	Fri Feb 8 01:35:	49 2019				
				Mast	er: brid	dge, VTEP: 10.0.0.112,
				MTU:	1500	

Monitor EVPN

EVPN (Ethernet Virtual Private Network) enables network administrators in the data center to deploy a virtual layer 2 bridge overlay on top of layer 3 IP networks creating access, or tunnel, between two locations. This connects devices in different layer 2 domains or sites running VXLANs and their associated underlays. With NetQ, you can monitor the configuration and status of the EVPN setup using the netq show evpn command. You can filter the EVPN information by a VNI (VXLAN network identifier), and view the current information or for a time in the past. The command also enables visibility into changes that have occurred in the configuration during a specific timeframe. The syntax for the command is:

netq [<hostname>] show evpn [vni <text-vni>] [mac-consistency] [around <texttime>] [ison]

netq [<hostname>] show events [level info|level error|level warning|level critical| level debug] type vxlan [between <text-time> and <text-endtime>] [json]

i NOTE

When entering a time value, you must include a numeric value *and* the unit of measure:

- d: day(s)
- h: hour(s)
- m: minute(s)
- s: second(s)
- now

For time ranges, the <text-time> is the most recent time and the <text-endtime> is the oldest time. The values do not have to have the same unit of measure.

For more information about and configuration of EVPN in your data center, refer to the Cumulus Linux EVPN topic.

View the Status of EVPN

You can view the configuration and status of your EVPN overlay across your network or for a particular device. This example shows the configuration and status for all devices, including the associated VNI, VTEP address, the import and export route (showing the

BGP ASN and VNI path), and the last time a change was made for each device running EVPN. Use the **hostname** option to view the configuration and status for a single device.

	vitcii.~‡	netq show e	vpn				
Matching ev	pn reco	rds:					
Hostname	VNI	VTEP IP]	in Kernel Export	RT	Import RT	Last
Changed							
leaf01	33	27.0.0.22	yes	197:33	197:33	Fri Feb	8
01:48:27 201	9						
leaf01	34	27.0.0.22	yes	197:34	197:34	Fri Feb	8
01:48:27 201	9						
leaf01	35	27.0.0.22	yes	197:35	197:35	Fri Feb	8
01:48:27 201	9						
leaf01	36	27.0.0.22	yes	197:36	197:36	Fri Feb	8
01:48:27 201	9						
leaf01	37	27.0.0.22	yes	197:37	197:37	Fri Feb	8
01:48:27 201	9						
leaf01	38	27.0.0.22	yes	197:38	197:38	Fri Feb	8
01:48:27 201	9						
leaf01	39	27.0.0.22	yes	197:39	197:39	Fri Feb	8
01:48:27 201	9						
leaf01	40	27.0.0.22	yes	197:40	197:40	Fri Feb	8
01:48:27 201	9						
leaf01	41	27.0.0.22	yes	197:41	197:41	Fri Feb	8
01:48:27 201	9						
leaf01	42	27.0.0.22	yes	197:42	197:42	Fri Feb	8
01:48:27 201	9						
leaf02	33	27.0.0.23	yes	198:33	198:33	Thu Fel	7

18:31:41 2	019					
leaf02	34	27.0.0.23	yes	198:34	198:34	Thu Feb 7
18:31:41 2	019					
leaf02	35	27.0.0.23	yes	198:35	198:35	Thu Feb 7
18:31:41 2	019					
leaf02	36	27.0.0.23	yes	198:36	198:36	Thu Feb 7
18:31:41 2	019					
leaf02	37	27.0.0.23	yes	198:37	198:37	Thu Feb 7
18:31:41 2	019					
leaf02	38	27.0.0.23	yes	198:38	198:38	Thu Feb 7
18:31:41 2	019					
leaf02	39	27.0.0.23	yes	198:39	198:39	Thu Feb 7
18:31:41 2	019					
leaf02	40	27.0.0.23	yes	198:40	198:40	Thu Feb 7
18:31:41 2	019					
leaf02	41	27.0.0.23	yes	198:41	198:41	Thu Feb 7
18:31:41 2	019					
leaf02	42	27.0.0.23	yes	198:42	198:42	Thu Feb 7
18:31:41 2	019					
•••						

View the Status of EVPN for a Given VNI

You can filter the full device view to focus on a single VNI. This example only shows the EVPN configuration and status for VNI 42.

cumulus@switch:~\$ netq show evpn vni 42

Matching evpn records:

Hostname VNI VTEP IP In Kernel Export RT Import RT Last

Changed						
leaf01	42	27.0.0.22	yes	197:42	197:42	Thu Feb 14
00:48:24 2	.019					
leaf02	42	27.0.0.23	yes	198:42	198:42	Wed Feb 13
18:14:49 2	.019					
leaf11	42	36.0.0.24	yes	199:42	199:42	Wed Feb 13
18:14:22 2	.019					
leaf12	42	36.0.0.24	yes	200:42	200:42	Wed Feb 13
18:14:27 2	.019					
leaf21	42	36.0.0.26	yes	201:42	201:42	Wed Feb 13
18:14:33 2	.019					
leaf22	42	36.0.0.26	yes	202:42	202:42	Wed Feb 13
18:14:37 2	.019					

View EVPN Events

You can view status and configuration change events for the EVPN protocol service using the **netq show events** command. This example shows the events that have occurred in the last 48 hours.

cumulus@	switch:/\$	neta sh	ow events type evpn between now and 48h					
Matching events records:								
Hostname	Mes	sage Ty	pe Severity Message Timestamp					
torc-21	evpn	info	VNI 33 state changed from down to u 1d:8h:16m:29s					
		р						
torc-12	evpn	info	VNI 41 state changed from down to u 1d:8h:16m:35s					
		р						

torc-11	evpn	info	VNI 39 state changed from down to u 1d:8h:16m:41s
		р	
tor-1	evpn	info	VNI 37 state changed from down to u 1d:8h:16m:47s
		р	
tor-2	evpn	info	VNI 42 state changed from down to u 1d:8h:16m:51s
		р	
torc-22	evpn	info	VNI 39 state changed from down to u 1d:8h:17m:40s
		р	
•••			
	evpn	info	VNI 39 state changed from down to u 1d:8h:17m:40s

Monitor LNV

Lightweight Network Virtualization (LNV) is a technique for deploying VXLANs without a central controller on bare metal switches. LNV enables data center network administrators and operators to create a data path between bridges on top of a layer 3 fabric. With NetQ, you can monitor the configuration and status of the LNV setup using the netq show lnv command. You can view the current information or for a time in the past. The command also enables visibility into changes that have occurred in the configuration during a specific timeframe. The syntax for the command is:

netq [<hostname>] show lnv [around <text-time>] [json]
netq [<hostname>] show events [level info|level error|level warning|level critical|
level debug] type lnv [between <text-time> and <text-endtime>] [json]

View LNV Status

You can view the configuration and status of your LNV overlay across your network or for a particular device. This example shows the configuration and status of LNV across

the network, including the role each node plays, replication mode, number of peers and VNIs, and the last time the configuration was changed.

cumulus@s	cumulus@switch:~\$ netq show lnv							
Matching Ll	NV sessi	ion reco	rds are	•				
Hostname	Rol	e Re _l	plMode	Stat	е	#Peers #VNIs Last Changed		
spine01	SND	HER	up	3	6	Thu Feb 7 18:31:31 2019		
spine02	SND	HER	up	3	6	Thu Feb 7 18:31:31 2019		
spine03	SND	HER	up	3	6	Thu Feb 7 18:31:31 2019		
leaf01	RD	HER	up	4	6	Thu Feb 7 18:31:31 2019		
leaf02	RD	HER	up	4	6	Thu Feb 7 18:31:31 2019		
leaf11	RD	HER	up	0	0	Thu Feb 7 18:31:31 2019		
leaf12	RD	HER	up	4	6	Thu Feb 7 18:31:31 2019		
leaf21	RD	HER	up	4	6	Thu Feb 7 18:31:31 2019		
leaf22	RD	HER	up	4	6	Thu Feb 7 18:31:31 2019		

View LNV Status in the Past

You can view the status in the past using the **around** option. This example shows the status of LNV about 30 minutes ago.

cumulus@switch:~\$ netq show lnv around 30m							
Matching LNV session records are:							
Hostname	Role	Repl	Mode	State	#	Peers #VNIs Last Changed	
spine01	SND	HER	up	3	6	Thu Feb 7 18:31:31 2019	
spine02	SND	HER	up	3	6	Thu Feb 7 18:31:31 2019	

spine03	SND	HEF	R up	3	3 6	5 Thu Feb 7 18:31:31 2019
leaf01	RD	HER	up	4	6	Thu Feb 7 18:31:31 2019
leaf02	RD	HER	up	4	6	Thu Feb 7 18:31:31 2019
leaf11	RD	HER	up	4	6	Thu Feb 7 18:31:31 2019
leaf12	RD	HER	up	4	6	Thu Feb 7 18:31:31 2019
leaf21	RD	HER	up	4	6	Thu Feb 7 18:31:31 2019
leaf22	RD	HER	up	4	6	Thu Feb 7 18:31:31 2019

For more information about and configuration of LNV, refer to the Cumulus Linux LNV Overview topic.

Monitor Linux Hosts

Running NetQ on Linux hosts provides unprecedented network visibility, giving the network operator a complete view of the entire infrastructure's network connectivity instead of just from the network devices.

The NetQ Agent is supported on the following Linux hosts:

- CentOS 7
- Red Hat Enterprise Linux 7.1
- Ubuntu 16.04

You need to install the OS-specific NetQ metapack on every host you want to monitor with NetQ.

The NetQ Agent monitors the following on Linux hosts:

- netlink
- Layer 2: LLDP and VLAN-aware bridge
- Layer 3: IPv4, IPv6
- Routing on the Host: BGP, OSPF
- systemctl for services
- Docker containers refer to the Monitor Container Environments topic

Using NetQ on a Linux host is the same as using it on a Cumulus Linux switch. For example, if you want to check LLDP neighbor information about a given host, run:

cumulus@switch:~\$ netq server01 show lldp Matching lldp records:

Monitor Linux Hosts

Hostname Changed	Interface	Peer I	Hostname Peer	Interface Last
server01 01:50:59 201	eth0 19	oob-mgn	nt-switch swp2	Fri Feb 8
server01 2019	eth1	leaf01	swp1	Fri Feb 8 01:50:59
server01 2019	eth2	leaf02	swp1	Fri Feb 8 01:50:59

Then, to see LLDP from the switch's perspective:

cumulus@switch:~\$ netq leaf01 show lldp Matching lldp records:						
Hostname	Interface	Peer Ho	ostname	Peer Interface Last		
Changed						
leaf01	eth0	oob-mgmt-s	switch sw	p6 Thu Feb 7		
18:31:26 20°	19					
leaf01	swp1	server01	eth1	Thu Feb 7 18:31:26		
2019						
leaf01	swp2	server02	eth1	Thu Feb 7 18:31:26		
2019						
leaf01	swp49	leaf02	swp49	Thu Feb 7 18:31:26		
2019						
leaf01	swp50	leaf02	swp50	Thu Feb 7 18:31:26		
2019						
leaf01	swp51	spine01	swp1	Thu Feb 7 18:31:26		

Monitor Linux Hosts

2019				
leaf01	swp52	spine02	swp1	Thu Feb 7 18:31:26
2019				

To get the routing table for a server:

Matching rout					
Origin VRF	igin VRF Prefix Hostname		Nexthops	Last	
Changed					
					
	10.2.4.0/24	server01	10.1.3.1: uplink	Fri	
Feb 8 01:50:4					
no default	172.16.1.0/24	server01	10.1.3.1: uplink		
Fri Feb 8 01:5	0:49 2019				
yes default	10.1.3.0/24	server01	uplink	Fri	
Feb 8 01:50:4	9 2019				
yes default	10.1.3.101/32	server01	uplink	Fri	
Feb 8 01:50:4	9 2019				
yes default	192.168.0.0/24	server01	eth0	Fri	
Feb 8 01:50:49 2019					
yes default	192.168.0.31/32	server01	eth0	Fri	

Monitor Container Environments

The NetQ Agent monitors container environments the same way it monitors physical servers. There is no special implementation. The NetQ Agent pulls data from the container as it would pull data from a Cumulus Linux switch or Linux host. It can be installed on a Linux server or in a Linux VM. NetQ Agent integrates with the Kubernetes container orchestrator.

NetQ monitors many aspects of containers on your network, including their:

- Identity: The NetQ agent tracks every container's IP and MAC address, name, image, and more. NetQ can locate containers across the fabric based on a container's name, image, IP or MAC address, and protocol and port pair.
- Port mapping on a network: The NetQ agent tracks protocol and ports exposed by a container. NetQ can identify containers exposing a specific protocol and port pair on a network.
- **Connectivity**: NetQ can provide information on network connectivity for a container, including adjacency, and can identify containers that can be affected by a top of rack switch.

NetQ helps answer questions such as:

- Where is this container located?
- Open ports? What image is being used?
- Which containers are part of this service? How are they connected?

Use NetQ with Kubernetes Clusters

The NetQ Agent interfaces with a Kubernetes API server and listens to Kubernetes events. The NetQ Agent monitors network identity and physical network connectivity of Kubernetes resources like Pods, Daemon sets, Service, and so forth. NetQ works with any container network interface (CNI), such as Calico or Flannel.

The NetQ Kubernetes integration enables network administrators to:

- Identify and locate pods, deployment, replica-set and services deployed within the network using IP, name, label, and so forth.
- Track network connectivity of all pods of a service, deployment and replica set.
- Locate what pods have been deployed adjacent to a top of rack (ToR) switch.
- Check what pod, services, replica set or deployment can be impacted by a specific ToR switch.

NetQ also helps network administrators identify changes within a Kubernetes cluster and determine if such changes had an adverse effect on the network performance (caused by a noisy neighbor for example). Additionally, NetQ helps the infrastructure administrator determine how Kubernetes workloads are distributed within a network.

Requirements

The NetQ Agent supports Kubernetes version 1.9.2 or later.

Due to the higher memory requirements to run containers, Cumulus Networks recommends you run the NetQ Platform on a host with at least 64G RAM.

Command Summary

There is a large set of commands available to monitor Kubernetes configurations, including the ability to monitor clusters, nodes, daemon-set, deployment, pods, replication, and services. Run netq show

kubernetes help to see all the possible commands:

netq [<hostname>] show kubernetes cluster [name <kube-cluster-name>] [around <text-time>] [json] netq [<hostname>] show kubernetes node [components] [name <kube-nodename>] [cluster <kube-cluster-name>] [label <kube-node-label>] [around <texttime>] [json] netq [<hostname>] show kubernetes daemon-set [name <kube-ds-name>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kube-ds-label>] [around <text-time>] [json] netq [<hostname>] show kubernetes daemon-set [name <kube-ds-name>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kube-ds-label>] connectivity [around <text-time>] [json] netq [<hostname>] show kubernetes deployment [name <kube-deploymentname>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kubedeployment-label>] [around <text-time>] [json] netq [<hostname>] show kubernetes deployment [name <kube-deploymentname>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kubedeployment-label>] connectivity [around <text-time>] [json] netq [<hostname>] show kubernetes pod [name <kube-pod-name>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kube-pod-label>] [podip <kube-pod-ipaddress>] [node <kube-node-name>] [around <text-time>] [json] netq [<hostname>] show kubernetes replication-controller [name <kube-rcname>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kube-rclabel>] [around <text-time>] [json] netq [<hostname>] show kubernetes replica-set [name <kube-rs-name>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kube-rs-label>] [around <text-time>] [json] netq [<hostname>] show kubernetes replica-set [name <kube-rs-name>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kube-rs-label>] connectivity [around <text-time>] [json]

netq [<hostname>] show kubernetes service [name <kube-service-name>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kube-service-label>] [service-cluster-ip <kube-service-cluster-ip>] [service-external-ip <kube-serviceexternal-ip>] [around <text-time>] [json] netq [<hostname>] show kubernetes service [name <kube-service-name>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kube-service-label>] [service-cluster-ip <kube-service-cluster-ip>] [service-external-ip <kube-serviceexternal-ip>] connectivity [around <text-time>] [json] netq <hostname> show impact kubernetes service [master <kube-master-node>] [name <kube-service-name>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kube-service-label>] [service-cluster-ip <kube-servicecluster-ip>] [service-external-ip <kube-service-external-ip>] [around <text-time>] [ison] netq <hostname> show impact kubernetes replica-set [master <kube-masternode>] [name <kube-rs-name>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kube-rs-label>] [around <text-time>] [json] netq <hostname> show impact kubernetes deployment [master <kube-masternode>] [name <kube-deployment-name>] [cluster <kube-cluster-name>] [namespace <namespace>] [label <kube-deployment-label>] [around <text-time>] [ison] netg config add agent kubernetes-monitor [poll-period <text-duration-period>] netg config del agent kubernetes-monitor netq config show agent kubernetes-monitor [json]

Enable Kubernetes Monitoring

For NetQ to monitor the containers on a host, you must configure the following on the Kubernetes master node:

- Configure the host to point to the NetQ Platform by its IP address. See the Install NetQ topic for details.
- 2. Enable Kubernetes monitoring by NetQ. You can specify a polling period between 10 and 120 seconds; 15 seconds is the default.

cumulus@host:~\$ netq config add agent kubernetes-monitor poll-period 20 Successfully added kubernetes monitor. Please restart netq-agent.

3. Restart the NetQ agent:

cumulus@server01:~\$ netq config restart agent

Next, you must enable the NetQ Agent on all the worker nodes, as described in the Install NetQ Agents and CLI on Switches, for complete insight into your container network.

View Status of Kubernetes Clusters

You can get the status of all Kubernetes clusters in the fabric using the **netq show** kubernetes cluster command:

cumulus@switch:~\$ netq show kubernetes cluster

Matching kube_cluster records:

Master Cluster Name Controller Status Scheduler Status Nodes

server11:3.0.0.68	default	Healthy	Healthy	server11 server13
se				
			rver22 server11	serv
			er12 server23 se	erver
			24	
server12:3.0.0.69	default	Healthy	Healthy	server12 server21
se				
			rver23 server13	serv
			er14 server21 se	erver
			22	

To filter the list, you can specify the hostname of the master before the show command:

cumulus@switch:~\$ netq server11 show kubernetes cluster					
Matching kube_cluster records:					
Master	Cluster Name	Controlle	er Status	Schedu	ler Status Nodes
server11:3.0.0.68	3 default	Healthy	Не	althy	server11 server13
se					
			rver22 s	erver11	serv
			er12 ser	ver23 se	rver
			24		

Optionally, you can output the results in JSON format:

```
cumulus@server11:~$ netq show kubernetes cluster json
  "kube_cluster":[
      "clusterName":"default",
      "schedulerStatus":"Healthy",
      "master":"server12:3.0.0.69",
      "nodes": "server12 server21 server23 server13 server14 server21 server22",
      "controllerStatus":"Healthy"
    },
    {
      "clusterName":"default",
      "schedulerStatus":"Healthy",
      "master":"server11:3.0.0.68",
      "nodes": "server11 server13 server22 server11 server12 server23 server24",
      "controllerStatus":"Healthy"
  }
  1,
  "truncatedResult":false
}
```

View Changes to a Cluster

If data collection from the NetQ Agents is not occurring as it once was, you can verify that no changes have been made to the Kubernetes cluster configuration using the around option. This example shows the changes that have been made in the last hour.

cumulus@server11:~\$ netq show kubernetes cluster around 1h Matching kube cluster records: Cluster Name Controller Status Scheduler Status Master Nodes DBState Last changed server11:3.0.0.68 default Healthy Healthy server11 server13 server22 server11 serv Add Fri Feb 8 01:50:50 2019 er12 server23 server24 server12:3.0.0.69 default Healthy Healthy server12 server21 server23 server13 serv Add Fri Feb 8 01:50:50 2019 er14 server21 server22 Healthy server12 server21 server12:3.0.0.69 default Healthy server23 server13 Add Fri Feb 8 01:50:50 2019 server11:3.0.0.68 default Healthy Healthy Add Fri Feb 8 01:50:50 2019 server11 server12:3.0.0.69 default Healthy Healthy server12 Add Fri Feb 8 01:50:50 2019

View Kubernetes Pod Information

You can show configuration and status of the pods in a cluster, including the names, labels, addresses, associated cluster and containers, and whether the pod is running. This example shows pods for FRR, Nginx, Calico, various Kubernetes components sorted by master node.

cumulus@server11:~\$ netq show kubernetes pod Matching kube_pod records: Master Namespace Name IP Node Labels

Status Containers Last Changed

server11:3.0.0.68 default cumulus-frr-8vssx 3.0.0.70 server13 podtemplate-generat Running cumulus-frr:f8cac70bb217 Fri Feb 8 01:50:50 2019

ion:1 name:cumulus-f

rr controller-revisi

on-hash:3710533951

server11:3.0.0.68 default cumulus-frr-dkkgp 3.0.5.135 server24 podtemplate-generat Running cumulus-frr:577a60d5f40c Fri Feb 8 01:50:50 2019

ion:1 name:cumulus-f

rr controller-revisi

on-hash:3710533951

server11:3.0.0.68 default cumulus-frr-f4bgx 3.0.3.196 server11 podtemplate-generat Running cumulus-frr:1bc73154a9f5 Fri Feb 8 01:50:50 2019

ion:1 name:cumulus-f

rr controller-revisi

on-hash:3710533951

server11:3.0.0.68 default cumulus-frr-gqqxn 3.0.2.5 server22 podtemplate-generat Running cumulus-frr:3ee0396d126a Fri Feb 8 01:50:50 2019

ion:1 name:cumulus-f

rr controller-revisi

on-hash:3710533951

server11:3.0.0.68 default cumulus-frr-kdh9f 3.0.3.197 server12 podtemplate-generat Running cumulus-frr:94b6329ecb50 Fri Feb 8 01:50:50 2019

ion:1 name:cumulus-f

rr controller-revisi

on-hash:3710533951

server11:3.0.0.68 default cumulus-frr-mvv8m 3.0.5.134 server23 pod-template-generat Running cumulus-frr:b5845299ce3c Fri Feb 8 01:50:50 2019

ion:1 name:cumulus-f

rr controller-revisi

on-hash:3710533951

server11:3.0.0.68 default httpd-5456469bfd-bq9 10.244.49.65 server22 app:httpd Running httpd:79b7f532be2d Fri Feb 8 01:50:50 2019

zm

server11:3.0.0.68 default influxdb-6cdb566dd-8 10.244.162.128 server13 app:influx Running influxdb:15dce703cdec Fri Feb 8 01:50:50 2019

9lwn

server11:3.0.0.68 default nginx-8586cf59-26pj5 10.244.9.193 server24 Running nginx:6e2b65070c86 Fri Feb 8 01:50:50 2019 run:nginx server11:3.0.0.68 default nginx-8586cf59-c82ns 10.244.40.128 server12 Running nginx:01b017c26725 Fri Feb 8 01:50:50 2019 run:nginx nginx-8586cf59-wjwgp 10.244.49.64 server22 server11:3.0.0.68 default Running nginx:ed2b4254e328 Fri Feb 8 01:50:50 2019 run:nginx kube-system calico-etcd-pfg9r 3.0.0.68 server11:3.0.0.68 app:calico-etcd Running calico-etcd:f95f44b745a7 Fri Feb 8 01:50:50 2019

pod-template-generat

ion:1 controller-rev

ision-hash:142071906

5

server11:3.0.0.68 kube-system calico-kube-controll 3.0.2.5 server22 k8s-app:calico-kube- Running calico-kube-controllers: Fri Feb 8 01:50:50 2019 ers-d669cc78f-4r5t2 controllers

3688b0c5e9c5

server11:3.0.0.68 kube-system calico-node-4px69 3.0.2.5 server22

k8s-app:calico-node Running calico-node:1d01648ebba4 Fri Feb 8 01:50:50 2019

pod-template-generat

install-cni:da350802a3d2

ion:1 controller-rev

ision-hash:324404111

9

server11:3.0.0.68 kube-system calico-node-bt8w6 3.0.3.196 server11

k8s-app:calico-node Running calico-node:9b3358a07e5e Fri Feb 8 01:50:50 2019

pod-template-generat

install-cni:d38713e6fdd8

ion:1 controller-rev

ision-hash:324404111

9

server11:3.0.0.68 kube-system calico-node-gtmkv 3.0.3.197 server12

k8s-app:calico-node Running calico-node:48fcc6c40a6b Fri Feb 8 01:50:50 2019

pod-template-generat

install-cni:f0838a313eff

ion:1 controller-rev

ision-hash:324404111

9

server11:3.0.0.68 kube-system calico-node-mvslq 3.0.5.134 server23

k8s-app:calico-node Running calico-node:7b361aece76c Fri Feb 8 01:50:50 2019

pod-template-generat

install-cni:f2da6bc36bf8

ion:1 controller-rev

ision-hash:324404111

9

server11:3.0.0.68 kube-system calico-node-sjj2s 3.0.5.135 server24

k8s-app:calico-node Running calico-node:6e13b2b73031 Fri Feb 8 01:50:50 2019

pod-template-generat

install-cni:fa4b2b17fba9

ion:1 controller-rev

ision-hash:324404111

9

server11:3.0.0.68 kube-system calico-node-vdkk5 3.0.0.70 server13

k8s-app:calico-node Running calico-node:fb3ec9429281 Fri Feb 8 01:50:50 2019

pod-template-generat

install-cni:b56980da7294

ion:1 controller-rev

ision-hash:324404111

9

server11:3.0.0.68 kube-system calico-node-zzfkr 3.0.0.68 server11 k8s-app:calico-node Running calico-node:c1ac399dd862 Fri Feb 8 01:50:50 2019

pod-template-generat

install-cni:60a779fdc47a

ion:1 controller-rev

ision-hash:324404111

9

server11:3.0.0.68 kube-system etcd-server11 3.0.0.68 server11

tier:control-plane c Running etcd:dde63d44a2f5 Fri Feb 8 01:50:50 2019

omponent:etcd

server11:3.0.0.68 kube-system kube-apiserver-hostd 3.0.0.68 server11

tier:control-plane c Running kube-apiserver:0cd557bbf Fri Feb 8 01:50:50 2019

-11 omponent:kube-apiser 2fe

ver

server11:3.0.0.68 kube-system kube-controller-mana 3.0.0.68 server11

tier:control-plane c Running kube-controller-manager: Fri Feb 8 01:50:50 2019

ger-server11 omponent:kube-

contro 89b2323d09b2

ller-manager

server11:3.0.0.68 kube-system kube-dns-6f4fd4bdf-p 10.244.34.64 server23 k8s-app:kube-dns Running dnsmasq:284d9d363999 kub Fri Feb 8 01:50:50 2019

lv7p

edns:bd8bdc49b950 sideca

r:fe10820ffb19

server11:3.0.0.68 kube-system kube-proxy-4cx2t 3.0.3.197 server12 k8s-app:kube-proxy p Running kube-proxy:49b0936a4212 Fri Feb 8 01:50:50 2019

od-template-generati

on:1 controller-revi

sion-hash:3953509896

server11:3.0.0.68 kube-system kube-proxy-7674k 3.0.3.196 server11 k8s-app:kube-proxy p Running kube-proxy:5dc2f5fe0fad Fri Feb 8 01:50:50 2019

od-template-generati

on:1 controller-revi

sion-hash:3953509896

server11:3.0.0.68 kube-system kube-proxy-ck5cn 3.0.2.5 server22 k8s-app:kube-proxy p Running kube-proxy:6944f7ff8c18 Fri Feb 8 01:50:50 2019

od-template-generati

on:1 controller-revi

sion-hash:3953509896

server11:3.0.0.68 kube-system kube-proxy-f9dt8 3.0.0.68 server11 k8s-app:kube-proxy p Running kube-proxy:032cc82ef3f8 Fri Feb 8 01:50:50 2019

od-template-generati

on:1 controller-revi

sion-hash:3953509896

server11:3.0.0.68 kube-system kube-proxy-j6qw6 3.0.5.135 server24

k8s-app:kube-proxy p Running kube-proxy:10544e43212e Fri Feb 8 01:50:50 2019

od-template-generati

on:1 controller-revi

sion-hash:3953509896

server11:3.0.0.68 kube-system kube-proxy-lq8zz 3.0.5.134 server23

k8s-app:kube-proxy p Running kube-proxy:1bcfa09bb186 Fri Feb 8 01:50:50 2019

od-template-generati

on:1 controller-revi

sion-hash:3953509896

server11:3.0.0.68 kube-system kube-proxy-vg7kj 3.0.0.70 server13

k8s-app:kube-proxy p Running kube-proxy:8fed384b68e5 Fri Feb 8 01:50:50 2019

od-template-generati

on:1 controller-revi

sion-hash:3953509896

server11:3.0.0.68 kube-system kube-scheduler-hostd 3.0.0.68 server11

tier:control-plane c Running kube-scheduler:c262a8071 Fri Feb 8 01:50:50 2019

-11 omponent:kube-schedu

3cb

ler

server12:3.0.0.69 default cumulus-frr-2gkdv 3.0.2.4 server21 podtemplate-generat Running cumulus-frr:25d1109f8898 Fri Feb 8 01:50:50 2019

ion:1 name:cumulus-f

rr controller-revisi

on-hash:3710533951

server12:3.0.0.69 default cumulus-frr-b9dm5 3.0.3.199 server14

pod-template-generat Running cumulus-frr:45063f9a095f Fri Feb 8 01:50:50 2019

ion:1 name:cumulus-f

rr controller-revisi

on-hash:3710533951

server12:3.0.0.69 default cumulus-frr-rtqhv 3.0.2.6 server23 podtemplate-generat Running cumulus-frr:63e802a52ea2 Fri Feb 8 01:50:50 2019

ion:1 name:cumulus-f

rr controller-revisi

on-hash:3710533951

server12:3.0.0.69 default cumulus-frr-tddrg 3.0.5.133 server22 podtemplate-generat Running cumulus-frr:52dd54e4ac9f Fri Feb 8 01:50:50 2019

ion:1 name:cumulus-f

rr controller-revisi

on-hash:3710533951

server12:3.0.0.69 default cumulus-frr-vx7jp 3.0.5.132 server21 podtemplate-generat Running cumulus-frr:1c20addfcbd3 Fri Feb 8 01:50:50 2019

ion:1 name:cumulus-f

rr controller-revisi

on-hash:3710533951

server12:3.0.0.69 default cumulus-frr-x7ft5 3.0.3.198 server13 podtemplate-generat Running cumulus-frr:b0f63792732e Fri Feb 8 01:50:50 2019

ion:1 name:cumulus-f

rr controller-revisi

on-hash:3710533951

server12:3.0.0.69 kube-system calico-etcd-btqgt 3.0.0.69 server12 k8s-app:calico-etcd Running calico-etcd:72b1a16968fb Fri Feb 8 01:50:50 2019

pod-template-generat

ion:1 controller-rev

ision-hash:142071906

5

server12:3.0.0.69 kube-system calico-kube-controll 3.0.5.132 server21 k8s-app:calico-kube- Running calico-kube-controllers: Fri Feb 8 01:50:50 2019

ers-d669cc78f-bdnzk controllers

6821bf04696f

server12:3.0.0.69 kube-system calico-node-4g6vd 3.0.3.198 server13

k8s-app:calico-node Running calico-node:1046b559a50c Fri Feb 8 01:50:50 2019

pod-template-generat

install-cni:0a136851da17

ion:1 controller-rev

ision-hash:490828062

server12:3.0.0.69 kube-system calico-node-4hg6l 3.0.0.69 server12

k8s-app:calico-node Running calico-node:4e7acc83f8e8 Fri Feb 8 01:50:50 2019

pod-template-generat

install-cni:a26e76de289e

ion:1 controller-rev

ision-hash:490828062

server12:3.0.0.69 kube-system calico-node-4p66v 3.0.2.6 server23

k8s-app:calico-node Running calico-node:a7a44072e4e2 Fri Feb 8 01:50:50 2019

pod-template-generat

install-cni:9a19da2b2308

ion:1 controller-rev

ision-hash:490828062

server12:3.0.0.69 kube-system calico-node-5z7k4 3.0.5.133 server22

k8s-app:calico-node Running calico-node:9878b0606158 Fri Feb 8 01:50:50 2019

pod-template-generat

install-cni:489f8f326cf9

ion:1 controller-rev

ision-hash:490828062

• • •

You can filter this information to focus on a particular pod:

cumulus@server11:~\$ netq show kubernetes pod node server11 Matching kube pod records: Labels Master Namespace Name IP Node Status Containers Last Changed server11:3.0.0.68 kube-system calico-etcd-pfq9r 3.0.0.68 server11 k8sapp:calico-etcd Running calico-etcd:f95f44b745a7 2d:14h:0m:59s pod-template-generat ion:1 controller-rev ision-hash:142071906 server11:3.0.0.68 kube-system calico-node-zzfkr 3.0.0.68 server11 k8sapp:calico-node Running calico-node:c1ac399dd862 2d:14h:0m:59s pod-template-generat install-cni:60a779fdc47a ion:1 controller-rev ision-hash:324404111 9 server11:3.0.0.68 kube-system etcd-server11 3.0.0.68 server11 tier:control-plane c Running etcd:dde63d44a2f5 2d:14h:1m:44s omponent:etcd server11:3.0.0.68 kube-system kube-apiserver-serve 3.0.0.68 server11 tier:control-plane c Running kube-apiserver:0cd557bbf 2d:14h:1m:44s r11 omponent:kube-apiser 2fe ver server11:3.0.0.68 kube-system kube-controller-mana 3.0.0.68 server11 tier:control-plane c Running kube-controller-manager: 2d:14h:1m:44s ger-server11 omponent:kubecontro 89b2323d09b2

ller-manager

server11:3.0.0.68 kube-system kube-proxy-f9dt8 3.0.0.68 server11

k8s-app:kube-proxy p Running kube-proxy:032cc82ef3f8 2d:14h:0m:59s

od-template-generati

on:1 controller-revi

sion-hash:3953509896

server11:3.0.0.68 kube-system kube-scheduler-serve 3.0.0.68 server11

tier:control-plane c Running kube-scheduler:c262a8071 2d:14h:1m:44s

r11 omponent:kube-schedu

3cb

ler

View Kubernetes Node Information

You can view a lot of information about a node, including the pod CIDR and kubelet status.

cumulus@host:~\$ netq server11 show kubernetes node

Matching kube_cluster records:

Master Cluster Name Node Name Role Status

Labels Pod CIDR Last Changed

.....

server11:3.0.0.68 default server11 master KubeletReady node-

role.kubernetes 10.224.0.0/24 14h:23m:46s

.io/master: kubernet

es.io/hostname:hostd

-11 beta.kubernetes.

io/arch:amd64 beta.k

ubernetes.io/os:linu

Χ

server11:3.0.0.68 default server13 worker KubeletReady

kubernetes.io/hostna 10.224.3.0/24 14h:19m:56s

me:server13 beta.kub

ernetes.io/arch:amd6

4 beta.kubernetes.io

/os:linux

server11:3.0.0.68 default server22 worker KubeletReady

kubernetes.io/hostna 10.224.1.0/24 14h:24m:31s

me:server22 beta.kub

ernetes.io/arch:amd6

4 beta.kubernetes.io

/os:linux

server11:3.0.0.68 default server11 worker KubeletReady

kubernetes.io/hostna 10.224.2.0/24 14h:24m:16s

me:server11 beta.kub

ernetes.io/arch:amd6

4 beta.kubernetes.io

/os:linux

server11:3.0.0.68 default server12 worker KubeletReady

kubernetes.io/hostna 10.224.4.0/24 14h:24m:16s

me:server12 beta.kub

ernetes.io/arch:amd6

4 beta.kubernetes.io

/os:linux

server11:3.0.0.68 default server23 worker KubeletReady

kubernetes.io/hostna 10.224.5.0/24 14h:24m:16s

			me:server23 beta.kub	
			ernetes.io/arch:amd6	
			4 beta.kubernetes.io	
			/os:linux	
server11:3.0.0.68	default	server24	worker KubeletReady	
kubernetes.io/hostna 10.224.6.0/24 14h:24m:1s				
			me:server24 beta.kub	
			ernetes.io/arch:amd6	
			4 beta.kubernetes.io	
			/os:linux	

To display the kubelet or Docker version, use the **components** option on the above command. This example lists all the details of all master and worker nodes because the master's hostname — *server11* in this case — was included in the query.

Matching kube_clus	ster records:				
Maste	er Clus	ter Name	Node Name	Kubelet	
KubeProxy Cor	ntainer Runt				
				ime	
server11:3.0.0.68	default	server11	v1.9.2	v1.9.2	docker://
17.3.2 KubeletRea	dy				
server11:3.0.0.68	default	server13	v1.9.2	v1.9.2	docker://
17.3.2 KubeletRea	dy				
server11:3.0.0.68	default	server22	v1.9.2	v1.9.2	docker://
17.3.2 KubeletReady					

17.3.2 KubeletRea	dy				
server11:3.0.0.68	default	server12	v1.9.2	v1.9.2	docker://
17.3.2 KubeletRea	dy				
server11:3.0.0.68	default	server23	v1.9.2	v1.9.2	docker://
17.3.2 KubeletRea	dy				
server11:3.0.0.68	default	server24	v1.9.2	v1.9.2	docker://
17.3.2 KubeletRea	dy				

To view only the details for a worker node, specify the hostname at the end of the command using the name option:

```
cumulus@server11:~$ netq server11 show kubernetes node components name server13

Matching kube_cluster records:

Master Cluster Name Node Name Kubelet

KubeProxy Container Runt

ime

server11:3.0.0.68 default server13 v1.9.2 v1.9.2 docker://

17.3.2 KubeletReady
```

You can view information about the replica set:

cumulus@server11:~\$ netq server11 show kubernetes replica-set

Matching kube_replica records:

Master Cluster Name Namespace Replication Name

Labels	Replicas	Ready	y Replicas Last Changed
server11:3.0.0).68 default	default	influxdb-6cdb566dd
app:influx	1	1	14h:19m:28s
server11:3.0.0).68 default	default	nginx-8586cf59
run:nginx	3	3	14h:24m:39s
server11:3.0.0).68 default	default	httpd-5456469bfd
app:httpd	1	1	14h:19m:28s
server11:3.0.0).68 default	kube-syster	m kube-dns-6f4fd4bdf k8s-
app:kube-dns	1	1	14h:27m:9s
server11:3.0.0).68 default	kube-syster	m calico-kube-controllers-d669cc
k8s-app:calico	-kube- 1	1	14h:27m:9s
		78f	controllers

You can view information about the daemon set:

Monitor Container Environments

Use NetQ with Kubernetes Clusters

You can view information about the pod:

cumulus@server11:~\$ netq server11 show kubernetes pod namespace default label nginx Matching kube_pod records: Master Namespace Name IP Node Labels Status Containers Last Changed server11:3.0.0.68 default nginx-8586cf59-26pj5 10.244.9.193 server24 Running nginx:6e2b65070c86 14h:25m:24s run:nginx default nginx-8586cf59-c82ns 10.244.40.128 server12 server11:3.0.0.68 Running nginx:01b017c26725 14h:25m:24s run:nginx default nginx-8586cf59-wjwgp 10.244.49.64 server22 server11:3.0.0.68 Running nginx:ed2b4254e328 14h:25m:24s run:nginx cumulus@server11:~\$ netq server11 show kubernetes pod namespace default label app Matching kube_pod records: Master Namespace Name IP Node Labels Status Containers Last Changed server11:3.0.0.68 default httpd-5456469bfd-bq9 10.244.49.65 server22 Running httpd:79b7f532be2d 14h:20m:34s app:httpd zm default influxdb-6cdb566dd-8 10.244.162.128 server13 server11:3.0.0.68 app:influx Running influxdb:15dce703cdec 14h:20m:34s 9lwn

You can view information about the replication controller:

cumulus@server11:~\$ netq server11 show kubernetes replication-controller

No matching kube_replica records found

You can view information about a deployment:

You can search for information using labels as well. The label search is similar to a "contains" regular expression search. In the following example, we are looking for all nodes that contain *kube* in the replication set name or label:

```
server11:3.0.0.68 default kube-system kube-dns-6f4fd4bdf k8s-
app:kube-dns 1 1 14h:30m:41s
server11:3.0.0.68 default kube-system calico-kube-controllers-d669cc
k8s-app:calico-kube-1 1 14h:30m:41s
78f controllers
```

View Container Connectivity

You can view the connectivity graph of a Kubernetes pod, seeing its replica set, deployment or service level. The impact/connectivity graph starts with the server where the pod is deployed, and shows the peer for each server interface.

```
cumulus@server11:~$ netq server11 show kubernetes deployment name nginx connectivity
nginx -- nginx-8586cf59-wjwgp -- server22:swp1:torbond1 -- swp7:hostbond3:torc-21
-- server22:swp2:torbond1 -- swp7:hostbond3:torc-22
-- server22:swp3:NetQBond-2 -- swp20:NetQBond-20:edge01
-- server22:swp4:NetQBond-2 -- swp20:NetQBond-20:edge02
-- nginx-8586cf59-c82ns -- server12:swp2:NetQBond-1 -- swp23:NetQBond-23:edge01
-- server12:swp3:NetQBond-1 -- swp23:NetQBond-23:edge02
-- server12:swp1:swp1 -- swp6:VlanA-1:tor-1
-- nginx-8586cf59-26pj5 -- server24:swp2:NetQBond-1 -- swp29:NetQBond-29:edge01
-- server24:swp3:NetQBond-1 -- swp29:NetQBond-29:edge02
-- server24:swp1:swp1 -- swp8:VlanA-1:tor-2
```

View Kubernetes Service Connectivity and Impact

You can show the Kubernetes services in a cluster:

Aastor N	lamosnaco Co	onvica Nama	Labola Typo Cluster
	•		Labels Type Cluster
IP External IP	Ports	Last Ch	anged
server11:3.0.0.68	default k	ubernetes	ClusterIP
10.96.0.1	TCP:443	2d:	13h:45m:30s
server11:3.0.0.68	kube-system	calico-etcd	k8s-app:cali ClusterIP
10.96.232.136	TCP:6666		2d:13h:45m:27s
	C	o-etcd	
server11:3.0.0.68	kube-system	kube-dns	k8s-app:kube ClusterIP
10.96.0.10	UDP:53 TCP:	:53	2d:13h:45m:28s
	-1	dns	
server12:3.0.0.69	default k	ubernetes	ClusterIP
10.96.0.1	TCP:443	2d:	13h:46m:24s
server12:3.0.0.69	kube-system	calico-etcd	k8s-app:cali ClusterIP
10.96.232.136	TCP:6666		2d:13h:46m:20s
	C	o-etcd	
server12:3.0.0.69	kube-system	kube-dns	k8s-app:kube ClusterIP
10.96.0.10	UDP:53 TCP:	:53	2d:13h:46m:20s

And get detailed information about a Kubernetes service:

cumulus@server11:~\$ netq show kubernetes service name calico-etcd							
Matching kube_service records:							
Namespace	Service Name	Labels	Type	Cluster			
IP Ports	Last Ch	nanged					
68 kube-syste	m calico-etcd	k8s-app	:cali Clu	sterIP			
TCP:66	666	2d:13h:48m	n:10s				
	co otad						
	co-etca						
69 kube-syste	m calico-etcd	k8s-app	:cali Clu	sterIP			
10.96.232.136 TCP:6666 2d:13h:49m:3s							
	co-etcd						
	Namespace IP Ports 68 kube-syste TCP:66	Namespace Service Name IP Ports Last Ch	Namespace Service Name Labels IP Ports Last Changed	Namespace Service Name Labels Type IP Ports Last Changed			

To see the connectivity of a given Kubernetes service, run:

```
cumulus@server11:~$ netq show kubernetes service name calico-etcd connectivity calico-etcd -- calico-etcd-pfg9r -- server11:swp1:torbond1 -- swp6:hostbond2:torc-11
-- server11:swp2:torbond1 -- swp6:hostbond2:torc-12
-- server11:swp3:NetQBond-2 -- swp16:NetQBond-16:edge01
-- server11:swp4:NetQBond-2 -- swp16:NetQBond-16:edge02
calico-etcd -- calico-etcd-btqgt -- server12:swp1:torbond1 -- swp7:hostbond3:torc-11
-- server12:swp2:torbond1 -- swp7:hostbond3:torc-12
-- server12:swp3:NetQBond-2 -- swp17:NetQBond-17:edge01
-- server12:swp4:NetQBond-2 -- swp17:NetQBond-17:edge02
```

To see the impact of a given Kubernetes service, run:

```
cumulus@server11:~$ netq server11 show impact kubernetes service name calico-etcd
etcd
calico-etcd -- calico-etcd-pfg9r -- server11:swp1:torbond1 --
swp6:hostbond2:torc-11
-- server11:swp2:torbond1 -- swp6:hostbond2:torc-12
-- server11:swp3:NetQBond-2 -- swp16:NetQBond-16:edge01
-- server11:swp4:NetQBond-2 -- swp16:NetQBond-16:edge02
```

View Kubernetes Cluster Configuration in the Past

You can use the "time machine" features of NetQ on a Kubernetes cluster, using the around option to go back in time to check the network status and identify any changes that occurred on the network.

This example shows the current state of the network. Notice there is a node named *server23*. server23 is there because the node *server22* went down and Kubernetes spun up a third replica on a different host to satisfy the deployment requirement.

```
cumulus@redis-1:~$ netq server11 show kubernetes deployment name nginx connectivity
nginx -- nginx-8586cf59-fqtnj -- server12:swp2:NetQBond-1 -- swp23:NetQBond-23:edge01
-- server12:swp3:NetQBond-1 -- swp23:NetQBond-23:edge02
-- server12:swp1:swp1 -- swp6:VlanA-1:tor-1
-- nginx-8586cf59-8g487 -- server24:swp2:NetQBond-1 -- swp29:NetQBond-29:edge01
```

```
-- server24:swp3:NetQBond-1 -- swp29:NetQBond-29:edge02
-- server24:swp1:swp1 -- swp8:VlanA-1:tor-2
-- nginx-8586cf59-2hb8t -- server23:swp1:swp1 -- swp7:VlanA-1:tor-2
-- server23:swp2:NetQBond-1 -- swp28:NetQBond-28:edge01
-- server23:swp3:NetQBond-1 -- swp28:NetQBond-28:edge02
```

You can see this by going back in time 10 minutes. *server23* was not present, whereas *server22* was present:

You can determine the impact on the Kubernetes deployment in the event a host or switch goes down. The output is color coded (not shown in the example below) so you

can clearly see the impact: green shows no impact, yellow shows partial impact, and red shows full impact.

```
cumulus@server11:~$ netq torc-21 show impact kubernetes deployment name
nginx
nginx -- nginx-8586cf59-wjwgp -- server22:swp1:torbond1 --
swp7:hostbond3:torc-21
               -- server22:swp2:torbond1 -- swp7:hostbond3:torc-22
               -- server22:swp3:NetQBond-2 -- swp20:NetQBond-20:edge01
               -- server22:swp4:NetQBond-2 -- swp20:NetQBond-20:edge02
   -- nginx-8586cf59-c82ns -- server12:swp2:NetQBond-1 --
swp23:NetQBond-23:edge01
               -- server12:swp3:NetQBond-1 -- swp23:NetQBond-23:edge02
               -- server12:swp1:swp1 -- swp6:VlanA-1:tor-1
   -- nginx-8586cf59-26pj5 -- server24:swp2:NetQBond-1 --
swp29:NetQBond-29:edge01
               -- server24:swp3:NetQBond-1 -- swp29:NetQBond-29:edge02
               -- server24:swp1:swp1 -- swp8:VlanA-1:tor-2
cumulus@server11:~$ netq server12 show impact kubernetes deployment name
nginx
nginx -- nginx-8586cf59-wjwgp -- server22:swp1:torbond1 --
swp7:hostbond3:torc-21
               -- server22:swp2:torbond1 -- swp7:hostbond3:torc-22
               -- server22:swp3:NetQBond-2 -- swp20:NetQBond-20:edge01
               -- server22:swp4:NetQBond-2 -- swp20:NetQBond-20:edge02
   -- nginx-8586cf59-c82ns -- server12:swp2:NetQBond-1 --
swp23:NetQBond-23:edge01
               -- server12:swp3:NetQBond-1 -- swp23:NetQBond-23:edge02
               -- server12:swp1:swp1 -- swp6:VlanA-1:tor-1
```

```
-- nginx-8586cf59-26pj5 -- server24:swp2:NetQBond-1 -- swp29:NetQBond-29:edge01
-- server24:swp3:NetQBond-1 -- swp29:NetQBond-29:edge02
```

Manage NetQ Agents

At various points in time, you might want to change which network nodes are being monitored by NetQ or look more closely at a network node for troubleshooting purposes. Adding the NetQ Agent to a switch or host is described in Install NetQ. Viewing the status of an Agent, disabling an Agent, and managing NetQ Agent logging are presented.

View NetQ Agent Status

To view the health of your NetQ Agents, use the netq show agents command:

netq [<hostname>] show agents [fresh | dead | rotten | opta] [around <texttime>] [json]

You can view the status for a given switch, host or NetQ server. You can also filter by the status as well as view the status at a time in the past.

To view the current status of all NetQ Agents:

cumulus@switch~:\$ netq show agents

Matching agents records:

Hostname Status NTP Sync Version Sys Uptime

Agent Uptime Reinitialize Time Last Changed

exit-1	Fresh	yes	2.2.1-cl3u19~15645	03011.e3b463d	1d:4h:35m:
11s	1d:4h:14m:3	4s	1d:4h:14m:34s	Wed Jul 31 16	5:50:40 2019
exit-2	Fresh	yes	2.2.1-cl3u19~15645	03011.e3b463d	1d:4h:35m:
11s	1d:4h:14m:3	0s	1d:4h:14m:30s	Wed Jul 31 16	5:51:07 2019
firewall-1	Fresh	yes	2.2.1-ub16.04u19	~1564494614.6fec	l81f 1d:4h:
35m:11s	1d:4h:1	4m:24s	1d:4h:14m:2	4s Wed Jul	31 16:51:13
2019					
firewall-2	Fresh	yes	2.2.1-rh7u19~156	4496494.6fed81f	1d:4h:34m:
35s	1d:4h:14m:1	8s	1d:4h:14m:18s	Wed Jul 31 16	5:51:06 2019
hostd-11	Fresh	yes	2.2.1-ub16.04u19	%1564494614.6fe	d81f 1d:4h:
35m:6s	1d:4h:14	4m:6s	1d:4h:14m:6s	Wed Jul 3	1 16:51:16 2019
hostd-12	Fresh	yes	2.2.1-rh7u19~156	64496494.6fed81f	1d:4h:34m:
40s	1d:4h:14m:2	S	1d:4h:14m:2s	Wed Jul 31 16:	51:40 2019
•••					

To view NetQ Agents that are not communicating:

cumulus@switch~:\$ netq show agents rotten

No matching agents records found

To view NetQ Agent status on the NetQ Server or Appliance, run the following command from a node:

cumulus@leaf01~:\$ netq show agents opta Matching agents records:

Hostname	Status	NTP S	Sync Ver	sion		Sys Uptim	е
Agent Uptime	Reini	tialize ⁻	Time	Last Cha	nged		
					-		
10-20-14-157	Fresh	yes	2.2.1-c	l3u19~156	1299612.	73c7ab4	1d:5h:
40m:41s	6m:34.417s	;	6m:34	1.417s	Wed	Jul 31 22:1	2:40 2019

Modify the Configuration of the NetQ Agent on a Node

The agent configuration commands enable you to add and remove agents from switches and hosts, start and stop agent operations, add and remove Kubernetes container monitoring, add or remove sensors, debug the agent, and add or remove FRR (FRRouting).



(i) NOTE

Commands apply to one agent at a time, and are run from the switch or host where the NetQ Agent resides.

The agent configuration commands include:

netq config add agent frr-monitor [<text-frr-docker-name>] netq config add agent kubernetes-monitor [poll-period <text-duration-period>] netq config add agent loglevel [debug|error|info|warning] netq config add agent sensors netq config add agent server <text-opta-ip> [port <text-opta-port>] [vrf <text-vrfname>]

netq config (start|stop|status|restart) agent
netq config del agent (agent-url|frr-monitor|kubernetes-monitor|loglevel|
sensors|server)
netq config show agent [frr-monitor|kubernetes-monitor|loglevel|sensors] [json]

This example shows how to specify the IP address and optionally a specific port on the NetQ Platform where agents should send their data.

cumulus@switch~:\$ netq config add agent server 10.0.0.23

This example shows how to configure the agent to send sensor data.

cumulus@switch~:\$ netq config add agent sensors

This example shows how to start monitoring with Kubernetes.

cumulus@switch:~\$ netq config add kubernetes-monitor



After making configuration changes to your agents, you must restart the agent for the changes to take effect. Use the netq config restart agent command.

Disable the NetQ Agent on a Node

You can temporarily disable NetQ Agent on a node. Disabling the agent maintains the activity history in the NetQ database.

To disable NetQ Agent on a node, run the following command from the node:

cumulus@switch:~\$ netq config stop agent

Remove the NetQ Agent from a Node

You can decommission a NetQ Agent on a given node. You might need to do this when you:

- RMA the switch or host being monitored
- Change the hostname of the switch or host being monitored
- Move the switch or host being monitored from one data center to another

(i) NOTE

Decommissioning the node removes the agent server settings from the local configuration file.

To decommission a node from the NetQ database:

1. On the given node, stop and disable the NetQ Agent service.

cumulus@switch:~\$ sudo systemctl stop netq-agent cumulus@switch:~\$ sudo systemctl disable netq-agent

2. On the NetQ Appliance or Platform, decommission the node.

cumulus@netq-appliance:~\$ netq decommission <hostname>

Configure Logging for a NetQ Agent

The logging level used for a NetQ Agent determines what types of events are logged about the NetQ Agent on the switch or host.

First, you need to decide what level of logging you want to configure. You can configure the logging level to be the same for every NetQ Agent, or selectively increase or decrease the logging level for a NetQ Agent on a problematic node.

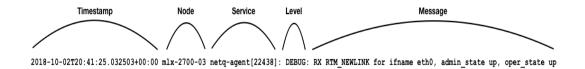
Logging Level	Description
debug	Sends notifications for all debugging-related, informational, warning, and error messages.
info	Sends notifications for informational, warning, and error messages (default).
warning	Sends notifications for warning and error messages.
error	Sends notifications for errors messages.

You can view the NetQ Agent log directly. Messages have the following structure:

<timestamp> <node> <service>[PID]: <level>: <message>

Element	Description
timestamp	Date and time event occurred in UTC format
node	Hostname of network node where event occurred
service [PID]	Service and Process IDentifier that generated the event
level	Logging level in which the given event is classified; <i>debug</i> , <i>error</i> , <i>info</i> , or <i>warning</i>
message	Text description of event, including the node where the event occurred

For example:



This example shows a portion of a NetQ Agent log with debug level logging.

2019-02-16T18:45:53.951124+00:00 spine-1 netq-agent[8600]: INFO: OPTA

Discovery exhibit url hydra-09.cumulusnetworks.com port 4786

2019-02-16T18:45:53.952035+00:00 spine-1 netq-agent[8600]: INFO: OPTA

Discovery Agent ID spine-1

2019-02-16T18:45:53.960152+00:00 spine-1 netq-agent[8600]: INFO: Received

Discovery Response 0

2019-02-16T18:46:54.054160+00:00 spine-1 netq-agent[8600]: INFO: OPTA

Discovery exhibit url hydra-09.cumulusnetworks.com port 4786

2019-02-16T18:46:54.054509+00:00 spine-1 netq-agent[8600]: INFO: OPTA

Discovery Agent ID spine-1

2019-02-16T18:46:54.057273+00:00 spine-1 netq-agent[8600]: INFO: Received

Discovery Response 0

2019-02-16T18:47:54.157985+00:00 spine-1 netq-agent[8600]: INFO: OPTA

Discovery exhibit url hydra-09.cumulusnetworks.com port 4786

2019-02-16T18:47:54.158857+00:00 spine-1 netq-agent[8600]: INFO: OPTA

Discovery Agent ID spine-1

2019-02-16T18:47:54.171170+00:00 spine-1 netq-agent[8600]: INFO: Received

Discovery Response 0

2019-02-16T18:48:54.260903+00:00 spine-1 netq-agent[8600]: INFO: OPTA

Discovery exhibit url hydra-09.cumulusnetworks.com port 4786

...

Example: Configure debug-level logging

1. Set the logging level to *debug*.

cumulus@switch:~\$ netq config add agent loglevel debug

2. Restart the NetQ Agent.

cumulus@switch:~\$ netq config restart agent

Optionally, verify connection to the NetQ platform by viewing the netq-agent.log messages.

Example: Configure warning-level logging

cumulus@switch:~\$ netq config add agent loglevel warning cumulus@switch:~\$ netq config restart agent

Example: Disable Agent Logging

If you have set the logging level to *debug* for troubleshooting, it is recommended that you either change the logging level to a less heavy mode or completely disable agent logging altogether when you are finished troubleshooting.

To change the logging level, run the following command and restart the agent service:

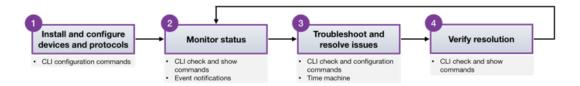
cumulus@switch:~\$ netq config add agent loglevel <LOG_LEVEL> cumulus@switch:~\$ netq config restart agent

To disable all logging:

cumulus@switch:~\$ netq config del agent loglevel cumulus@switch:~\$ netq config restart agent

Investigate NetQ Issues

Monitoring of systems inevitably leads to the need to troubleshoot and resolve the issues found. In fact network management follows a common pattern as shown in this diagram.



This topic describes some of the tools and commands you can use to troubleshoot issues with the network and NetQ itself. Some example scenarios are included here: - viewing configuration and log files, - verifying NetQ Agent health, - investigating recent events, and - investigating events from the past - running a trace.

Try looking at the specific protocol or service, or particular devices as well. If none of these produce a resolution, you can **capture a log** to use in discussion with the Cumulus Networks support team.

Browse Configuration and Log Files

To aid in troubleshooting issues with NetQ, there are the following configuration and log files that can provide insight into the root cause of the issue:

File	Description
/etc/netq/ netq.yml	The NetQ configuration file. This file appears only if you installed either the netq-apps package or the NetQ Agent on the system.

File	Description
/var/log/ netqd.log	The NetQ daemon log file for the NetQ CLI. This log file appears only if you installed the netq-apps package on the system.
/var/log/netq- agent.log	The NetQ Agent log file. This log file appears only if you installed the NetQ Agent on the system.

Check NetQ Agent Health

Checking the health of the NetQ Agents is a good way to start troubleshooting NetQ on your network. If any agents are rotten, meaning three heartbeats in a row were not sent, then you can investigate the rotten node. In the example below, the NetQ Agent on *server01* is rotten, so you know where to start looking for problems:

```
cumulus@switch:$ netq check agents

Checked nodes: 12,

Rotten nodes: 1

netq@446c0319c06a:/$ netq show agents

Node Status Sys Uptime Agent Uptime

------
exit01

Fresh

8h ago 4h ago
exit02

Fresh

8h ago 4h ago
leaf01

Fresh
```

8h ago 4h ago leaf02 Fresh 8h ago 4h ago leaf03 Fresh 8h ago 4h ago leaf04 Fresh 8h ago 4h ago server01 Rotten 4h ago 4h ago server02 Fresh 4h ago 4h ago server03 Fresh 4h ago 4h ago server04 Fresh 4h ago 4h ago spine01 Fresh 4h ago 8h ago spine02 Fresh 8h ago 4h ago

Diagnose an Event after It Occurs

NetQ provides users with the ability to go back in time to replay the network state, see fabric-wide event change logs and root cause state deviations. The NetQ Telemetry Server maintains data collected by NetQ agents in a time-series database, making fabric-wide events available for analysis. This enables you to replay and analyze network-wide events for better visibility and to correlate patterns. This allows for root-cause analysis and optimization of network configs for the future.

NetQ provides a number of commands for diagnosing past events.

NetQ records network events and stores them in its database. You can view the events through a third-party notification application like PagerDuty or Slack or use **netq show events** to look for any changes made to the runtime configuration that may have triggered the alert, then use **netq trace** to track the connection between the nodes.

The netq trace command traces the route of an IP or MAC address from one endpoint to another. It works across bridged, routed and VXLAN connections, computing the path using available data instead of sending real traffic — this way, it can be run from anywhere. It performs MTU and VLAN consistency checks for every link along the path.

For example, say you get an alert about a BGP session failure. You can quickly run netq check bgp to determine what sessions failed:

firewall-1 swp6.2: AFI/ 1d:7h:56m:9s

SAFI evpn not activated on peer

exit-1 DataVrf1080 swp7.2 firewall-2 BGP session with peer

firewall-2 (swp7.2 vrf 1d:7h:49m:31s

DataVrf1080) failed,

reason: Peer not configured

exit-1 DataVrf1081 swp6.3 firewall-1 BGP session with peer

firewall-1 swp6.3: AFI/ 1d:7h:56m:9s

SAFI evpn not activated on peer

exit-1 DataVrf1081 swp7.3 firewall-2 BGP session with peer

firewall-2 (swp7.3 vrf 1d:7h:49m:31s

DataVrf1081) failed,

reason: Peer not configured

You can run a trace from spine01 to leaf02, which has the IP address 10.1.20.252:

cumulus@switch:~\$ netq trace 10.1.20.252 from spine01 around 5m spine01 -- spine01:swp1 -- leaf01:vlan20

-- spine01:swp2 -- leaf02:vlan20

Then you can check what's changed on the network to help you identify the problem.

cumulus@switch:~\$ netq show events type bgp

Matching events records:

Hostname Message Type Severity Message Timestamp

leaf21	bgp	info BGP session with peer spine-1 swp3. 1d:8h:35m:19s
		3 vrf DataVrf1081 state changed fro
		m failed to Established
leaf21	bgp	info BGP session with peer spine-2 swp4. 1d:8h:35m:19s
		3 vrf DataVrf1081 state changed fro
		m failed to Established
leaf21	bgp	info BGP session with peer spine-3 swp5. 1d:8h:35m:19s
		3 vrf DataVrf1081 state changed fro
		m failed to Established
leaf21	bgp	info BGP session with peer spine-1 swp3. 1d:8h:35m:19s
		2 vrf DataVrf1080 state changed fro
		m failed to Established
leaf21	bgp	info BGP session with peer spine-3 swp5. 1d:8h:35m:19s
		2 vrf DataVrf1080 state changed fro
		m failed to Established
•••		

Use NetQ as a Time Machine

With NetQ, you can travel back to a specific point in time or a range of times to help you isolate errors and issues.

For example, if you think you had an issue with your sensors last night, you can check the sensors on all your nodes around the time you think the issue occurred:

cumulus@leaf01:~\$ netq check sensors around 12h

Total Nodes: 25, Failed Nodes: 0, Checked Sensors: 221, Failed Sensors: 0

Or you can specify a range of times using the **between** option. The units of time you can specify are second (*s*), minutes (*m*), hours (*h*) and days (*d*). Always specify the most recent time first, then the more distant time. For example, to see the changes made to the network between the past minute and 5 minutes ago, you'd run:

Matching e	vents records	s:		
Hostname	J		Severity Message	Timestamp
leaf21			leaf21 config file ptm was r	
leaf21	configdiff	info	leaf21 config file lldpd was	modifi 1d:8h:38m:6s
	ec	k		
leaf21	configdiff	info	leaf21 config file interfaces	was m 1d:8h:38m:6s
	00	dified		
leaf21	configdiff	info	leaf21 config file frr was mo	odified 1d:8h:38m:6s
leaf12	configdiff	info	leaf12 config file ptm was r	modified 1d:8h:38m:11s
leaf12	configdiff	info	leaf12 config file lldpd was	modifi 1d:8h:38m:11s
	ec	k		
leaf12	configdiff	info	leaf12 config file interfaces	was m 1d:8h:38m:11s
	00	dified		
leaf12	configdiff	info	leaf12 config file frr was mo	odified 1d:8h:38m:11s
leaf11	configdiff	info	leaf11 config file ptm was r	modified 1d:8h:38m:22s

You can travel back in time 5 minutes and run a trace from spine02 to exit01, which has the IP address 27.0.0.1:

```
cumulus@leaf01:~$ netq trace 27.0.0.1 from spine02 around 5m

Detected Routing Loop. Node exit01 (now via Local Node exit01 and Ports swp6

<==> Remote Node/s spine01 and Ports swp3) visited twice.
```

Detected Routing Loop. Node spine02 (now via mac:00:02:00:00:05) visited twice.

```
spine02 -- spine02:swp3 -- exit01:swp6.4 -- exit01:swp3 -- exit01
-- spine02:swp7 -- spine02
```

Trace Paths in a VRF

The netq trace command works with VRFs as well:

```
cumulus@leaf01:~$ netq trace 10.1.20.252 from spine01 vrf default around 5m spine01 -- spine01:swp1 -- leaf01:vlan20 -- spine01:swp2 -- leaf02:vlan20
```

Generate a Support File

The opta-support command generates an archive of useful information for troubleshooting issues with NetQ. It is an extension of the cl-support command in Cumulus Linux. It provides information about the NetQ Platform configuration and runtime statistics as well as output from the docker ps command. The Cumulus Networks support team may request the output of this command when assisting with any issues that you could not solve with your own troubleshooting. Run the following command:

cumulus@switch:~\$ opta-support

Resolve MLAG Issues

This topic outlines a few scenarios that illustrate how you use NetQ to troubleshoot MLAG on Cumulus Linux switches. Each starts with a log message that indicates the current MLAG state.

NetQ can monitor many aspects of an MLAG configuration, including:

- Verifying the current state of all nodes
- Verifying the dual connectivity state
- Checking that the peer link is part of the bridge
- Verifying whether MLAG bonds are not bridge members
- Verifying whether the VXLAN interface is not a bridge member
- Checking for remote-side service failures caused by systemctl
- Checking for VLAN-VNI mapping mismatches
- Checking for layer 3 MTU mismatches on peerlink subinterfaces
- Checking for VXLAN active-active address inconsistencies
- Verifying that STP priorities are the same across both peers

Scenario: All Nodes Are Up

When the MLAG configuration is running smoothly, NetQ sends out a message that all nodes are up:

2017-05-22T23:13:09.683429+00:00 noc-pr netq-notifier[5501]: INFO: CLAG: All nodes are up

Scenario: All Nodes Are Up

Running netq show clag confirms this:

latching cl	ag records:									
Hostname	Peer	SysMac	Stat	e	Bac	kup :	#Bor	nd #D	ual Last	
Changed										
		S								
spine01(P)	spine02	00:01:01:10:	00:01	 up		 up	24	24	Thu Feb	7
18:30:49 20	19									
spine02	spine01(P)	00:01:01:10:	00:01	up		up	24	24	Thu Feb	7
18:30:53 20	19									
leaf01(P)	leaf02	44:38:39:ff:ff:0	1 up	ι	лр	12	12	Thu	Feb 7	
18:31:15 20	19									
leaf02	leaf01(P)	44:38:39:ff:ff:0	1 up	ι	лр	12	12	Thu	Feb 7	
18:31:20 20	19									
leaf03(P)	leaf04	44:38:39:ff:ff:02	2 up	L	лр	12	12	Thu	Feb 7	
18:31:26 20	19									
leaf04	leaf03(P)	44:38:39:ff:ff:0	2 up	L	au	12	12	Thu	Feb 7	

You can also verify a specific node is up:

cumulus@switch:~\$ netq spine01 show clag

Matching clag records:

Hostname Peer SysMac State Backup #Bond #Dual Last

Changed

Scenario: All Nodes Are Up

s ------spine01(P) spine02 00:01:01:10:00:01 up up 24 24 Thu Feb 7 18:30:49 2019

Similarly, checking the MLAG state with NetQ also confirms this:

cumulus@switch:~\$ netq check clag

Checked Nodes: 6, Failed Nodes: 0

When you are logged directly into a switch, you can run clagctl to get the state:

cumulus@switch:/var/log# sudo clagctl

The peer is alive

Peer Priority, ID, and Role: 4096 00:02:00:00:4e primary

Our Priority, ID, and Role: 8192 44:38:39:00:a5:38 secondary

Peer Interface and IP: peerlink-3.4094 169.254.0.9

VxLAN Anycast IP: 36.0.0.20

Backup IP: 27.0.0.20 (active)

System MAC: 44:38:39:ff:ff:01

CLAG Interfaces

Our Interface Peer Interface CLAG Id Conflicts Proto-Down Reason

```
      vx-33
      -
      -
      -

      hostbond4
      1
      -
      -

      hostbond5
      2
      -
      -

      vx-37
      vx-37
      -
      -

      vx-36
      vx-36
      -
      -

      vx-35
      vx-35
      -
      -

      vx-34
      vx-34
      -
      -
```

Scenario: Dual-connected Bond Is Down

When dual connectivity is lost in an MLAG configuration, you receive messages from NetQ similar to the following:

```
2017-05-22T23:14:40.290918+00:00 noc-pr netq-notifier[5501]: WARNING: LINK: 1 link(s) are down. They are: spine01 hostbond5
2017-05-22T23:14:53.081480+00:00 noc-pr netq-notifier[5501]: WARNING: CLAG: 1 node(s) have failures. They are: spine01
2017-05-22T23:14:58.161267+00:00 noc-pr netq-notifier[5501]: WARNING: CLAG: 2 node(s) have failures. They are: spine01, leaf01
```

To begin your investigation, show the status of the clagd service:

```
cumulus@switch:~$ netq spine01 show services clagd

Matching services records:

Hostname Service PID VRF Enabled Active Monitored

Status Uptime Last Changed
```

Checking the MLAG status provides the reason for the failure:

```
cumulus@switch:~$ netq check clag

Checked Nodes: 6, Warning Nodes: 2

Node Reason

------

spine01 Link Down: hostbond5

leaf01 Singly Attached Bonds: hostbond5
```

You can retrieve the output in JSON format for export to another tool:

```
"reason": "Singly Attached Bonds: hostbond5"

}

],

"failedNodes":[

],

"summary":{

"checkedNodeCount":6,

"failedNodeCount":0,

"warningNodeCount":2

}

}
```

After you fix the issue, you can show the MLAG state to see if all the nodes are up. The notifications from NetQ indicate all nodes are UP, and the netq check flag also indicates there are no failures.

```
Cumulus@switch:~$ netq show clag

Matching clag records:

Hostname Peer SysMac State Backup #Bond #Dual Last

Changed

s

spine01(P) spine02 00:01:01:10:00:01 up up 24 24 Thu Feb 7

18:30:49 2019

spine02 spine01(P) 00:01:01:10:00:01 up up 24 24 Thu Feb 7

18:30:53 2019

leaf01(P) leaf02 44:38:39:ff:ff:01 up up 12 12 Thu Feb 7
```

18:31:15 20)19					
leaf02	leaf01(P)	44:38:39:ff:ff:01 up	up	12	12	Thu Feb 7
18:31:20 20)19					
leaf03(P)	leaf04	44:38:39:ff:ff:02 up	up	12	12	Thu Feb 7
18:31:26 20)19					
leaf04	leaf03(P)	44:38:39:ff:ff:02 up	up	12	12	Thu Feb 7
18:31:30 20)19					

When you are logged directly into a switch, you can run clagctl to get the state:

```
cumulus@switch:/var/log# sudo clagctl
The peer is alive
Peer Priority, ID, and Role: 4096 00:02:00:00:00:4e primary
Our Priority, ID, and Role: 8192 44:38:39:00:a5:38 secondary
Peer Interface and IP: peerlink-3.4094 169.254.0.9
VxLAN Anycast IP: 36.0.0.20
Backup IP: 27.0.0.20 (active)
System MAC: 44:38:39:ff:ff:01
CLAG Interfaces
Our Interface Peer Interface CLAG Id Conflicts Proto-Down Reason
vx-38 vx-38 - -
hostbond4 hostbond4 1 -
hostbond5 - 2 - -
vx-37 vx-37 - -
```

Resolve MLAG Issues Scenario: VXLAN Active-active Device or Interface Is Down

Scenario: VXLAN Active-active Device or Interface Is Down

When a VXLAN active-active device or interface in an MLAG configuration is down, log messages also include VXLAN and LNV checks.

```
2017-05-22T23:16:51.517522+00:00 noc-pr netq-notifier[5501]: WARNING: VXLAN: 2 node(s) have failures. They are: spine01, leaf01 2017-05-22T23:16:51.525403+00:00 noc-pr netq-notifier[5501]: WARNING: LINK: 2 link(s) are down. They are: leaf01 vx-37, spine01 vx-37 2017-05-22T23:16:54.194681+00:00 noc-pr netq-notifier[5501]: WARNING: LNV: 1 node(s) have failures. They are: leaf02 2017-05-22T23:16:59.448755+00:00 noc-pr netq-notifier[5501]: WARNING: LNV: 3 node(s) have failures. They are: leaf01, leaf03, leaf04 2017-05-22T23:17:04.703044+00:00 noc-pr netq-notifier[5501]: WARNING: CLAG: 2 node(s) have failures. They are: spine01, leaf01
```

To begin your investigation, show the status of the clagd service:

```
cumulus@switch:~$ netq spine01 show services clagd

Matching services records:

Hostname Service PID VRF Enabled Active Monitored
```

Checking the MLAG status provides the reason for the failure:

```
cumulus@switch:~$ netq check clag

Checked Nodes: 6, Warning Nodes: 2, Failed Nodes: 2

Node Reason

------
spine01 Protodown Bonds: vx-37:vxlan-single

leaf01 Protodown Bonds: vx-37:vxlan-single
```

You can retrieve the output in JSON format for export to another tool:

After you fix the issue, you can show the MLAG state to see if all the nodes are up:

```
cumulus@switch:~$ netq show clag
Matching clag session records are:
Hostname
            Peer
                      SysMac State Backup #Bond #Dual Last
Changed
                               S
spine01(P) spine02 00:01:01:10:00:01 up up 24 24 Thu Feb 7
18:30:49 2019
spine02
          spine01(P) 00:01:01:10:00:01 up up 24 24 Thu Feb 7
18:30:53 2019
leaf01(P) leaf02 44:38:39:ff:ff:01 up up 12 12 Thu Feb 7
18:31:15 2019
         leaf01(P) 44:38:39:ff:ff:01 up up 12 12 Thu Feb 7
leaf02
18:31:20 2019
leaf03(P) leaf04 44:38:39:ff:ff:02 up up 12 12 Thu Feb 7
```

```
18:31:26 2019
leaf04 leaf03(P) 44:38:39:ff:ff:02 up up 12 12 Thu Feb 7
18:31:30 2019
```

When you are logged directly into a switch, you can run clagctl to get the state:

```
cumulus@switch:/var/log# sudo clagctl
The peer is alive
Peer Priority, ID, and Role: 4096 00:02:00:00:00:4e primary
Our Priority, ID, and Role: 8192 44:38:39:00:a5:38 secondary
Peer Interface and IP: peerlink-3.4094 169.254.0.9
VxLAN Anycast IP: 36.0.0.20
Backup IP: 27.0.0.20 (active)
System MAC: 44:38:39:ff:ff:01
CLAG Interfaces
Our Interface Peer Interface CLAG Id Conflicts Proto-Down Reason
vx-38 - -
vx-33 - - -
hostbond4 hostbond4 1 -
hostbond5 hostbond5 2
vx-37 - - vxlan-single
vx-36 vx-36 - -
vx-35 vx-35
vx-34 vx-34 -
```

Scenario: Remote-side clagd Stopped by systemctl Command

In the event the **clagd** service is stopped via the **systemctl** command, NetQ Notifier sends messages similar to the following:

2017-05-22T23:51:19.539033+00:00 noc-pr netq-notifier[5501]: WARNING: VXLAN: 1 node(s) have failures. They are: leaf01 2017-05-22T23:51:19.622379+00:00 noc-pr netq-notifier[5501]: WARNING: LINK: 2 link(s) flapped and are down. They are: leaf01 hostbond5, leaf01 hostbond4 2017-05-22T23:51:19.622922+00:00 noc-pr netq-notifier[5501]: WARNING: LINK: 23 link(s) are down. They are: leaf01 VlanA-1-104-v0, leaf01 VlanA-1-101-v0, leaf01 VlanA-1, leaf01 vx-33, leaf01 vx-36, leaf01 vx-37, leaf01 vx-34, leaf01 vx-35, leaf01 swp7, leaf01 VlanA-1-102-v0, leaf01 VlanA-1-103-v0, leaf01 VlanA-1-100-v0, leaf01 VlanA-1-106-v0, leaf01 swp8, leaf01 VlanA-1.106, leaf01 VlanA-1.105, leaf01 VlanA-1.104, leaf01 VlanA-1.103, leaf01 VlanA-1.102, leaf01 VlanA-1.101, leaf01 VlanA-1.100, leaf01 VlanA-1-105-v0, leaf01 vx-38 2017-05-22T23:51:27.696572+00:00 noc-pr netq-notifier[5501]: INFO: LINK: 15 link(s) are up. They are: leaf01 VlanA-1.106, leaf01 VlanA-1-104-v0, leaf01 VlanA-1.104, leaf01 VlanA-1.103, leaf01 VlanA-1.101, leaf01 VlanA-1-100-v0, leaf01 VlanA-1.100, leaf01 VlanA-1.102, leaf01 VlanA-1-101-v0, leaf01 VlanA-1-102-v0, leaf01 VlanA-1.105, leaf01 VlanA-1-103-v0, leaf01 VlanA-1-106-v0, leaf01 VlanA-1, leaf01 VlanA-1-105-v0 2017-05-22T23:51:30.863789+00:00 noc-pr netq-notifier[5501]: WARNING: LNV: 1 node(s) have failures. They are: leaf01 2017-05-22T23:51:36.156708+00:00 noc-pr netq-notifier[5501]: WARNING: CLAG: 2 node(s) have failures. They are: spine01, leaf01 2017-05-22T23:51:36.183638+00:00 noc-pr netq-notifier[5501]: WARNING: LNV: 2 node(s) have failures. They are: spine02, leaf01

```
2017-05-22T23:51:41.444670+00:00 noc-pr netq-notifier[5501]: WARNING: LNV: 1 node(s) have failures. They are: leaf01
```

Showing the MLAG state reveals which nodes are down:

Checking the MLAG status provides the reason for the failure:

```
spine01 Peer Connectivity failed
leaf01 Peer Connectivity failed
```

You can retrieve the output in JSON format for export to another tool:

```
cumulus@switch:~$ netq check clag json
{
  "failedNodes": [
    {
      "node": "spine01",
      "reason": "Peer Connectivity failed"
    }
      "node": "leaf01",
      "reason": "Peer Connectivity failed"
    }
  ],
  "summary":{
    "checkedNodeCount": 6,
    "failedNodeCount": 2,
    "warningNodeCount": 1
 }
}
```

When you are logged directly into a switch, you can run clagctl to get the state:

cumulus@switch:~\$ sudo clagctl

The peer is not alive

Our Priority, ID, and Role: 8192 44:38:39:00:a5:38 primary

Peer Interface and IP: peerlink-3.4094 169.254.0.9

VxLAN Anycast IP: 36.0.0.20

Backup IP: 27.0.0.20 (inactive)

System MAC: 44:38:39:ff:ff:01

CLAG Interfaces

Our Interface Peer Interface CLAG Id Conflicts Proto-Down Reason

vx-38 - - - -

vx-33 - - - -

hostbond4 - 1 - -

hostbond5 - 2 - -

vx-37 - - - -

vx-36 - - -

vx-35 - - - -

vx-34 - - -

CLI Early Access Features

NetQ has early access features that provide advanced access to new functionality before it becomes generally available. Two features are available as early access features in NetQ 2.3.0:

- Detailed netg check validation output (CLI only), described here
- Topology view (GUI Cloud deployments only), refer to UI Early Access Features

These features are bundled into the netq-apps package; there is no specific EA package like there typically is with Cumulus Linux.

These features are provided as is, and are subject to change before they become generally available.

Enable/Disable Early Access Features

You enable early access features by running the **netq config add experimental** command on any node running NetQ.

cumulus@switch:~\$ netq config add experimental Experimental config added

You disable the early access features by running the **netq config del experimental** command on any node running NetQ.

cumulus@switch:~\$ netq config del experimental Experimental config deleted

Validation Commands

If you enable the experimental commands, you are able to access the updated validation commands. Two key differences are present with these updated commands:

- You can view more detail about the validation tests that are run with each command
- You can filter these tests to run only those tests of interest

The syntax of the updated netq check commands is:

netq check agents [around <text-time>] [json]

netq check bgp [vrf <vrf>] [include <bgp-number-range-list> | exclude <bgp-number-range-list>] [around <text-time>] [json]

netq check clag [include <clag-number-range-list> | exclude <clag-number-range-list>] [around <text-time>] [json]

netq check evpn [mac-consistency] [include <evpn-number-range-list> | exclude <evpn-number-range-list> | exclude <evpn-number-range-list> | faround <text-time>] [json]

netq check interfaces [include <interface-number-range-list> | exclude <interface-number-range-list> | exclude

list>] [around <text-time>] [json]

netq check ospf [include <ospf-number-range-list> | exclude <ospf-number-range-list>] [around <text-time>] [json]

netq check sensors [include <sensors-number-range-list> | exclude <sensors-number-range-list>] [around <text-time>] [json]

netq check vlan [unverified] [include <vlan-number-range-list> | exclude <vlan-number-range-list>] [around <text-time>] [json]

netq check vxlan [include <vxlan-number-range-list> | exclude <vxlan-number-range-list>] [around <text-time>] [json]

Each of the check commands provides a starting point for troubleshooting configuration and connectivity issues within your network in real time.

A summary of the validation results is achieved by running the **netq check** commands without any options; for example, **netq check agents** or **netq check evpn**. This summary displays such data as the total number of nodes checked, how many failed a test, total number of sessions checked, how many of these that failed, and so forth.

With the NetQ 2.3.0 release, you have can view more information about the individual tests that are run as part of the validation, with the exception of agents and LNV.

You can run validations for a time in the past and output the results in JSON format if desired. The **around** option enables users to view the network state at an earlier time. The **around** option value requires an integer *plus* a unit of measure (UOM), with no space between them. The following are valid UOMs:

UOM	Command Value	Example
day(s)	<#>d	3d
hour(s)	<#>h	6h

UOM	Command Value	Example
minute(s)	<#>m	30m
second(s)	<#>s	20s



If you want to go back in time by months or years, use the equivalent number of days.

For validation commands that have the include <protocol-number-range-list> and exclude <protocol-number-range-list> options, you can include or exclude one or more of the various tests performed during the validation. Each test is assigned a number, which is used to identify which tests to run. By default, all tests are run. The value of commas, or a range using a dash, or a combination of these. Do not use spaces after commas. For example:

- include 1,3,5
- include 1-5
- include 1,3-5
- exclude 6,7
- exclude 6-7
- exclude 3,4-7,9

The output indicates whether a given test passed, failed, or was skipped.



TIP

Output from the netq check commands are color-coded; green for successful results and red for failures, warnings, and errors. Use the netq config add color command to enable the use of color.

What the NetQ Validation System Checks

Each of the **netq check** commands perform a set of validation tests appropriate to the protocol or element being validated. This section describes these tests.

NetQ Agent Validation Tests

The netq check agents command looks for an agent status of Rotten for each node in the network. A *Fresh* status indicates the Agent is running as expected. The Agent sends a heartbeat every 30 seconds, and if three consecutive heartbeats are missed, its status changes to *Rotten*.

BGP Validation Tests

The netq check bgp command runs the following tests to establish session sanity:

Test Number	Test Name	Description
0	Session Establishment	Checks that BGP sessions are in an established state
1	Address Families	Checks if transmit and receive address family advertisement is consistent between peers of a BGP session
2	Router ID	Checks for BGP router ID conflict in the network

Validation Commands

CLAG Validation Tests

The netq check clag command runs the following tests:

Test Number	Test Name	Description
0	Peering	 Checks if: CLAG peerlink is up CLAG peerlink bond slaves are down (not in full capacity and redundancy) Peering is established between two nodes in a CLAG pair
1	Backup IP	 Checks if: CLAG backup IP configuration is missing on a CLAG node CLAG backup IP is correctly pointing to the CLAG peer and its connectivity is available
2	Clag Sysmac	 Checks if: CLAG Sysmac is consistently configured on both nodes in a CLAG pair there is any duplication of a CLAG sysmac within a bridge domain
3	VXLAN Anycast IP	Checks if the VXLAN anycast IP address is consistently configured on both nodes in a CLAG pair
4	Bridge Membership	Checks if the CLAG peerlink is part of bridge

Test Number	Test Name	Description
5	Spanning Tree	 Checks if: STP is enabled and running on the CLAG nodes CLAG peerlink role is correct from STP perspective the bridge ID is consistent between two nodes of a CLAG pair the VNI in the bridge has BPDU guard and BPDU filter enabled
6	Dual Home	 Checks for: CLAG bonds that are not in dually connected state dually connected bonds have consistent VLAN and MTU configuration on both sides STP has consistent view of bonds' dual connectedness
7	Single Home	Checks for:singly connected bondsSTP has consistent view of bond's single connectedness
8	Conflicted Bonds	Checks for bonds in CLAG conflicted state and shows the reason
9	ProtoDown Bonds	Checks for bonds in protodown state and shows the reason

Test Number	Test Name	Description
10	SVI	 Checks if: an SVI is configured on both sides of a CLAG pair SVI on both sides have consistent MTU setting

Cumulus Linux Version Tests

The netq check cl-version command runs the following tests:

Test Number	Test Name	Description
0	Cumulus Linux Image Version	 Checks the following: no version specified, checks that all switches in the network have consistent version match-version specified, checks that a switch's OS version is equals the specified version min-version specified, checks that a switch's OS version is equal to or greater than the specified version

EVPN Validation Tests

The netq check evpn command runs the following tests to establish session sanity:

Test Number	Test Name	Description
0	EVPN BGP Session	Checks if:BGP EVPN sessions are establishedthe EVPN address family advertisement is consistent
1	EVPN VNI Type Consistency	Because a VNI can be of type L2 or L3, checks that for a given VNI, its type is consistent across the network
2	EVPN Type 2	Checks for consistency of IP-MAC binding and the location of a given IP-MAC across all VTEPs
3	EVPN Type 3	Checks for consistency of replication group across all VTEPs
4	EVPN Session	 For each EVPN session, checks if: adv_all_vni is enabled FDB learning is disabled on tunnel interface
5	Vlan Consistency	Checks for consistency of VLAN to VNI mapping across the network
6	Vrf Consistency	Checks for consistency of VRF to L3 VNI mapping across the network

Interface Validation Tests

The netq check interfaces command runs the following tests:

Test Number	Test Name	Description
0	Admin State	Checks for consistency of administrative state on two sides of a physical interface
1	Oper State	Checks for consistency of operational state on two sides of a physical interface
2	Speed	Checks for consistency of the speed setting on two sides of a physical interface
3	Autoneg	Checks for consistency of the auto-negotiation setting on two sides of a physical interface

License Validation Tests

The netq check license command runs the following test:

Test Number	Test Name	Description
0	License Validity	Checks for validity of license on all switches

LNV Validation Tests

The **netq check lnv** command checks for VXRD peer database, VXSND peer database, VNI operational state and head end replication list consistency.

Link MTU Validation Tests

The netq check mtu command runs the following tests:

Test Number	Test Name	Description
0	Link MTU Consistency	Checks for consistency of MTU setting on two sides of a physical interface

Test Number	Test Name	Description
1	VLAN interface	Checks if the MTU of an SVI is no smaller than the parent interface, substracting the VLAN tag size
2	Bridge interface	Checks if the MTU on a bridge is not arbitrarily smaller than the smallest MTU among its members

NTP Validation Tests

The netq check ntp command runs the following test:

Test Number	Test Name	Description
O	NTP Sync	Checks if the NTP service is running and in sync state

OSPF Validation Tests

The netq check ospf command runs the following tests to establish session sanity:

Test Number	Test Name	Description
0	Router ID	Checks for OSPF router ID conflicts in the network
1	Adjacency	Checks or OSPF adjacencies in a down or unknown state
2	Timers	Checks for consistency of OSPF timer values in an OSPF adjacency
3	Network Type	Checks for consistency of network type configuration in an OSPF adjacency

Test Number	Test Name	Description
4	Area ID	Checks for consistency of area ID configuration in an OSPF adjacency
5	Interface MTU	Checks for MTU consistency in an OSPF adjacency
6	Service Status	Checks for OSPF service health in an OSPF adjacency

Sensor Validation Tests

The netq check sensors command runs the following tests:

Test Number	Test Name	Description
0	PSU sensors	Checks for power supply unit sensors that are not in ok state
1	Fan sensors	Checks for fan sensors that are not in ok state
2	Temperature sensors	Checks for temperature sensors that are not in ok state

VLAN Validation Tests

The netq check vlan command runs the following tests:

Test Number	Test Name	Description
0	Link Neighbor VLAN Consistency	Checks for consistency of VLAN configuration on two sides of a port or a bond

Test Number	Test Name	Description
1	CLAG Bond VLAN Consistency	Checks for consistent VLAN membership of a CLAG bond on each side of the CLAG pair

VXLAN Validation Tests

The netq check vxlan command runs the following tests:

Test Number	Test Name	Description
0	VLAN Consistency	Checks for consistent VLAN to VXLAN mapping across all VTEPs
1	BUM replication	Checks for consistent replication group membership across all VTEPs

Validation Examples

This section provides validation examples for a variety of protocols and elements.

Perform a NetQ Agent Validation

The default validation confirms that the NetQ Agent is running on all monitored nodes and provides a summary of the validation results. This example shows the results of a fully successful validation.

cumulus@switch:~\$ netq check agents

Checked nodes: 12, Rotten nodes: 0

This example shows representative results when one or more of the NetQ Agents do not pass the validation check.

cumulus@switch:~\$ netq check agents

Checked nodes: 25, Rotten nodes: 1

Hostname Status Last Changed

leaf01 Rotten 8d:13h:34m:51s

Perform a BGP Validation

The default validation runs a network-wide BGP connectivity and configuration check on all nodes running the BGP service:

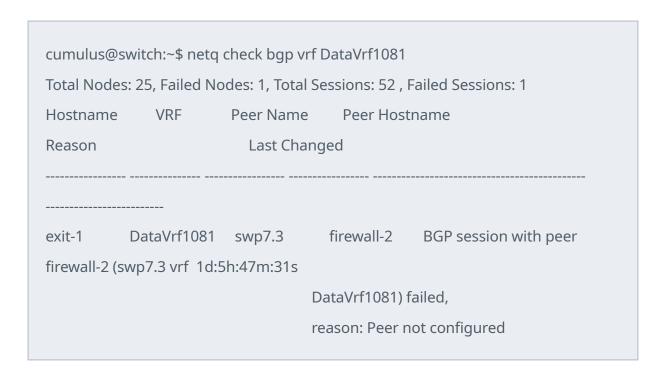
cumulus@switch:~\$ netq check bgp

Total Nodes: 15, Failed Nodes: 0, Total Sessions: 16, Failed Sessions: 0

This example indicates that all nodes running BGP and all BGP sessions are running properly. If there were issues with any of the nodes, NetQ would provide information about each node to aid in resolving the issues.

Perform a BGP Validation for a Particular VRF

Using the vrf <vrf> option of the netq check bgp command, you can validate the BGP service where communication is occurring through a particular virtual route. In this example, the VRF of interest is named *DataVrf1081*.



Perform a BGP Validation with Selected Tests

Using the include <bgp-number-range-list> and exclude <bgp-number-range-list> options, you can include or exclude one or more of the various checks performed during the validation. You can select from the following BGP validation tests:

Test Number	Test Name
0	Session Establishment
1	Address Families
2	Router ID

Refer to BGP Validation Tests for a description of these tests.

To include only the session establishment and router ID tests during a validation, run either of these commands:

cumulus@switch:~\$ netq check bgp include 0,2

cumulus@switch:~\$ netq check bgp exclude 1

Either way, a successful validation output would be similar to the following:

bgp check result summary:

Checked nodes : 8

Total nodes : 8

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Total Sessions : 30

Failed Sessions : 0

Session Establishment Test : passed,

Address Families Test : skipped

Router ID Test : passed,

Perform a BGP Validation and Output Results to JSON File

This example shows the default BGP validation results as it appears in a JSON file.

```
cumulus@switch:~$ netq check bgp json
{
  "tests":{
    "Session Establishment":{
      "errors":[
        {
           "hostname":"exit-1",
           "error":{
             "peerHostname":"firewall-2",
             "lastChanged":"Tue Jun 11 00:00:26 2019",
             "hostname":"exit-1",
             "peerName":"swp7.3",
             "reason":"BGP session with peer firewall-2 (swp7.3 vrf DataVrf1081)
failed, reason: Peer not configured",
             "vrf":"DataVrf1081"
           }
        },
      ],
      "enabled":true,
      "passed":false,
      "warnings":[
      ]
    },
    "Address Families":{
      "errors":[
           "hostname":"exit-1",
```

```
"error":{
             "peerHostname":"firewall-1",
             "lastChanged":"Sat Jun 1 03:34:10 2019",
             "hostname":"exit-1",
             "peerName":"swp6.3",
             "reason": "BGP session with peer firewall-1 swp6.3: AFI/SAFI evpn not
activated on peer",
             "vrf":"DataVrf1081"
          }
        },
      ],
      "enabled":true,
      "passed":false,
      "warnings":[
      ]
    },
    "Router ID":{
      "errors":[
      ],
      "enabled":true,
      "passed":true,
      "warnings":[
      ]
    }
  },
```

```
"failed_node_set":[
  "exit-1",
  "torc-12",
  "spine-1",
  "spine-3",
  "spine-2",
  "torc-21",
  "firewall-1"
],
"summary":{
  "total_cnt":14,
  "rotten_node_cnt":5,
  "failed_node_cnt":7,
  "warn_node_cnt":0,
  "checked_cnt":9,
  "total_sessions":174,
  "failed_sessions":42
},
"rotten_node_set":[
  "exit-2",
  "torc-22",
  "hostd-11",
  "torc-11",
  "firewall-2"
],
"warn_node_set":[
],
```

```
"validation":"BGP"
}
```

Perform a CLAG Validation

The default validation runs a network-wide CLAG connectivity and configuration check on all nodes running the CLAD service. This example shows results for a fully successful validation.

```
cumulus@switch:~$ netq check clag
clag check result summary:
Checked nodes : 4
Total nodes : 4
Rotten nodes : 0
Failed nodes : 0
Warning nodes : 0
Peering Test : passed,
Backup IP Test : passed,
Clag SysMac Test : passed,
VXLAN Anycast IP Test : passed,
Bridge Membership Test : passed,
Spanning Tree Test : passed,
Dual Home Test : passed,
Single Home Test : passed,
Conflicted Bonds Test : passed,
ProtoDown Bonds Test : passed,
SVI Test : passed,
```

This example shows representative results for one or more failures, warnings, or errors. In particular, you can see that you have duplicate system MAC addresses.

```
cumulus@switch:~$ netq check clag
clag check result summary:
```

Checked nodes : 4

Total nodes : 4

Rotten nodes : 0

Failed nodes : 2

Warning nodes : 0

Peering Test : passed,

Backup IP Test : passed,

Clag SysMac Test : 0 warnings, 2 errors,

VXLAN Anycast IP Test : passed,

Bridge Membership Test : passed,

Spanning Tree Test : passed,

Dual Home Test : passed,

Single Home Test : passed,

Conflicted Bonds Test : passed,

ProtoDown Bonds Test : passed,

SVI Test : passed,

Claq SysMac Test details:

Hostname Reason

leaf01 Duplicate sysmac with leaf02/None

leaf03 Duplicate sysmac with leaf04/None

Perform a CLAG Validation with Selected Tests

Using the include <clag-number-range-list> and exclude <clag-number-range-list> options, you can include or exclude one or more of the various checks performed during the validation. You can select from the following CLAG validation tests:

Test Number	Test Name
0	Peering
1	Backup IP
2	Clag Sysmac
3	VXLAN Anycast IP
4	Bridge Membership
5	Spanning Tree
6	Dual Home
7	Single Home
8	Conflicted Bonds
9	ProtoDown Bonds
10	SVI

Refer to CLAG Validation Tests for descriptions of these tests.

To include only the CLAG SysMAC test during a validation:

cumulus@switch:~\$ netq check clag include 2 clag check result summary:

Checked nodes : 4

Total nodes : 4

Rotten nodes : 0

Failed nodes : 2

Warning nodes : 0

Peering Test : skipped

Backup IP Test : skipped

Clag SysMac Test : 0 warnings, 2 errors,

VXLAN Anycast IP Test : skipped

Bridge Membership Test: skipped

Spanning Tree Test : skipped

Dual Home Test : skipped

Single Home Test : skipped

Conflicted Bonds Test : skipped

ProtoDown Bonds Test : skipped

SVI Test : skipped

Claq SysMac Test details:

Hostname Reason

leaf01 Duplicate sysmac with leaf02/None

leaf03 Duplicate sysmac with leaf04/None

To exclude the backup IP, CLAG SysMAC, and VXLAN anycast IP tests during a validation:

cumulus@switch:~\$ netq check clag exclude 1-3

clag check result summary:

Checked nodes : 4

Total nodes : 4

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Peering Test : passed,

Backup IP Test : skipped

Clag SysMac Test : skipped

VXLAN Anycast IP Test : skipped

Bridge Membership Test : passed,

Spanning Tree Test : passed,

Dual Home Test : passed,

Single Home Test : passed,

Conflicted Bonds Test : passed,

ProtoDown Bonds Test : passed,

SVI Test : passed,

Perform a Cumulus Linux Version Validation

The default validation (using no options) checks that all switches in the network have a consistent version.

cumulus@switch:/\$ netq check cl-version version check result summary:

Checked nodes : 12

Total nodes : 12

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Cumulus Linux Image Version Test : passed

Perform an EVPN Validation

The default validation runs a network-wide EVPN connectivity and configuration check on all nodes running the EVPN service. This example shows results for a fully successful validation.

cumulus@switch:~\$ netq check evpn

evpn check result summary:

Checked nodes : 6

Total nodes : 6

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Failed BGP Sessions: 0

Total Sessions : 16

Total VNIs : 3

EVPN BGP Session Test : passed,

EVPN VNI Type Consistency Test : passed,

EVPN Type 2 Test : passed,

EVPN Type 3 Test : passed,

EVPN Session Test : passed,

Vlan Consistency Test : passed,

Vrf Consistency Test : passed,

Perform an EVPN MAC Consistency Validation

Using the mac-consistency option, you can view any inconsistencies in the usage of MAC addresses in the EVPN overlay network.



/ IMPORTANT

The NetQ 2.3.x release is the last release that will support the mac-consistency option. However, this is equivalent to running only the EVPN Type 2 validation test. Refer to Perform an EVPN Validation with Selected Tests for details. As of Cumulus NetQ 2.4, the mac-consistency option will be removed.

cumulus@oob-mgmt-server:~\$ netq check evpn mac-consistency evpn check result summary:

Checked nodes : 6 Total nodes : 6

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Failed BGP Sessions: 0

Total Sessions : 16

Total VNIs : 3

EVPN BGP Session Test : passed,

EVPN VNI Type Consistency Test : passed,

EVPN Type 2 Test : passed,

EVPN Type 3 Test : passed,

EVPN Session Test : passed,

Vlan Consistency Test : passed,

Vrf Consistency Test : passed,

Perform an EVPN Validation for a Time in the Past

Using the **around** option, you can view the state of the EVPN service at a time in the past. Be sure to include the UOM.

cumulus@oob-mgmt-server:~\$ netq check evpn around 4d

evpn check result summary:

Checked nodes : 6

Total nodes : 6

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Failed BGP Sessions: 0

Total Sessions : 16

Total VNIs : 3

EVPN BGP Session Test : passed,

EVPN VNI Type Consistency Test : passed,

EVPN Type 2 Test : passed,

EVPN Type 3 Test : passed,

EVPN Session Test : passed,

Vlan Consistency Test : passed,

Vrf Consistency Test : passed,

Perform an EVPN Validation with Selected Tests

Using the include <evpn-number-range-list> and exclude <evpn-number-range-list> options, you can include or exclude one or more of the various checks performed during the validation. You can select from the following EVPN validation tests:

Test Number	Test Name
0	EVPN BGP Session
1	EVPN VNI Type Consistency
2	EVPN Type 2
3	EVPN Type 3
4	EVPN Session

Test Number	Test Name
5	Vlan Consistency
6	Vrf Consistency

Refer to EVPN Validation Tests for descriptions of these tests.

To run only the EVPN Type 2 test:

cumulus@switch:~\$ netq check evpn include 2

evpn check result summary:

Checked nodes : 6

Total nodes : 6

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Failed BGP Sessions: 0

Total Sessions : 0

Total VNIs : 3

EVPN BGP Session Test : skipped

EVPN VNI Type Consistency Test : skipped

EVPN Type 2 Test : passed,

EVPN Type 3 Test : skipped

EVPN Session Test : skipped

Vlan Consistency Test : skipped

Vrf Consistency Test : skipped

To exclude the BGP session and VRF consistency tests:

cumulus@switch:~\$ netq check evpn exclude 0,6 evpn check result summary:

Checked nodes : 6

Total nodes : 6

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Failed BGP Sessions: 0

Total Sessions : 0

Total VNIs : 3

EVPN BGP Session Test : skipped

EVPN VNI Type Consistency Test : passed,

EVPN Type 2 Test : passed,

EVPN Type 3 Test : passed,

EVPN Session Test : passed,

Vlan Consistency Test : passed,

Vrf Consistency Test : skipped

To run only the first five tests:

cumulus@switch:~\$ netq check evpn include 0-4 evpn check result summary:

Checked nodes : 6

Total nodes : 6

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Failed BGP Sessions: 0

Total Sessions : 16

Total VNIs : 3

EVPN BGP Session Test : passed,

EVPN VNI Type Consistency Test : passed,

EVPN Type 2 Test : passed,

EVPN Type 3 Test : passed,

EVPN Session Test : passed,

Vlan Consistency Test : skipped

Vrf Consistency Test : skipped

Perform an Interfaces Validation

The default validation runs a network-wide connectivity and configuration check on all interfaces. This example shows results for a fully successful validation.

cumulus@switch:~\$ netq check interfaces

interface check result summary:

Checked nodes : 12

Total nodes : 12

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Unverified Ports: 56

Checked Ports: 108

Failed Ports : 0

Admin State Test : passed,

Oper State Test : passed,

Speed Test : passed,

Autoneg Test : passed,

Perform an Interfaces Validation for a Time in the Past

Using the **around** option, you can view the state of the interfaces at a time in the past. Be sure to include the UOM.

cumulus@oob-mgmt-server:~\$ netq check interfaces around 6h interface check result summary:

Checked nodes : 12

Total nodes : 12

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Unverified Ports: 56

Checked Ports : 108

Failed Ports : 0

Admin State Test : passed,

Oper State Test : passed,

Speed Test : passed,

Autoneg Test : passed,

Perform an Interfaces Validation with Selected Tests

Using the include <interface-number-range-list> and exclude <interface-number-range-list> options, you can include or exclude one or more of the various checks performed during the validation. You can select from the following interface validation tests:

Test Number	Test Name
0	Admin State
1	Oper State
2	Speed
3	Autoneg

Refer to Interface Validation Tests for descriptions of these tests.

Perform a License Validation

You can also check for any nodes that have invalid licenses without going to each node. Because switches do not operate correctly without a valid license you might want to verify that your Cumulus Linux licenses on a regular basis.

This example shows that all licenses on switches are valid.

cumulus@oob-mgmt-server:~\$ netq check license

license check result summary:

Checked nodes : 12

Total nodes : 12

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Checked Licenses: 8

Failed Licenses : 0

License validity Test : passed,



This command checks every node, meaning every switch and host in the network. Hosts do not require a Cumulus Linux license, so the number of licenses checked might be smaller than the total number of nodes checked.

Perform a Link MTU Validation

The default validate verifies that all corresponding interface links have matching MTUs. This example shows no mismatches.

cumulus@switch:~\$ netq check mtu

mtu check result summary:

Checked nodes : 12

Total nodes : 12

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Warn Links : 0

Failed Links : 0

Checked Links : 196

Link MTU Consistency Test : passed,

VLAN interface Test : passed,

Bridge interface Test : passed,

Perform an NTP Validation

The default validation checks for synchronization of the NTP server with all nodes in the network. It is always important to have your devices in time synchronization to ensure configuration and management events can be tracked and correlations can be made between events.

This example shows that server04 has an error.

cumulus@switch:~\$ netq check ntp
ntp check result summary:

Checked nodes : 12
Total nodes : 12
Rotten nodes : 0
Failed nodes : 1
Warning nodes : 0

Additional summary:
Unknown nodes : 0

NTP Servers : 3

NTP Sync Test : 0 warnings, 1 errors,

NTP Sync Test details:

Hostname NTP Sync Connect Time

server04 no 2019-09-17 19:21:47

Perform an OSPF Validation

The default validation runs a network-wide OSPF connectivity and configuration check on all nodes running the OSPF service. This example shows results several errors in the Timers and Interface MTU tests.

cumulus@switch:~# netq check ospf

Checked nodes: 8, Total nodes: 8, Rotten nodes: 0, Failed nodes: 4, Warning

nodes: 0, Failed Adjacencies: 4, Total Adjacencies: 24

Router ID Test : passed

Adjacency Test : passed

Timers Test : 0 warnings, 4 errors

Network Type Test : passed

Area ID Test : passed

Interface Mtu Test : 0 warnings, 2 errors

Service Status Test : passed

Timers Test details:

Hostname Interface PeerID Peer IP

Reason Last Changed

spine-1 downlink-4 torc-22 uplink-1 dead time

mismatch Mon Jul 1 16:18:33 2019

spine-1 downlink-4 torc-22 uplink-1 hello time

mismatch Mon Jul 1 16:18:33 2019

torc-22 uplink-1 spine-1 downlink-4 dead time

mismatch Mon Jul 1 16:19:21 2019

torc-22 uplink-1 spine-1 downlink-4 hello time

mismatch Mon Jul 1 16:19:21 2019

Interface Mtu Test details:

Hostname Interface PeerID Peer IP

Reason Last Changed

spine-2 downlink-6 0.0.0.22 27.0.0.22 mtu

mismatch Mon Jul 1 16:19:02 2019

tor-2 uplink-2 0.0.0.20 27.0.0.20 mtu

mismatch Mon Jul 1 16:19:37 2019

Perform a Sensors Validation

Hardware platforms have a number sensors to provide environmental data about the switches. Knowing these are all within range is a good check point for maintenance.

For example, if you had a temporary HVAC failure and you are concerned that some of your nodes are beginning to overheat, you can run this validation to determine if any switches have already reached the maximum temperature threshold.

cumulus@switch:~\$ netq check sensors

sensors check result summary:

Checked nodes : 8

Total nodes : 8

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Checked Sensors : 136

Failed Sensors : 0

PSU sensors Test : passed,

Fan sensors Test : passed,

Temperature sensors Test : passed,

Perform a VLAN Validation

Validate that VLANS are configured and operating properly:

cumulus@switch:~\$ netq check vlan

vlan check result summary:

Checked nodes : 12

Total nodes : 12

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Additional summary:

Failed Link Count : 0

Total Link Count : 196

Link Neighbor VLAN Consistency Test : passed,

Clag Bond VLAN Consistency Test : passed,

Perform a VXLAN Validation

Validate that VXLANs are configured and operating properly:

cumulus@switch:~\$ netq check vxlan vxlan check result summary:

Checked nodes : 6

Total nodes : 6

Rotten nodes : 0

Failed nodes : 0

Warning nodes : 0

Vlan Consistency Test : passed,

BUM replication Test : passed,



Both asymmetric and symmetric VXLAN configurations are validated with this command.